

Chapter 2: Bits and Bytes II

Topics

- **Byte Ordering: Little vs Big Endian**
 - Ints
 - Pointers
 - Characters
 - Strings
- **Bit-Level Operations in C**
 - $\&$, $|$, \sim , \wedge
 - Bit Masking
 - Logical Expressions

Announcements

- **Data Lab is due Friday Sept 12 by 11:55 pm**
 - TAs may offer extra office hours later in the week
- **Recitation Exercise #1 is available on moodle and is due Monday Sept 8**
 - Print and hand in a hard copy at the beginning of recitation
 - These problems are useful in helping to study for the midterm
- **Essential that you read the textbook in detail & do the practice problems**
 - Read Chapter 2 (2.1-2.3 this week)
 - Next week: begin Chapter 3

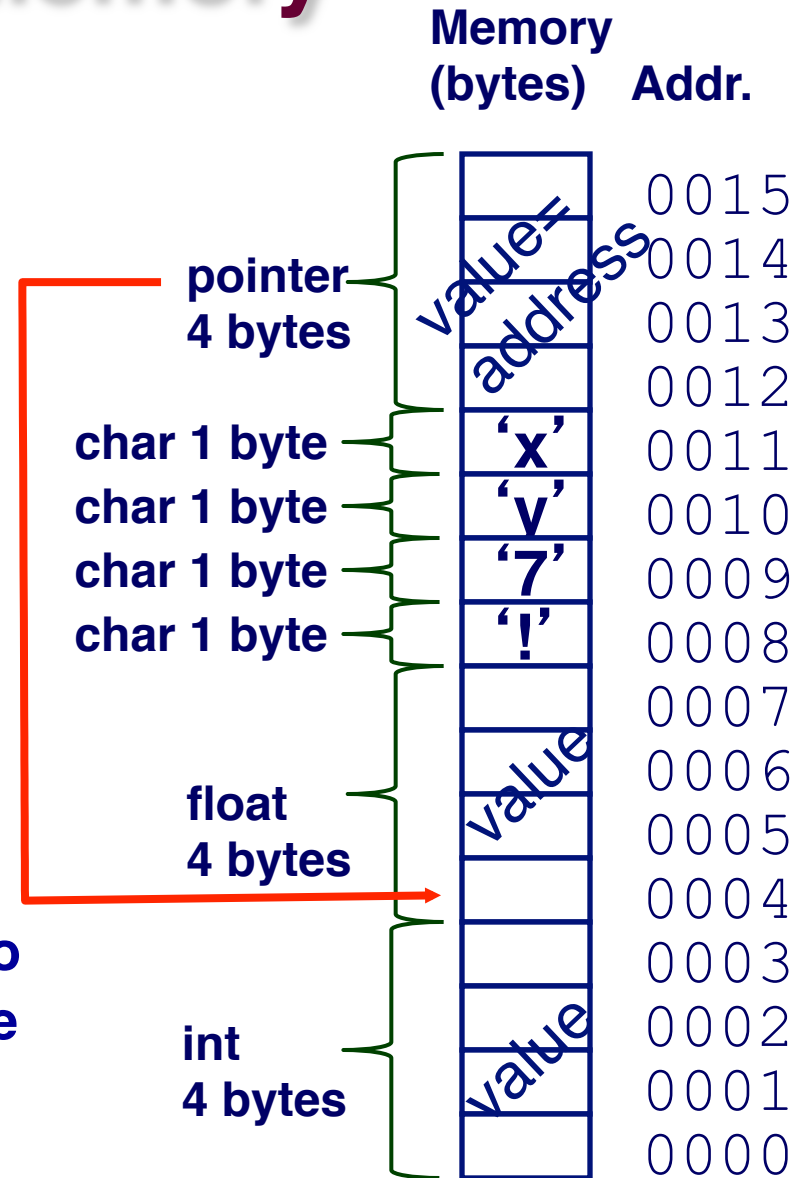
Recap...

- **Binary representations – base 2**
- **Binary digital logic**
- **Hexadecimal representations – base 16**
- **Byte-addressed memory**
- **Representing data in C**
 - **Ints, shorts, floats, doubles, chars, etc.**
- **Pointers in C**

Recap: Byte-based Memory

IA32 Example:

- Address of int is 0x00000000
- Address of float is 0x00000004
- Address of character = '7' is 0x00000009
- Address of pointer is 0x0000000c
 - Note: the pointer points to another memory location, i.e. stores a memory location *address*
e.g. if pointer = 0x00000004 it means the pointer is pointing to the float! (actually the first byte of the float)



Byte Ordering

In what order should bytes within a multi-byte word be stored in memory?

- Consider a 32-bit integer or int (4 bytes)

0000 0001 0010 0011 0100 0101 0110 0111₂

= 0x 0 1 2 3 4 5 6 7

= 0x 01 23 45 67

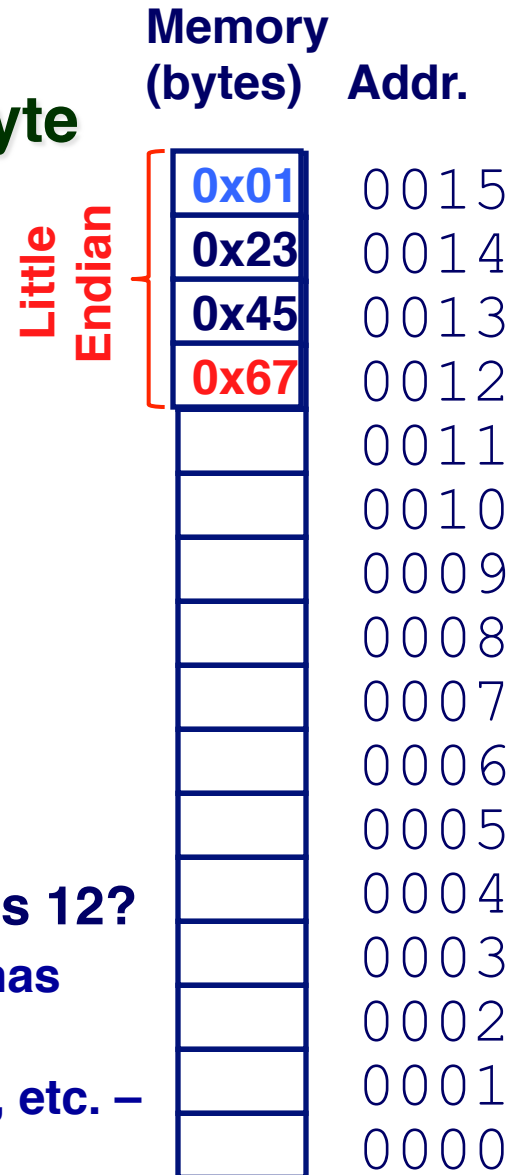
Most
Significant
Byte (MSB)

Least
Significant
Byte (LSB)

- In what order do we store these four bytes

0x01234567 in memory, say at starting address 12?

- Little Endian approach:** Least significant byte has lowest address (12), followed by the next least significant byte in the next lowest address (13), etc. – i.e. “little end first”



Byte Ordering Example

Big Endian

- Most significant byte has lowest address, i.e. “big end first”

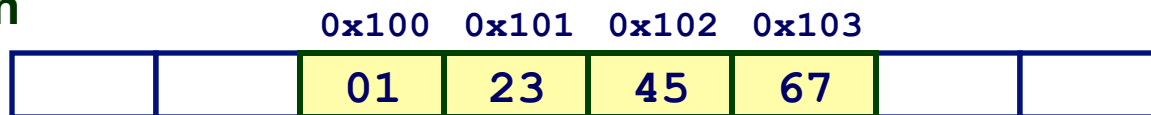
Little Endian

- Least significant byte has lowest address

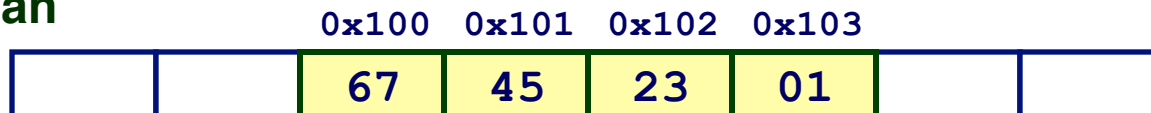
Example

- Variable `x` has 4-byte representation `0x01234567`
- Address given by `&x` is `0x100`

Big Endian



Little Endian



Byte Ordering in Hardware

Aside: Endianness came from Gulliver's Travels

CPU Conventions

- Sun SPARC, Motorola 68K, Power PC (PPC's) are “Big Endian” machines
 - Most significant byte has lowest address
- ISA32 are “Little Endian” machines
 - Least significant byte has lowest address
- Some are “bi-endian” (hardware supports both types of Endianness, which can improve performance)
 - MIPS, Alpha, ARM

Reading Byte-Reversed Listings

Disassembly

- Text representation of binary machine code
- Generated by program that reads the machine code

Example Fragment

Address	Instruction Code	Assembly Rendition
8048365:	5b	pop %ebx
8048366:	81 c3 ab 12 00 00	add \$0x12ab,%ebx
804836c:	83 bb 28 00 00 00 00	cmpl \$0x0,0x28(%ebx)

Deciphering Numbers

- Value: 0x12ab
- Pad to 4 bytes: 0x000012ab
- Split into bytes: 00 00 12 ab
- Reverse: ab 12 00 00

So this is
Little Endian

Examining Data Representations

Code to Print Byte Representation of Data

- Casting pointer to unsigned char * creates byte array

```
typedef unsigned char *pointer;  
  
void show_bytes(pointer start, int len)  
{  
    int i;  
    for (i = 0; i < len; i++)  
        printf("0x%p\t0x%.2x\n",  
               start+i, start[i]);  
    printf("\n");  
}
```

equivalent to ***(start+i)**

Printf directives:

%p: Print pointer

%x: Print Hexadecimal

Use show_bytes() to find if your machine is Little Endian or Big Endian

show_bytes Execution Example

```
int a = 15213;           // = 0x3B6D
printf("int a = 15213;\n");
show_bytes((pointer) &a, sizeof(int));
```

Result (Linux):

```
int a = 15213;
0x11ffffcb8    0x6d
0x11ffffcb9    0x3b
0x11ffffcba    0x00
0x11ffffcbb    0x00
```

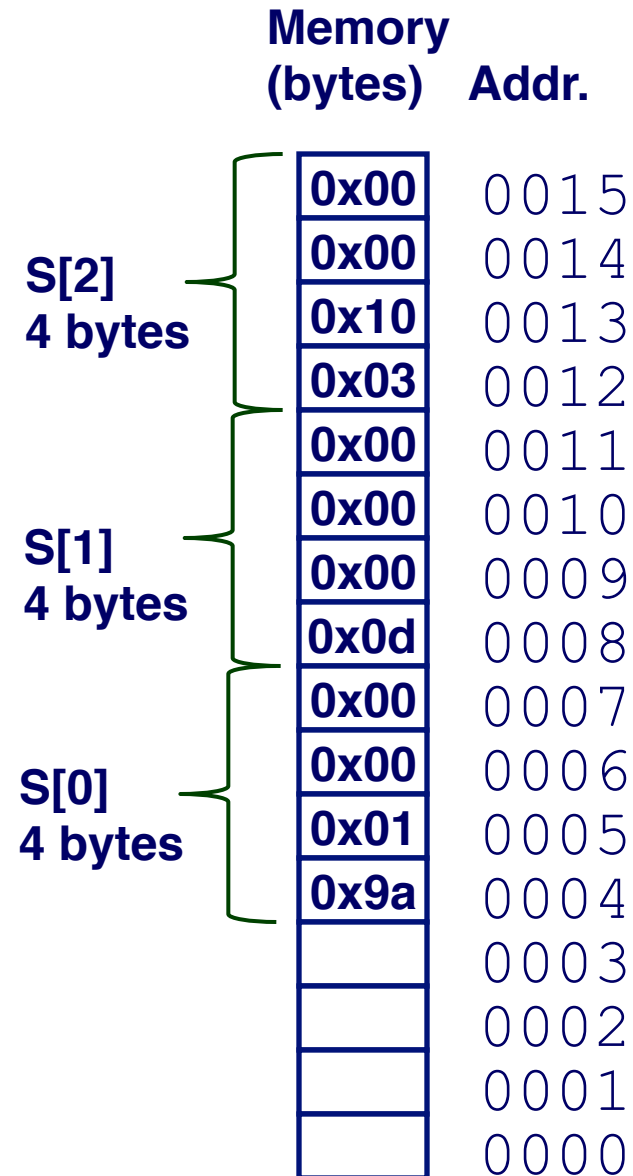
Little Endian with zero padding in most significant bytes

Endianness and Arrays

Integer Array Example:

```
int S[3];  
S[0] = 410;          /* = 0x0000019a */  
S[1] = 14;           /* = 0x0000000d */  
S[2] = 4099;         /* = 0x00001003 */
```

- Let the int array S be stored starting at memory address 4
- Lowest element of array S[0] is stored starting at the lowest address of the array (byte 4)
 - Next lowest element of array S[1] is stored at the next lowest memory address of int array (byte 8). And so on...
- Note how each int in the array is stored according to byte ordering rules (little endian)



Endianness and Pointers

```
int B = -15213;  
int *P = &B;
```

Sun Address (32-bit)

Hex: E F F F F B 2 C
Binary: 1110 1111 1111 1111 1111 1011 0010 1100

Big Endian Sun P

2C	High mem
FB	
FF	
EF	Low mem

Linux Address (32-bit)

Hex: B F F F F 8 D 4
Binary: 1011 1111 1111 1111 1111 1000 1101 0100

Little Endian Linux P

BF
FF
F8
D4

Alpha P

X64 Address (64-bit, higher order hex not shown)

Hex: 1 F F F F F C A 0
Binary: 0001 1111 1111 1111 1111 1111 1100 1010 0000

00
00
00
01
FF
FF
FC
A0

Different compilers & machines assign different locations to objects, i.e. not only are the addresses stored in the pointer of different widths, but they are also of different values (locations)

Endianness and Floats

Float F = 15213.0;

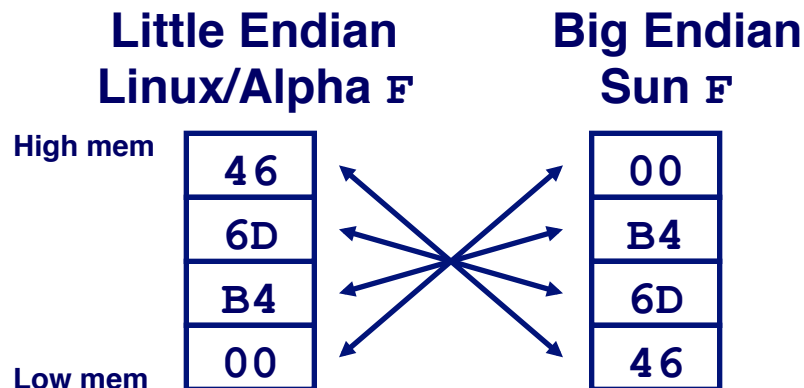
IEEE Single Precision Floating Point Representation

Hex:	4	6	6	D	B	4	0	0
Binary:	0100	0110	0110	1101	1011	0100	0000	0000
15213:			1110	1101	1011	01		



Not same as integer representation, but consistent across machines

Can see some relation to integer representation, but not obvious



Representing Characters

(hex)

- Actually unsigned characters, via ASCII (American Standard Code for Information Exchange)
- The alphabet, numbers, punctuation, and symbols are encoded via an 8-bit ASCII table:
- Example: numbers start with '0' = 0x30, capital letters start at 'A' = 0x41 = 65, lower case letter 'a' = 0x61
- Endianness doesn't affect character representations

	0	1	2	3	4	5	6	7
0	NUL	DLE	space	0	@	P	`	p
1	SOH	DC1 XON	!	1	A	Q	a	q
2	STX	DC2	"	2	B	R	b	r
3	ETX	DC3 XOFF	#	3	C	S	c	s
4	EOT	DC4	\$	4	D	T	d	t
5	ENQ	NAK	%	5	E	U	e	u
6	ACK	SYN	&	6	F	V	f	v
7	BEL	ETB	'	7	G	W	g	w
8	BS	CAN	(8	H	X	h	x
9	HT	EM)	9	I	Y	i	y
A	LF	SUB	*	:	J	Z	j	z
B	VT	ESC	+	;	K	[k	{
C	FF	FS	,	<	L	\	l	
D	CR	GS	-	=	M]	m	}
E	SO	RS	.	>	N	^	n	~
F	SI	US	/	?	O	_	o	del

Representing Strings

Strings in C

```
char S[6] = "15213";
```

- Represented by array of characters
- Each character encoded in ASCII format

- '1' = 0x31

- '5' = 0x35

- '2' = 0x32

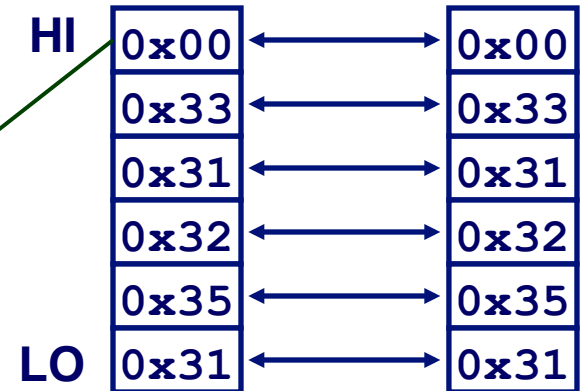
- '1' = 0x31

- '3' = 0x33

- String should be null-terminated

- Final character = 0

Little Endian Big Endian
Linux/Alpha s Sun s



Compatibility

- Byte ordering not an issue
 - Data are single byte quantities
- Text files generally platform independent
 - Except for different conventions of line termination character(s)!

Machine-Level Code Representation

Encode Program as Sequence of Instructions

- Each simple operation
 - Arithmetic operation
 - Read or write memory
 - Conditional branch
- Instructions encoded as bytes
 - Alpha's, Sun's, old Mac's use 4 byte instructions
 - » Reduced Instruction Set Computer (RISC)
 - PC's new Mac's use variable length instructions
 - » Complex Instruction Set Computer (CISC)
- Different instruction types and encodings for different machines
 - Most code not binary compatible

Programs are Byte Sequences Too!

Representing Instructions

```
int sum(int x, int y)
{
    return x+y;
}
```

- For this example, Alpha & Sun use two 4-byte instructions
 - Use differing numbers of instructions in other cases
- PC uses 7 instructions with lengths 1, 2, and 3 bytes
 - Same for NT and for Linux
 - NT / Linux not fully binary compatible

Alpha sum

00
00
30
42
01
80
FA
6B

Sun sum

81
C3
E0
08
90
02
00
09

PC sum

55
89
E5
8B
45
0C
03
45
08
89
EC
5D
C3

Different machines use totally different instructions and encodings

Relations Between Logic Operations

DeMorgan's Laws

- Express & in terms of |, and vice-versa
 - $A \& B = \sim(\sim A | \sim B)$
 - » A and B are true if and only if neither A nor B is false
 - $A | B = \sim(\sim A \& \sim B)$
 - » A or B are true if and only if A and B are not both false

Exclusive-Or using Inclusive Or

- $A \wedge B = (\sim A \& B) | (A \& \sim B)$
 - » Exactly one of A and B is true
 - » This is Shannon's circuit.
- $A \wedge B = (A | B) \& \sim(A \& B)$
 - » Either A is true, or B is true, but not both

General Boolean Algebras

Operate on Bit Vectors

■ Operations applied bitwise

$\&$	$\begin{array}{c} 01101001 \\ 01010101 \\ \hline 01000001 \end{array}$	$\begin{array}{c} 01101001 \\ \\ 01010101 \\ \hline 01111101 \end{array}$	$\begin{array}{c} 01101001 \\ ^\wedge \\ 01010101 \\ \hline 00111100 \end{array}$	$\begin{array}{c} 01101001 \\ \sim \\ 01010101 \\ \hline 10101010 \end{array}$
------	--	---	---	--

All of the Properties of Boolean Algebra Apply

Using Boolean Operators for Representing & Manipulating Sets

Representation

- Width w bit vector represents subsets of $\{0, \dots, w-1\}$

- $a_j = 1$ if $j \in A$

01101001

{ 0, 3, 5, 6 }

7 6 5 4 3 2 1 0

01010101

{ 0, 2, 4, 6 }

7 6 5 4 3 2 1 0

Operations

- & Intersection 01000001 { 0, 6 }
- | Union 01111101 { 0, 2, 3, 4, 5, 6 }
- ^ Symmetric difference 00111100 { 2, 3, 4, 5 }
- ~ Complement 10101010 { 1, 3, 5, 7 }

Bit-Level Operations in C

- Operations **&**, **|**, **~**, **^** Available in C

- Apply to any “integral” data type

- long, int, short, char

- View arguments as bit vectors

- Arguments applied bit-wise

- Examples (Char data type)

- $\sim 0x41 \rightarrow 0xBE$

- $\sim 01000001_2 \rightarrow 10111110_2$

- $\sim 0x00 \rightarrow 0xFF$

- $\sim 00000000_2 \rightarrow 11111111_2$

- $0x69 \ \& \ 0x55 \rightarrow 0x41$

- $01101001_2 \ \& \ 01010101_2 \rightarrow 01000001_2$

- $0x69 \ | \ 0x55 \rightarrow 0x7D$

- $01101001_2 \ | \ 01010101_2 \rightarrow 01111101_2$

Bit Masking

- **Example: Mask out the 4 least significant bits of a byte 0x69**
 - Let mask = 0xF0
 - $0x69 \ \& \ 0xF0 \ \rightarrow \ 0x60$
 - $01101001_2 \ \& \ 11110000_2 \ \rightarrow \ 01100000_2$
- **Example: Mask out all but the most significant bit of a byte 0x69**
 - Let mask = 0x80
 - $0x69 \ \& \ 0x80 \ \rightarrow \ 0x00$
 - $01101001_2 \ \& \ 10000000_2 \ \rightarrow \ 00000000_2$

Logical vs Bitwise Operations in C

Logical Operators

- `&&` (AND), `||` (OR), `!` (NOT or “bang”)
 - View 0 as “False”
 - Anything nonzero as “True”
 - Always return 0 or 1
 - Early termination

Example code:

```
■ int x,y,z;  
  
....  
if ( !((x==0) && (x>y)) || (z<256) ) {  
    ...  
    z = ~(x & y) | z;  
}
```

Each logical expression is either TRUE (1) or FALSE (0):

`(x==0)`

`(x>y)`

`((x==0) && (x>y))`

`!((x==0) && (x>y))`

`(z<256)`

`!((x==0) && (x>y)) || (z<256)`

Compare to the bit-wise logical operations, where input, intermediate, and final values don't have to be 0 or 1

By the way, parentheses are your friend, and good programming practice