

Floating Point Introduction

Topics

- Chapter 2.4
- IEEE Floating Point Standard

Standard Decimal Scientific Notation

- Real numbers expressed as $x \cdot 10^y$
 - e.g. $4.782 \cdot 10^{27}$, and $-1.396 \cdot 10^{-17}$, or $7.088\text{e-}6$, or $3.14\text{E}10$

- Expansion:

- $4.782 \cdot 10^{27} = 4 \cdot 10^{27} + 7 \cdot 10^{26} + 8 \cdot 10^{25} + 2 \cdot 10^{24}$

$$\begin{aligned} \text{decimal} &= d_m d_{m-1} \dots d_1 d_0 . d_{-1} d_{-2} \dots d_{-n} \\ &= \sum_{i=-n}^m d_i \cdot 10^i \end{aligned}$$

- Not all numbers can be expressed exactly in base 10
 - e.g. $1/3 = 0.33333\dots$, so it must be approximated
- Our goal is to represent real numbers using binary
 - We follow the approach of decimal scientific notation except using base 2

IEEE Floating Point

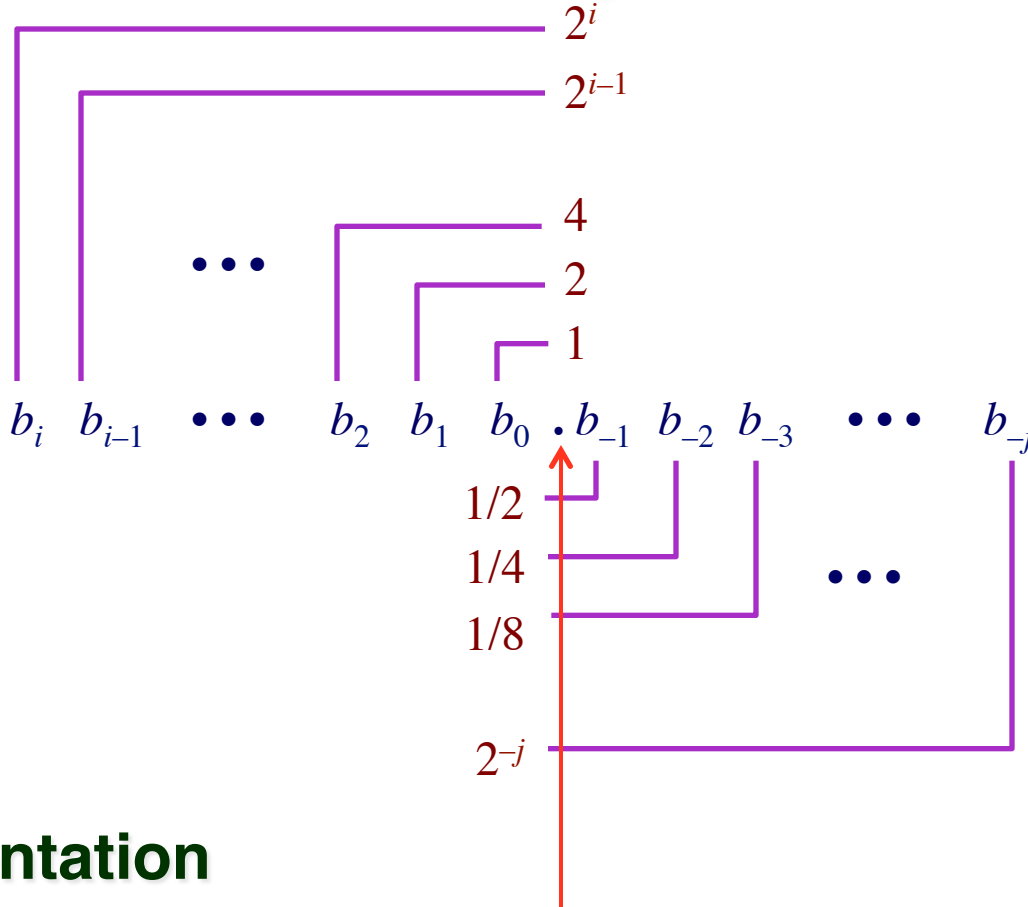
IEEE Standard 754

- Established in 1985 as uniform standard for floating point arithmetic
 - Before that, many idiosyncratic formats
- Supported by all major CPUs

Driven by Numerical Concerns

- Nice standards for rounding, overflow, underflow
- Hard to make go fast
 - Numerical analysts predominated over hardware types in defining standard

Fractional Binary Numbers



Representation

- Bits to right of “**binary point**” represent fractional powers of 2
- Represents rational number:

$$\sum_{k=-j}^i b_k \cdot 2^k$$

Fractional Binary Number Examples

Value

Representation

5 3/4

$$101.11_2 = 2^2 + 2^0 + 2^{-1} + 2^{-2}$$

2 7/8

$$10.111_2 = 2^1 + 2^{-1} + 2^{-2} + 2^{-3}$$

63/64

$$0.111111_2 = 2^{-1} + 2^{-2} + 2^{-3} + 2^{-4} + 2^{-5} + 2^{-6}$$

Observations

- Divide by 2 by shifting right
- Multiply by 2 by shifting left
- Numbers of form $0.111111\dots_2$ just below 1.0
 - $1/2 + 1/4 + 1/8 + \dots + 1/2^i + \dots \rightarrow 1.0$
 - Use notation $1.0 - \epsilon$

Representable Numbers

Limitation

- This style of binary point notation is not very good at representing larger numbers

e.g. 10100...
(followed by a 100 zeros)

- Can only exactly represent numbers of the form $x/2^k$
- Other numbers have repeating bit representations

Value

Representation

1/3

0.0101010101 [01] ...₂

1/5

0.001100110011 [0011] ...₂

1/10

0.0001100110011 [0011] ...₂

Floating Point Representation

Numerical Form

- Real number = $(-1^s) * M * 2^E$
 - Sign bit s determines whether number is negative or positive
 - Significand M normally a fractional value in range $[1.0, 2.0)$.
 - Exponent E weights value by power of two

Encoding



- MSB is sign bit
- exp field encodes E
- frac field encodes M

Floating Point Precisions



Encoding

- MSB is sign bit
- `exp` field encodes E
- `frac` field encodes M

$$\text{Real number} = (-1^s) * M * 2^E$$

Sizes

- Single precision: 8 `exp` bits, 23 `frac` bits
 - 32 bits total. Can represent from 2^{127} (1.7e38) down to 2^{-126}
- Double precision: 11 `exp` bits, 52 `frac` bits
 - 64 bits total
- Extended precision: 15 `exp` bits, 63 `frac` bits
 - Only found in Intel-compatible machines
 - Stored in 80 bits - 1 bit wasted
- Quad Precision (IEEE 754r - revised) - 15 `exp`, 112 `frac`
 - 128 bits total
- Half Precision (IEEE 754r) - 5 `exp`, 10 `frac`
 - 16 bits total