

Chapter 3:

Assembly Language Programming I

Topics

- Assembly Programmer's Execution Model
- Accessing Information
 - Registers
 - Memory
- Arithmetic operations

Announcements

- **Data Lab is due Friday Sept 12 by 11:55 pm**
 - Sign up for grading time slots with your TA
- **Prof. Han traveling next Tuesday to Washington DC**
 - Guest lecture next Tuesday by TA Yogesh Virkar
 - No office hours next Tuesday, but will be back for Wednesday office hours 3-4 pm
- **First midterm probably the week of Oct 6**
- **Essential that you read the textbook in detail & do the practice problems**
 - Read Chapter 3.1-3.14, skip 3.12 for now

Chapter Mapping

Chapter 3

Lines of
Source
code

Pre-processor
&
Compiler

add a,b
sub a,b
move a...

Lines of
Assembly
code

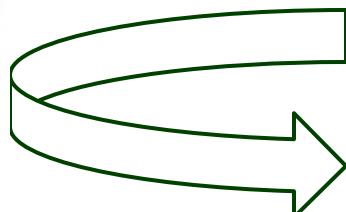
Assembler
&
Linker

Chapter 2

Operating
System

10101010

Memory



10101010
00001010
01010111

Lines of
Binary
code &
data

11101110
00111011
10000111

Chapters 3, 4 and 6

Intel x86 Processors

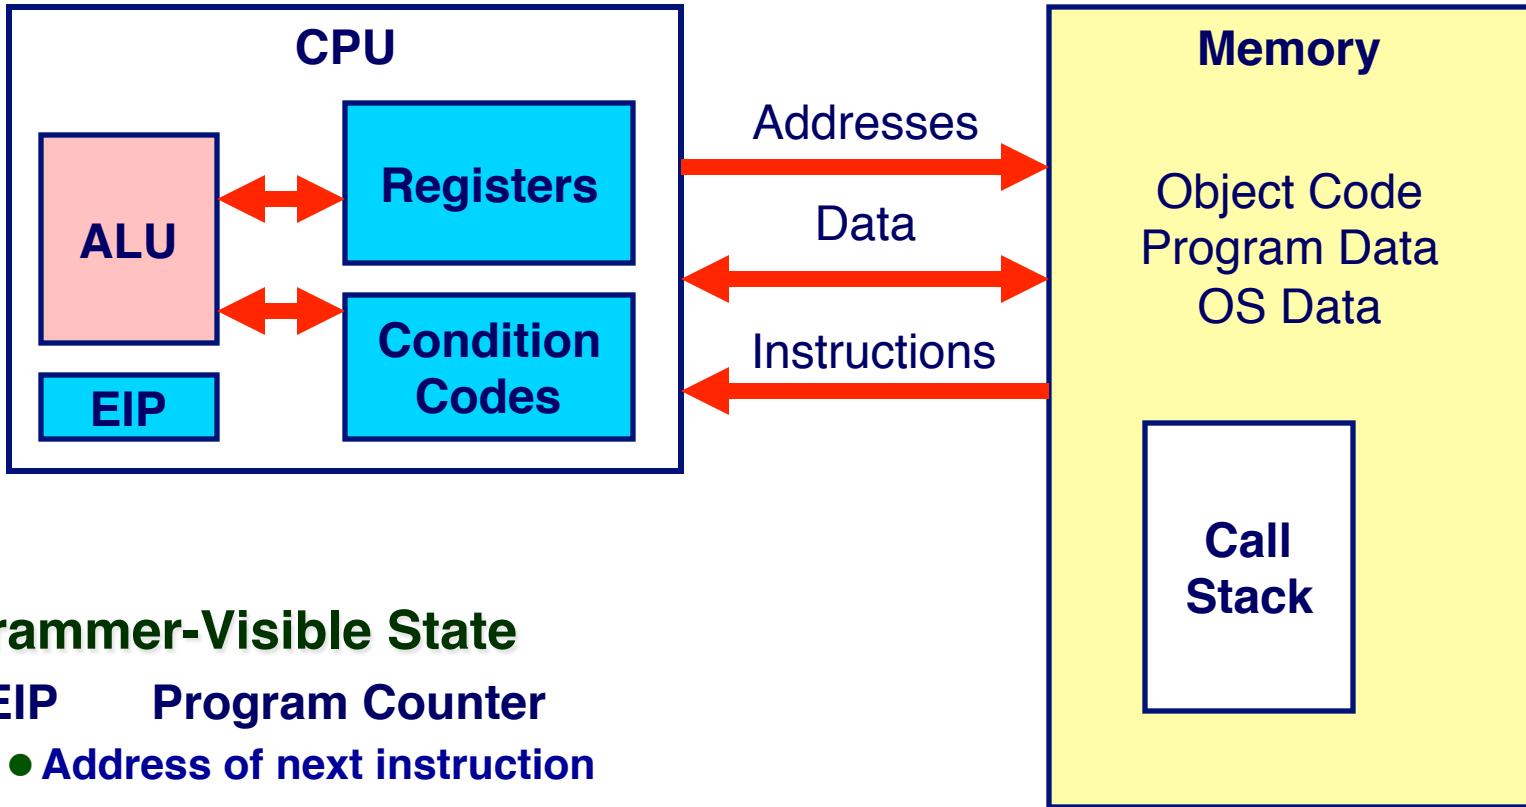
- **Totally dominate PC computer market**
- **Evolutionary design**
 - Backwards compatible up until 8086 16-bit CPU, introduced in 1978
 - Then 80286, 80386, 80486, Pentium, ..., Intel Core i7 – hence the name x86
 - Added more features as time goes on
- **Complex instruction set computer (CISC)**
 - Many different instructions with many different formats
 - But, only small subset encountered with Linux programs
 - Hard to match performance of Reduced Instruction Set Computers (RISC)
 - But, Intel has done just that!



Assembly Language

- **Specific to a CPU**
 - We will be using x86 assembly language
 - ARM processors will have a different assembly language, etc.
 - 32-bit processors will have different assembly language than 64-bit processors
- **Different styles for IA32 assembly code**
 - We will be using gcc/GNU-style assembly language
 - There is also the Intel style of assembly language
 - Example: switches order of source and destination compared to gcc/GNU

Assembly Programmer's View



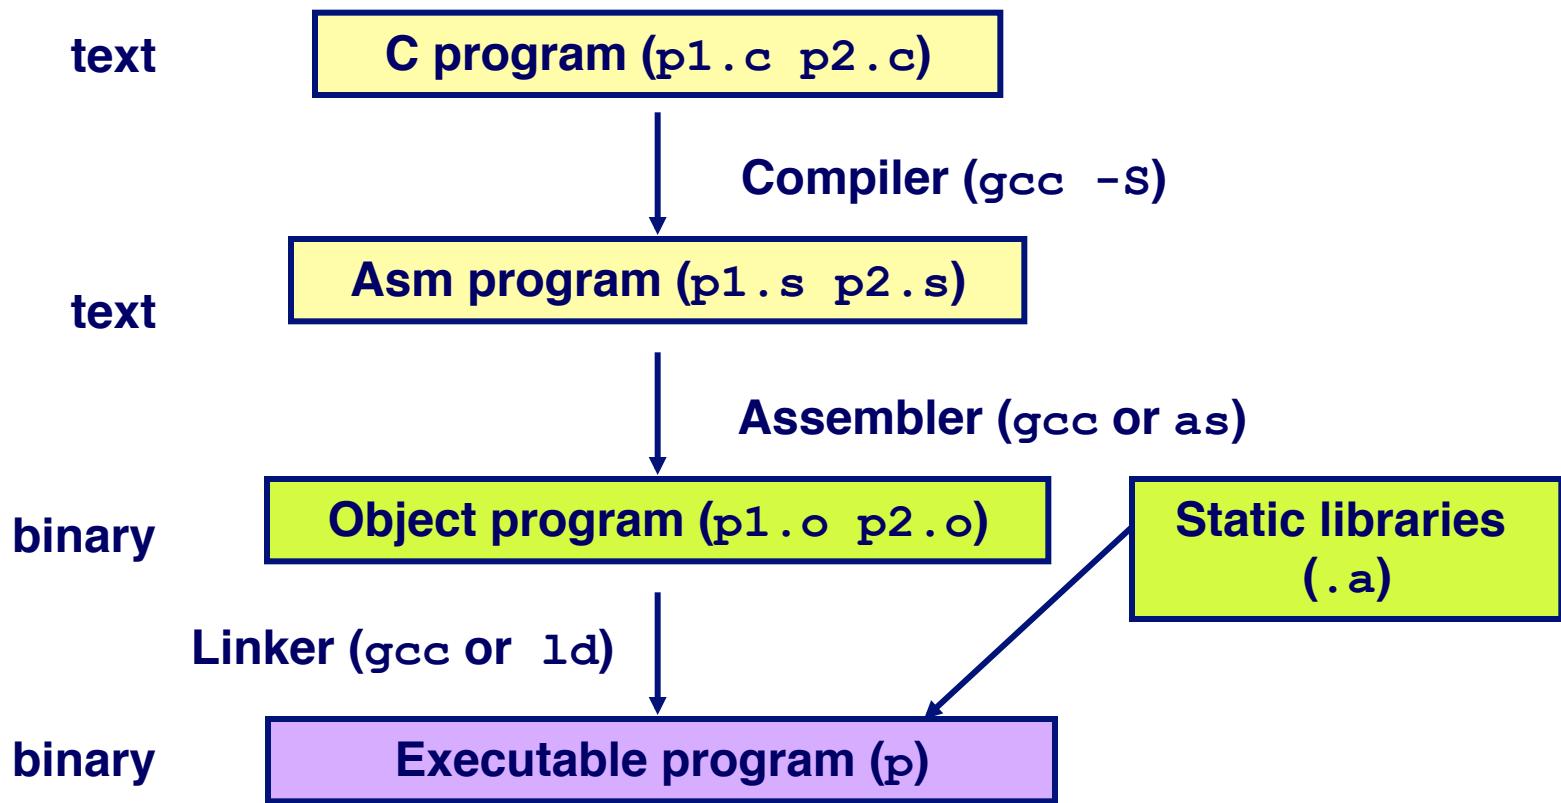
Programmer-Visible State

- **EIP Program Counter**
 - Address of next instruction
- **Register File**
 - Heavily used program data
- **Condition Codes**
 - Store status information about most recent arithmetic operation
 - Used for conditional branching

- **Memory**
 - Byte addressable array
 - Code, user data, (some) OS data
 - Includes call stack used to support procedures

Turning C into Object Code

- Code in files p1.c p2.c
- Compile with command: gcc -O p1.c p2.c -o p
 - Use optimizations (-O)
 - Put resulting binary in file p



Compiling Into Assembly

C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

Generated Assembly

```
_sum:
    pushl %ebp
    movl %esp,%ebp
    movl 12(%ebp),%eax
    addl 8(%ebp),%eax
    movl %ebp,%esp
    popl %ebp
    ret
```

Obtain with command

```
gcc -O -S code.c
```

Produces file code.s

Note how the assembly:

- Maps to C code (**t = x+y** maps to the **addl** instruction)
- Contains only simple data types (long ints with the “**l**” suffix)
- Moves data between memory and registers

Disassembling Object Code

Disassembled

```
00401040 <_sum>:  
 0:      55          push    %ebp  
 1:  89 e5         mov     %esp,%ebp  
 3:  8b 45 0c      mov     0xc(%ebp),%eax  
 6:  03 45 08      add     0x8(%ebp),%eax  
 9:  89 ec         mov     %ebp,%esp  
 b:  5d           pop    %ebp  
 c:  c3           ret  
 d:  8d 76 00      lea     0x0(%esi),%esi
```

Disassembler

`objdump -d p`

- Useful tool for examining object code
- Analyzes bit pattern of series of instructions
- Produces approximate rendition of assembly code
- Can be run on either a .out (complete executable) or .o file

Alternate Disassembly

Object

0x401040:
0x55
0x89
0xe5
0xb8
0x45
0x0c
0x03

Disassembled

0x401040 <sum>:	push	%ebp
0x401041 <sum+1>:	mov	%esp, %ebp
0x401043 <sum+3>:	mov	0xc(%ebp), %eax
0x401046 <sum+6>:	add	0x8(%ebp), %eax
0x401049 <sum+9>:	mov	%ebp, %esp
0x40104b <sum+11>:	pop	%ebp
0x40104c <sum+12>:	ret	
0x40104d <sum+13>:	lea	0x0(%esi), %esi

0x45
0x08
0x89
0xec
0x5d
0xc3

Within gdb Debugger

gdb p

disassemble sum

■ Disassemble procedure

x/13b sum

■ Examine the 13 bytes starting at sum

Moving Data

Moving Data

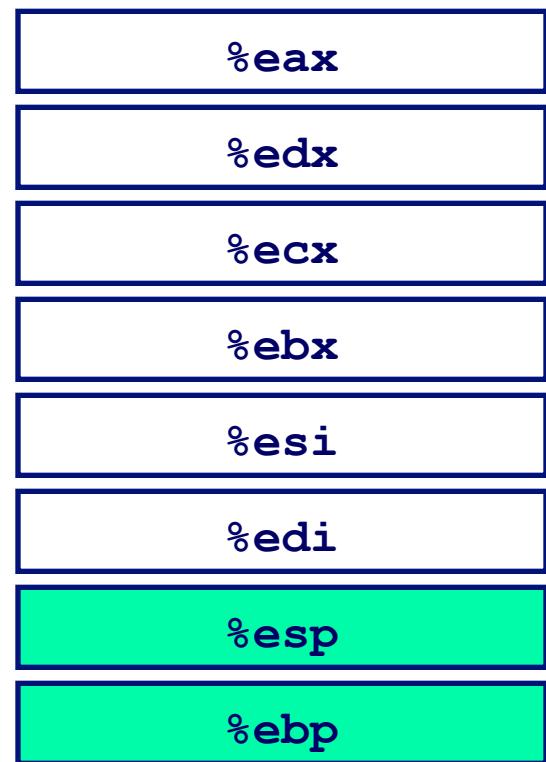
`movl Source, Dest`

- Move 4-byte (“long”) word
- Lots of these in typical code

Operand Types

- Register: One of 8 integer registers
 - But %esp and %ebp reserved for special use
 - Others have special uses for particular instructions
- Memory: 4 consecutive bytes of memory
 - Various “address modes”
- Immediate: Constant integer data
 - Like C constant, but prefixed with ‘\$’
 - E.g., \$0x400, \$-533
 - Encoded with 1, 2, or 4 bytes

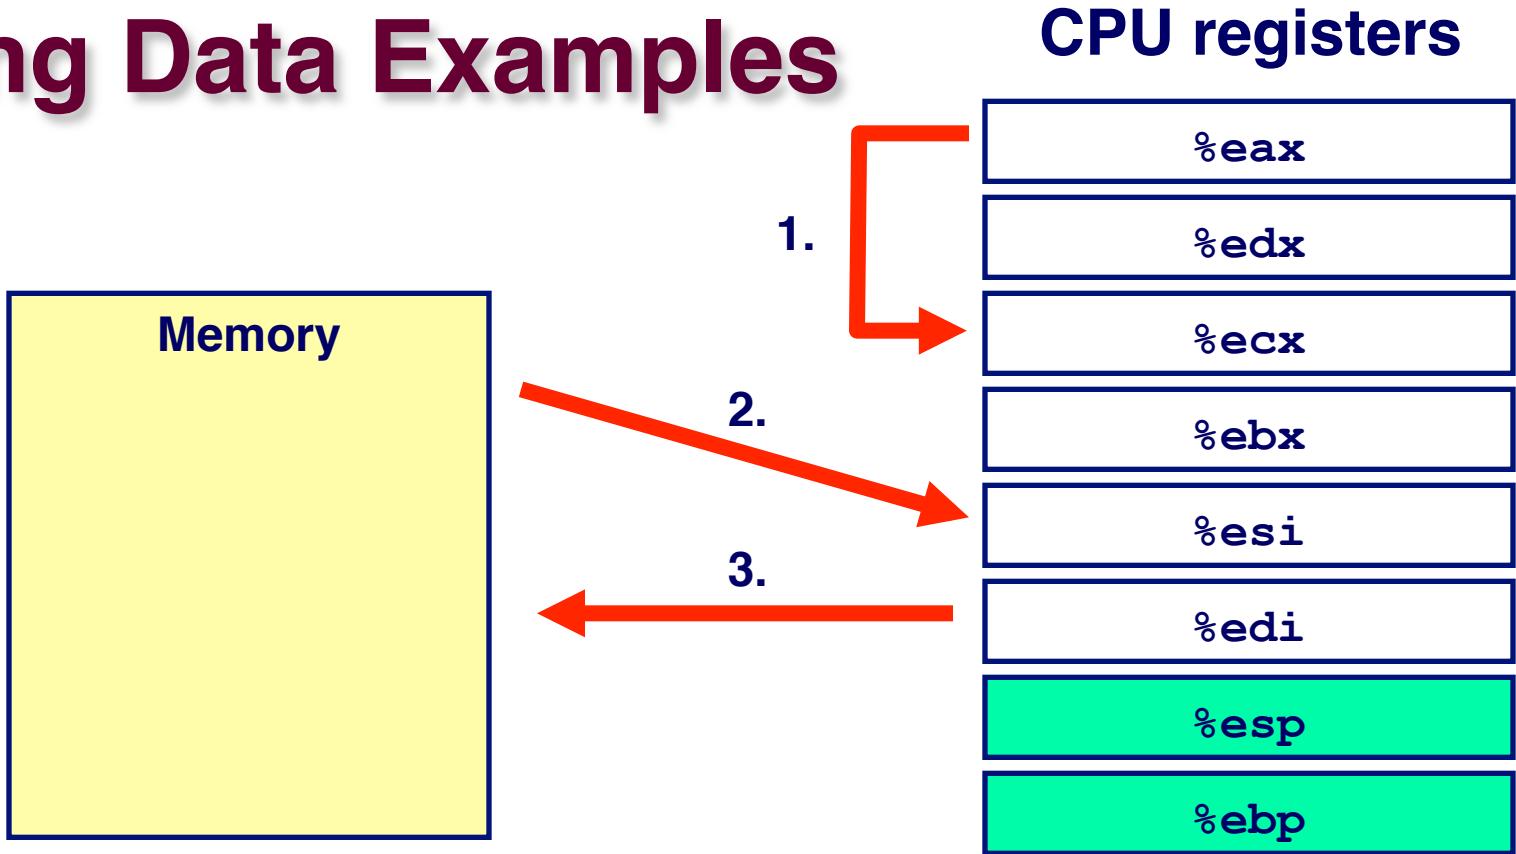
CPU registers



Representing Instructions

- **For historical reasons (16-bit processors), Intel terminology considers a “word” to be 16 bits long**
 - ‘`movw %ax, %dx`’, where the ‘`w`’ in `movw` implies a 16 bit quantity is about to be moved
 - ‘`movl %eax, %edx`’, the ‘`l`’ in `movl` implies a “long” 32-bit quantity is about to be moved.
 - **See text for more `mov` instructions:** `movb`, `movw`, `movl`, `movsbw`, `movsbl`, `movswl`, `movzbw`, `movzbl`, `movzwl`
- **Does the width of a C ‘long’ int == an x86 assembly language ‘long’ word?**
 - Yes, they’re both 32-bits or 4 bytes on a 32-bit x86 machine
 - No, a C ‘long’ is 64-bits on a 64-bit x86 machine, while an x86 assembly language ‘long’ is still 32-bits wide
 - i.e. a `movl` on a 64-bit machine is still going to move a 32-bit quantity

Moving Data Examples



Examples

1. Moving the value in one register to another
2. Moving a value at a memory location to a register
3. Moving a register value to a memory location

movl Operand Combinations

	Source	Destination	C Analog
movl	<i>Imm</i>	<i>Reg</i> movl \$0x4,%eax <i>Mem</i> movl \$-147,(%eax)	temp = 0x4; *p = -147;
	<i>Reg</i>	<i>Reg</i> movl %eax,%edx <i>Mem</i> movl %eax,(%edx)	temp2 = temp1; *p = temp;
	<i>Mem</i>	<i>Reg</i> movl (%eax),%edx	temp = *p;

- Cannot do memory-memory transfers with single instruction
 - i.e. can't do: `movl (%eax), (%edx)`

Simple Addressing Modes

Normal

(R)

Mem[Reg[R]]

- Register R specifies memory address

`movl (%ecx), %eax`

Displacement D(R)

Mem[Reg[R]+D]

- Register R specifies start of memory region
- Constant displacement D specifies offset

`movl 8(%ebp), %edx`



Go to memory address `%ebp+8`
and fetch the data located there

Using Simple Addressing Modes

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

```
pushl %ebp  
movl %esp,%ebp  
pushl %ebx
```

} Set Up

```
movl 12(%ebp),%ecx  
movl 8(%ebp),%edx  
movl (%ecx),%eax  
movl (%edx),%ebx  
movl %eax,(%edx)  
movl %ebx,(%ecx)
```

} Body

```
movl -4(%ebp),%ebx  
movl %ebp,%esp  
popl %ebp  
ret
```

} Clean Up

C Pointers – A Quick Recap

- **int count=1;**
 - Declare an integer named count
 - This allocates 4 bytes in memory for the variable count
 - Initialize count to the value 1
- **char *p1;**
 - Declare p1 as a *pointer* to a char, i.e. the value of p1 is interpreted as a memory address (4 bytes wide on 32-bit systems)
 - The pointer is allocated space in memory (4 bytes, not 1)
- **int *p2 = &count;**
 - Declare p2 as a pointer to an integer
 - Allocates 4 bytes in memory for the pointer (32-bit)
 - Initializes its value to the memory address of the count variable

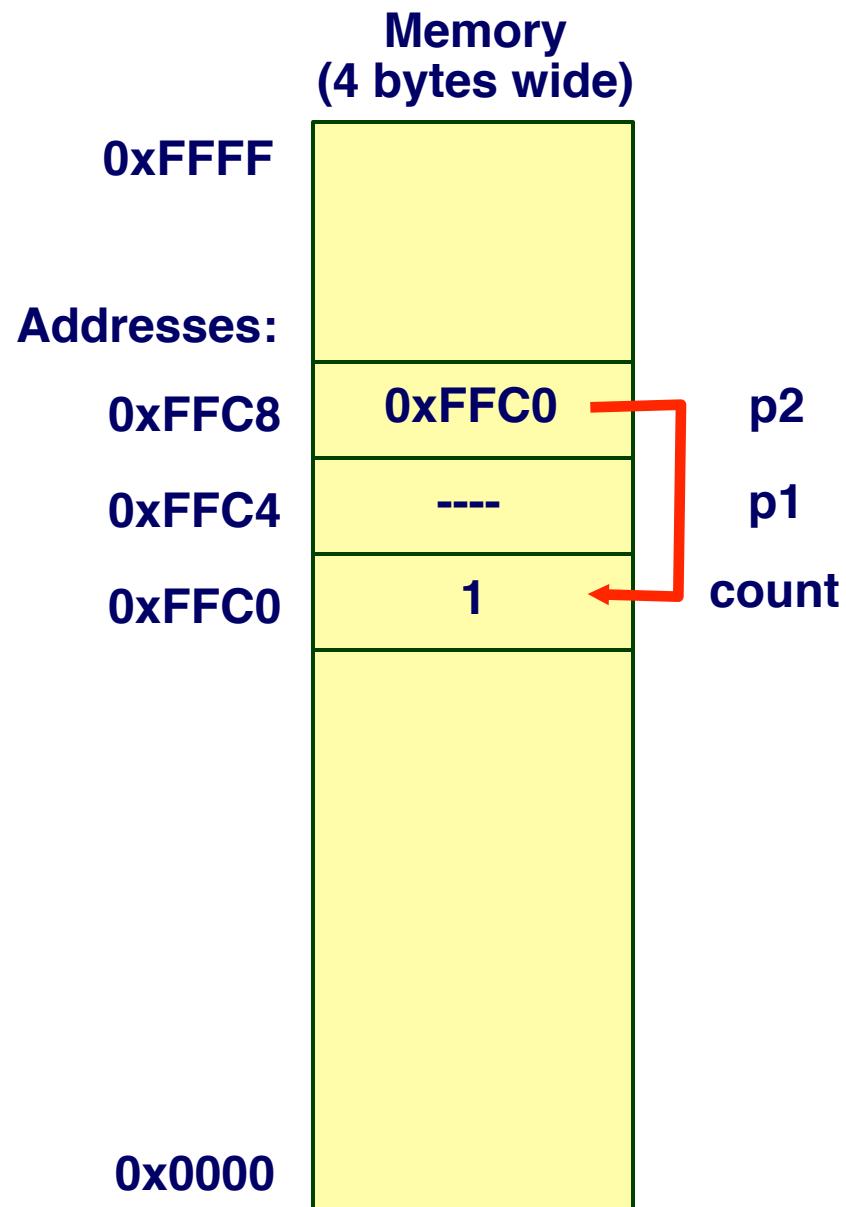
C Pointers (2)

```
int count=1;
```

```
char *p1;
```

```
int *p2 = &count;
```

- Assume the variables are laid out in memory as shown
- We see p2 storing the memory address of count, i.e. p2 is *pointing at* count
- p1 is uninitialized and not yet pointing at any character



For brevity, the two most significant bytes of address are not shown

Using Simple Addressing Modes

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

```
pushl %ebp  
movl %esp,%ebp  
pushl %ebx
```

} Set Up

```
movl 12(%ebp),%ecx  
movl 8(%ebp),%edx  
movl (%ecx),%eax  
movl (%edx),%ebx  
movl %eax,(%edx)  
movl %ebx,(%ecx)
```

} Body

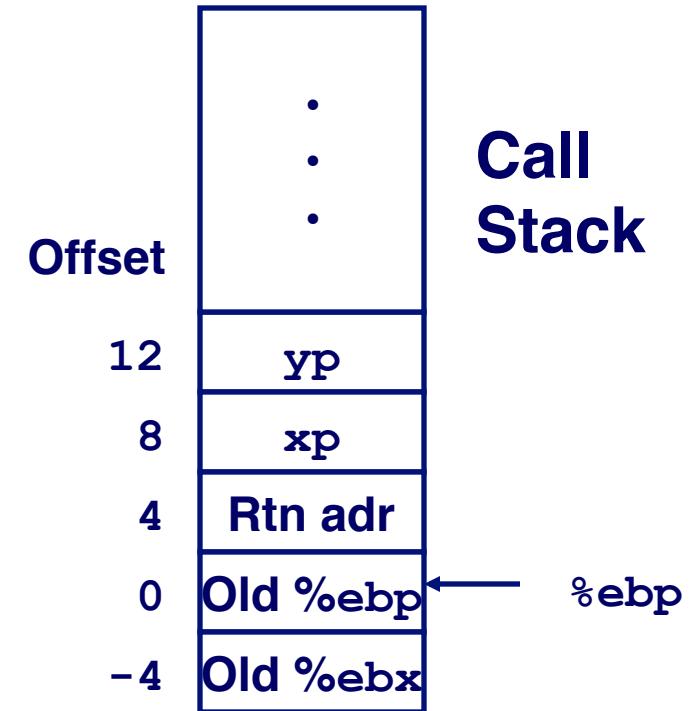
```
movl -4(%ebp),%ebx  
movl %ebp,%esp  
popl %ebp  
ret
```

} Finish

Understanding Swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

Memory
(shown as 4 bytes wide, but
in reality byte addressable)



Register	Variable
----------	----------

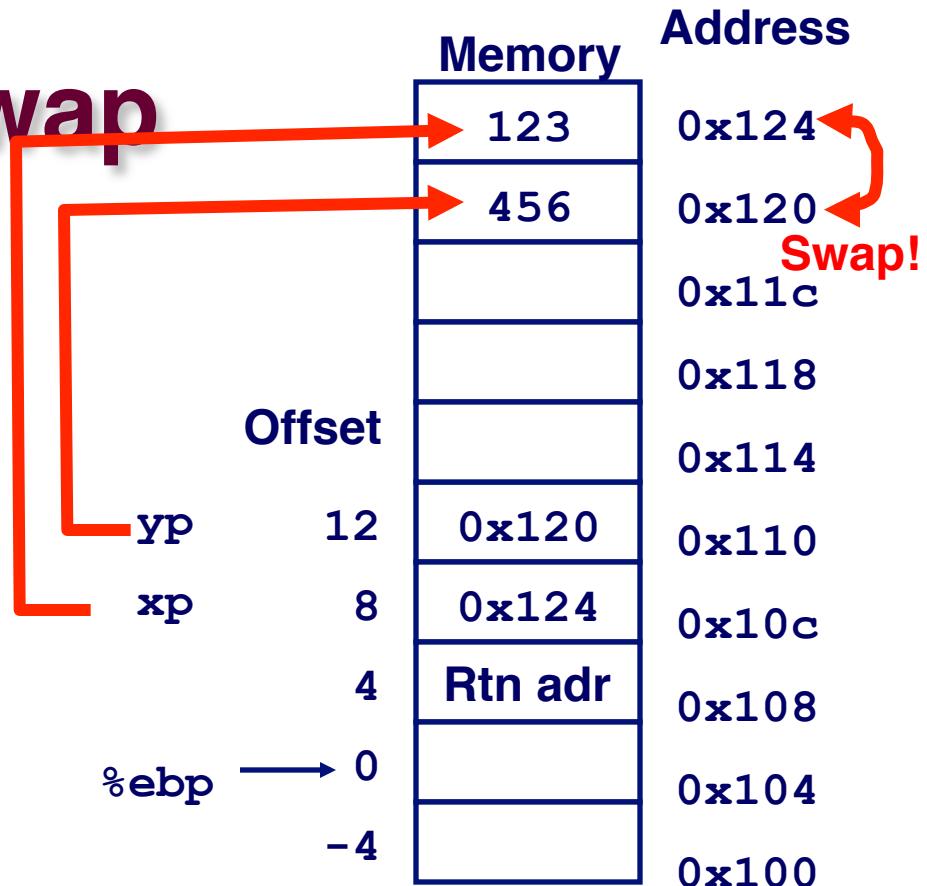
%ecx	yp
%edx	xp
%eax	t1
%ebx	t0

movl 12(%ebp), %ecx	# ecx = yp
movl 8(%ebp), %edx	# edx = xp
movl (%ecx), %eax	# eax = *yp (t1)
movl (%edx), %ebx	# ebx = *xp (t0)
movl %eax, (%edx)	# *xp = eax
movl %ebx, (%ecx)	# *yp = ebx

Understanding Swap

CPU Registers

%eax	
%edx	
%ecx	
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104



```

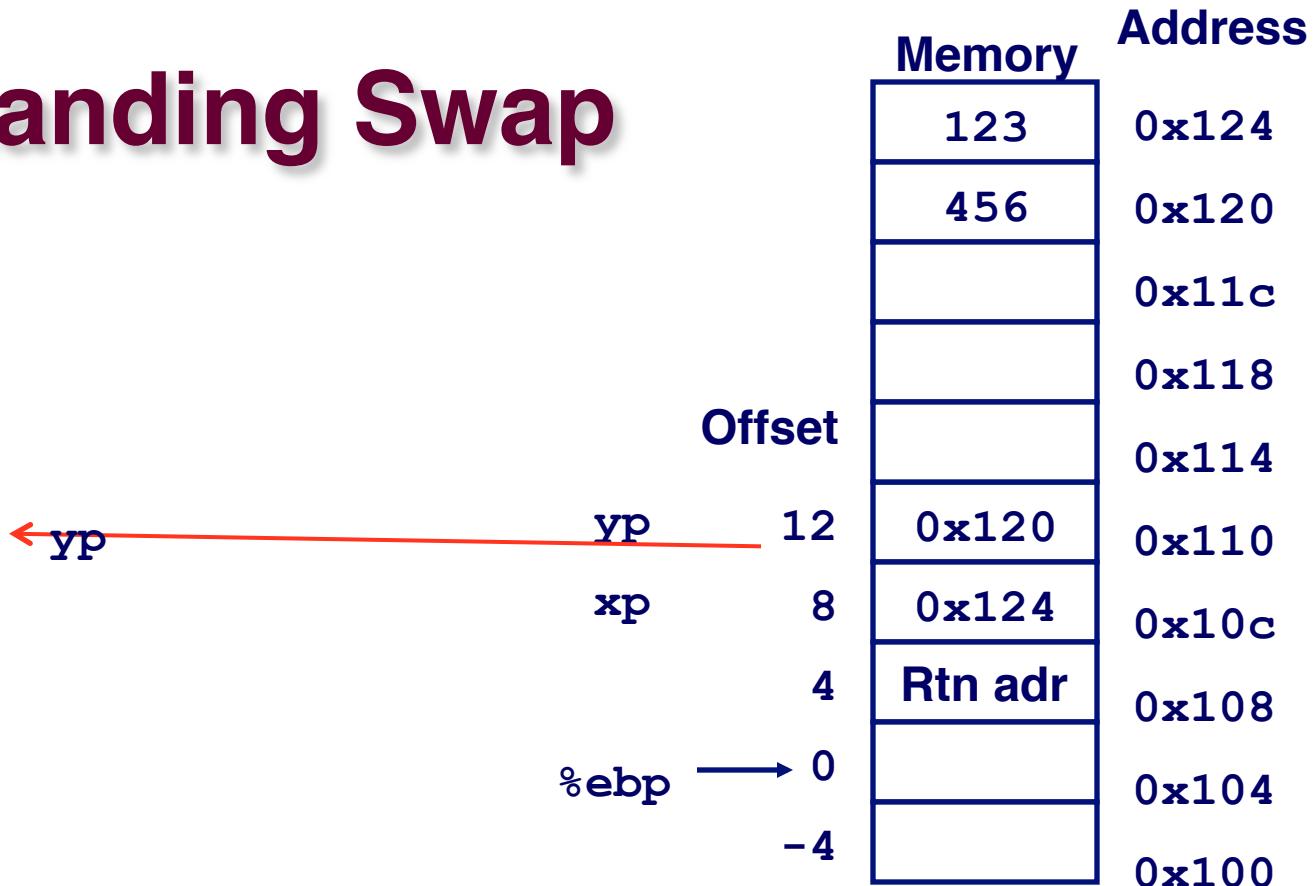
movl 12(%ebp), %ecx    # ecx = yp
movl 8(%ebp), %edx     # edx = xp
movl (%ecx), %eax      # eax = *yp (t1)
movl (%edx), %ebx      # ebx = *xp (t0)
movl %eax, (%edx)       # *xp = eax
movl %ebx, (%ecx)       # *yp = ebx

```

Understanding Swap

CPU Registers

%eax	
%edx	
%ecx	0x120
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104



```

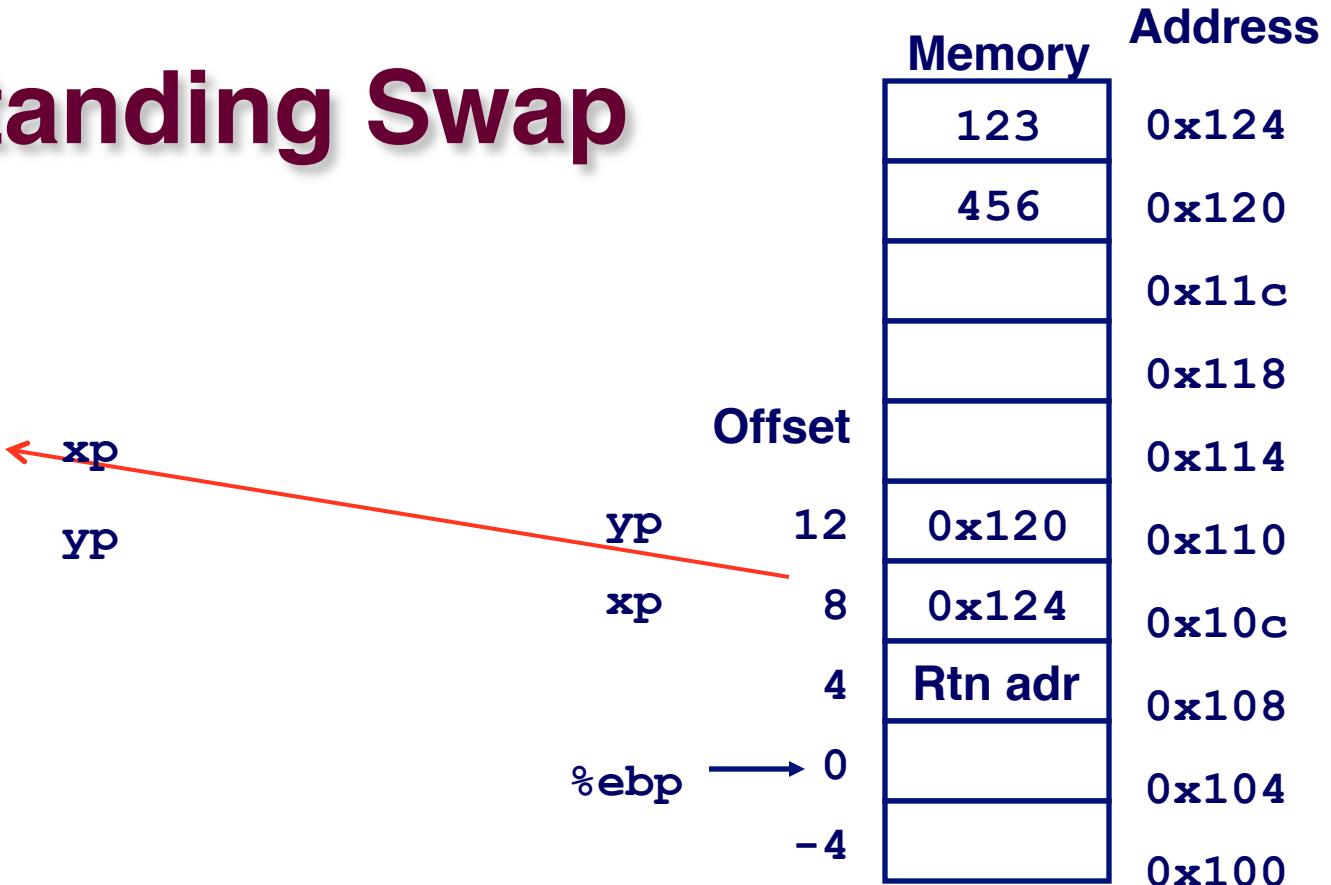
movl 12(%ebp),%ecx    # ecx = yp
movl 8(%ebp),%edx     # edx = xp
movl (%ecx),%eax      # eax = *yp (t1)
movl (%edx),%ebx      # ebx = *xp (t0)
movl %eax,(%edx)       # *xp = eax
movl %ebx,(%ecx)       # *yp = ebx

```

Understanding Swap

CPU Registers

%eax	
%edx	0x124
%ecx	0x120
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104



```

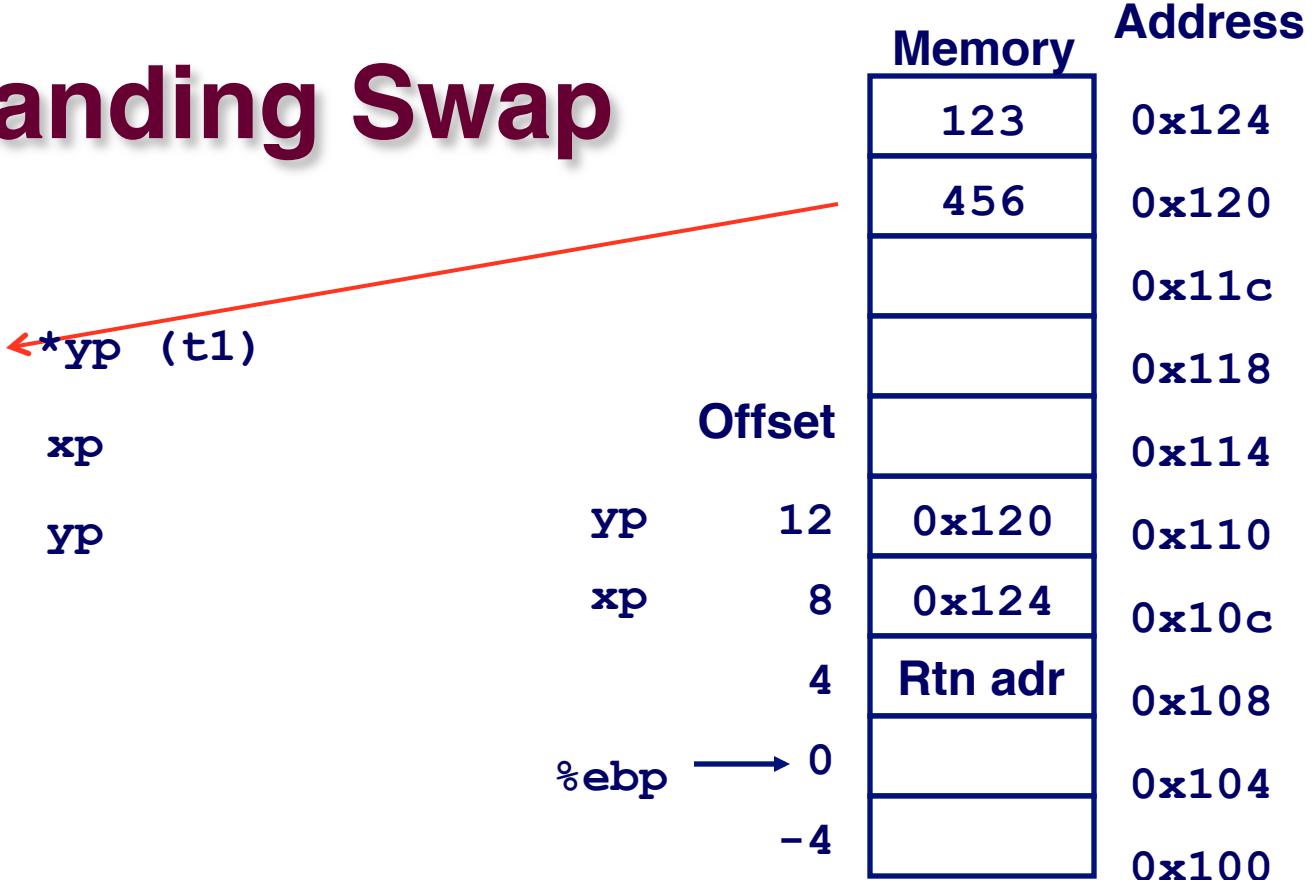
movl 12(%ebp),%ecx    # ecx = yp
movl 8(%ebp),%edx     # edx = xp
movl (%ecx),%eax      # eax = *yp (t1)
movl (%edx),%ebx      # ebx = *xp (t0)
movl %eax,(%edx)       # *xp = eax
movl %ebx,(%ecx)       # *yp = ebx

```

Understanding Swap

CPU Registers

%eax	456
%edx	0x124
%ecx	0x120
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104



```

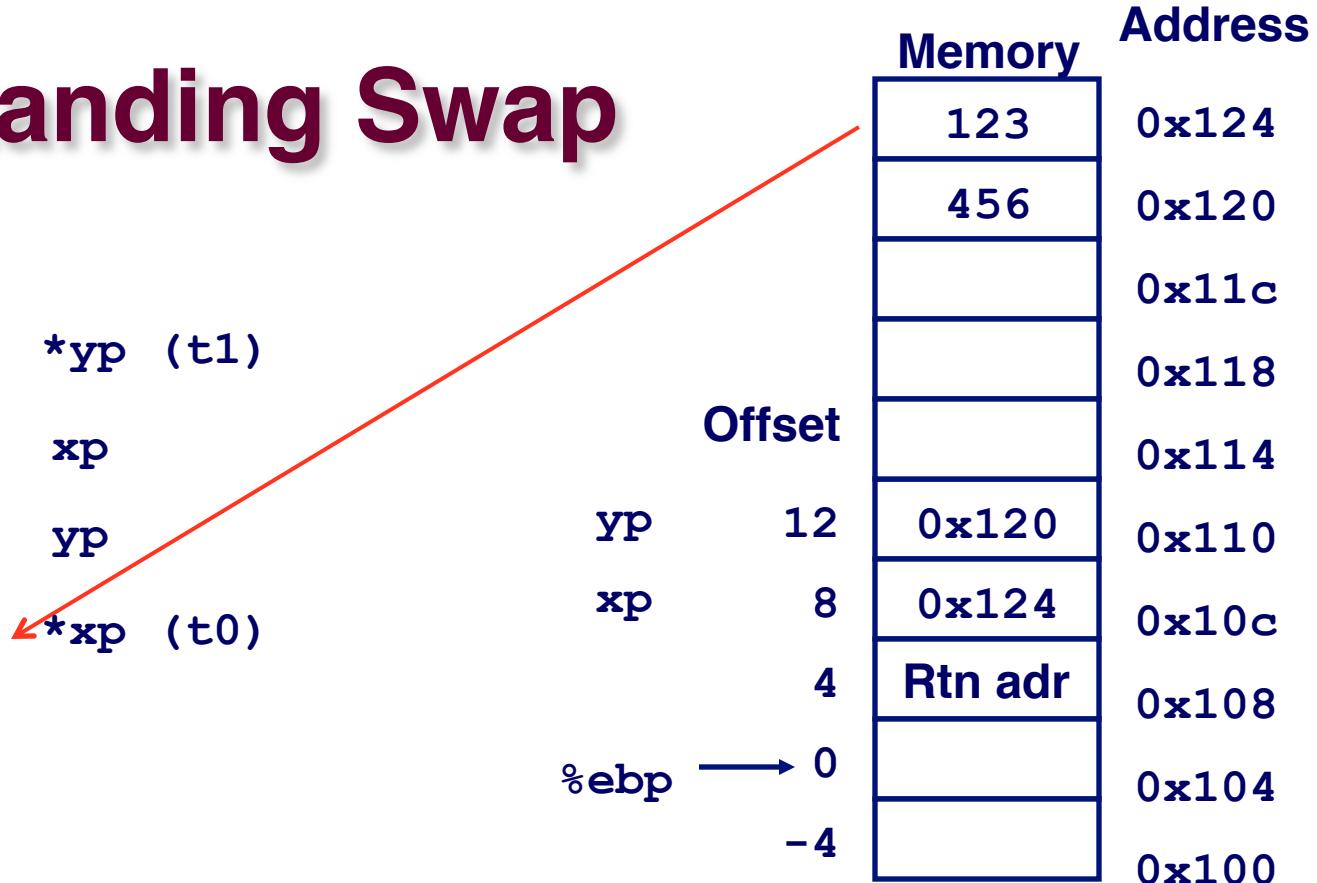
movl 12(%ebp), %ecx    # ecx = yp
movl 8(%ebp), %edx     # edx = xp
movl (%ecx), %eax      # eax = *yp (t1)
movl (%edx), %ebx      # ebx = *xp (t0)
movl %eax, (%edx)       # *xp = eax
movl %ebx, (%ecx)       # *yp = ebx

```

Understanding Swap

CPU Registers

%eax	456
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104



```

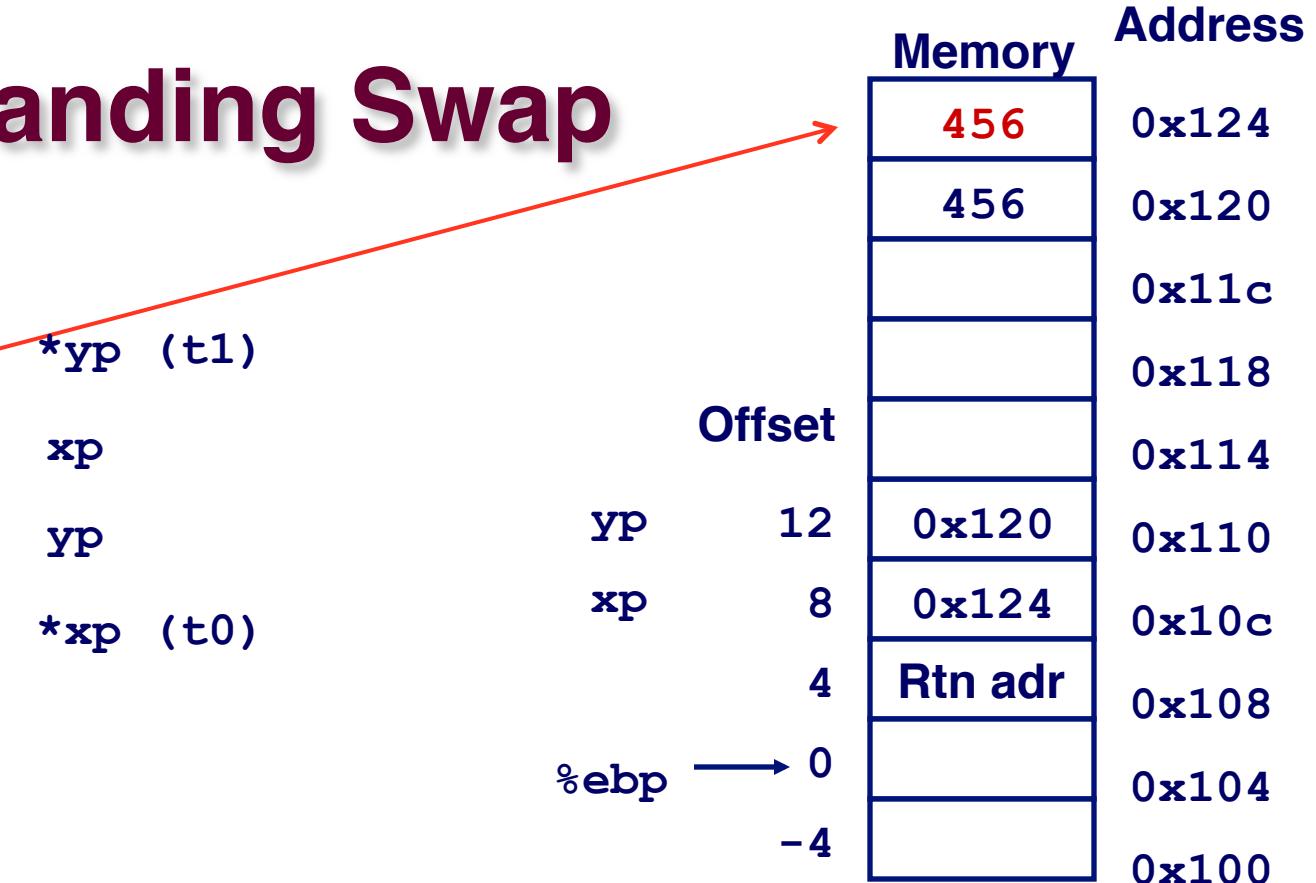
movl 12(%ebp), %ecx    # ecx = yp
movl 8(%ebp), %edx     # edx = xp
movl (%ecx), %eax      # eax = *yp (t1)
movl (%edx), %ebx      # ebx = *xp (t0)
movl %eax, (%edx)       # *xp = eax
movl %ebx, (%ecx)       # *yp = ebx

```

Understanding Swap

CPU Registers

%eax	456
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104



```

movl 12(%ebp), %ecx    # ecx = yp
movl 8(%ebp), %edx     # edx = xp
movl (%ecx), %eax      # eax = *yp (t1)
movl (%edx), %ebx      # ebx = *xp (t0)
movl %eax, (%edx)       # *xp = eax
movl %ebx, (%ecx)       # *yp = ebx

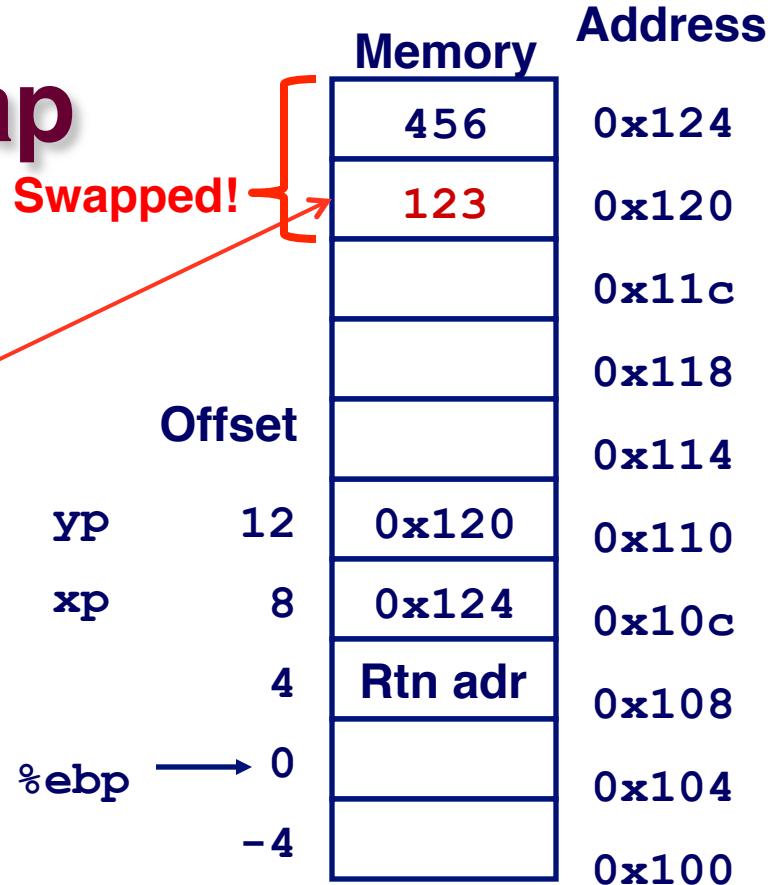
```

Understanding Swap

CPU Registers

%eax	456
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104

*yp (t1)
xp
yp
*xp (t0)



```

movl 12(%ebp), %ecx    # ecx = yp
movl 8(%ebp), %edx     # edx = xp
movl (%ecx), %eax      # eax = *yp (t1)
movl (%edx), %ebx      # ebx = *xp (t0)
movl %eax, (%edx)       # *xp = eax
movl %ebx, (%ecx)       # *yp = ebx

```

C operators – Assembly Equivalents?

Operators

() [] -> .
! ~ ++ -- + - * & (type) sizeof
* / %
+ -
<< >>
< <= > >=
== !=
&
^
|
&&
||
?:
= += -= *= /= %= &= ^= != <<= >>=
,

Many of these C operators have direct x86 assembly equivalents

Some Assembly Arithmetic Operations

Format Computation

Two Operand Instructions

<code>addl Src,Dest</code>	$Dest = Dest + Src$
<code>subl Src,Dest</code>	$Dest = Dest - Src$
<code>imull Src,Dest</code>	$Dest = Dest * Src$ Signed multiply
<code>sall Src,Dest</code>	$Dest = Dest \ll Src$ Also called shll
<code>sarl Src,Dest</code>	$Dest = Dest \gg Src$ Arithmetic
<code>shrl Src,Dest</code>	$Dest = Dest \gg Src$ Logical
<code>xorl Src,Dest</code>	$Dest = Dest \wedge Src$
<code>andl Src,Dest</code>	$Dest = Dest \& Src$
<code>orl Src,Dest</code>	$Dest = Dest \mid Src$

Some Arithmetic Operations

Format Computation

One Operand Instructions

incl Dest $Dest = Dest + 1$

decl Dest $Dest = Dest - 1$

negl Dest $Dest = - Dest$

notl Dest $Dest = \sim Dest$

Indexed Addressing Modes

Most General Form

$D(Rb, Ri, S)$

$\text{Mem}[\text{Reg}[Rb] + S * \text{Reg}[Ri] + D]$

- D : Constant “displacement” 1, 2, or 4 bytes
- Rb : Base register: Any of 8 integer registers
- Ri : Index register: Any, except for $\%esp$
 - Unlikely you’d use $\%ebp$, either
- S : Scale: 1, 2, 4, or 8

Example

- `movl 8(%eax,%edx,4), dst` is equivalent to
`movl (%eax+4*%edx+8), dst`

Special Cases (can leave out some elements)

(Rb, Ri)

$\text{Mem}[\text{Reg}[Rb] + \text{Reg}[Ri]]$

$D(Rb, Ri)$

$\text{Mem}[\text{Reg}[Rb] + \text{Reg}[Ri] + D]$

(Rb, Ri, S)

$\text{Mem}[\text{Reg}[Rb] + S * \text{Reg}[Ri]]$

Address Computation Examples

%edx	0xf000
%ecx	0x100

Expression	Computation	Address
0x8(%edx)	0xf000 + 0x8	0xf008
(%edx,%ecx)	0xf000 + 0x100	0xf100
(%edx,%ecx,4)	0xf000 + 4*0x100	0xf400
0x80(,%edx,2)	2*0xf000 + 0x80	0x1e080

`leal` Instruction for Address Computation

`leal Src, Dest`

- *Src* is indexed address mode expression
- Set *Dest* (must be register) to address denoted by expression

You should think of this as “Compute Using Effective Address”

- Example:

`leal 10(%edx, %edx, 4), %eax`

$$\underbrace{\quad}_{\%edx} \underbrace{+ 4 * \underbrace{\%edx}_{\%edx}}_{\%edx} + 10$$

$$= 5 * \%edx + 10$$

Therefore “ $\%eax = 5 * \%edx + 10$ ”

- Compare to:

`movl 10(%edx, %edx, 4), %eax`

means “ $\%eax = \text{Mem}[5 * \%edx + 10]$ ”

leal Instruction for Address Computation

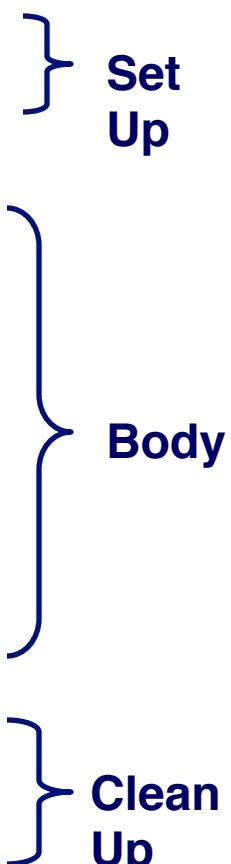
Uses

- Computing address without doing memory reference
 - E.g., translation of $p = \&x[i]$;
- Computing arithmetic expressions of the form $x + k^*y$
 - $k = 1, 2, 4$, or 8 .

An Arithmetic Example in Assembly

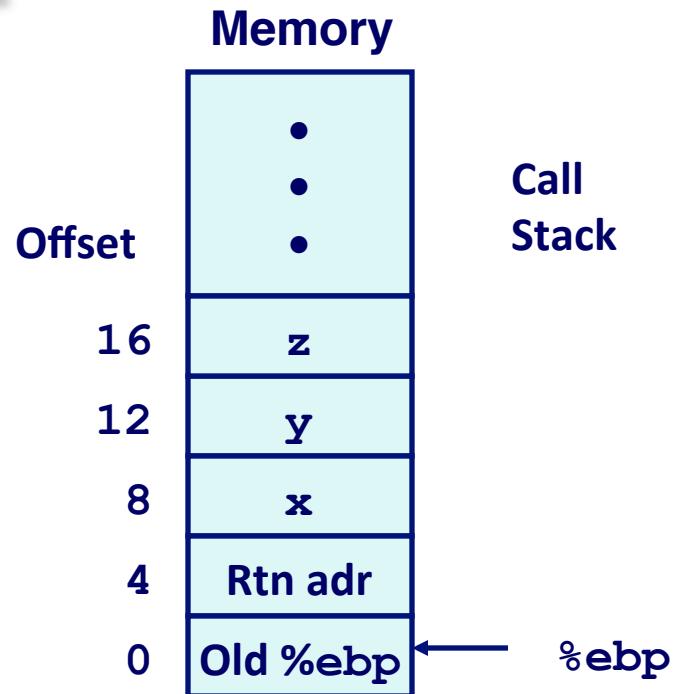
```
int arith  
  (int x, int y, int z)  
{  
    int t1 = x+y;  
    int t2 = z+t1;  
    int t3 = x+4;  
    int t4 = y * 48;  
    int t5 = t3 + t4;  
    int rval = t2 * t5;  
    return rval;  
}
```

```
arith:  
  pushl %ebp  
  movl %esp,%ebp  
  
  movl 8(%ebp),%eax  
  movl 12(%ebp),%edx  
  leal (%edx,%eax),%ecx  
  leal (%edx,%edx,2),%edx  
  sall $4,%edx  
  addl 16(%ebp),%ecx  
  leal 4(%edx,%eax),%eax  
  imull %ecx,%eax  
  
  movl %ebp,%esp  
  popl %ebp  
  ret
```



Understanding arith

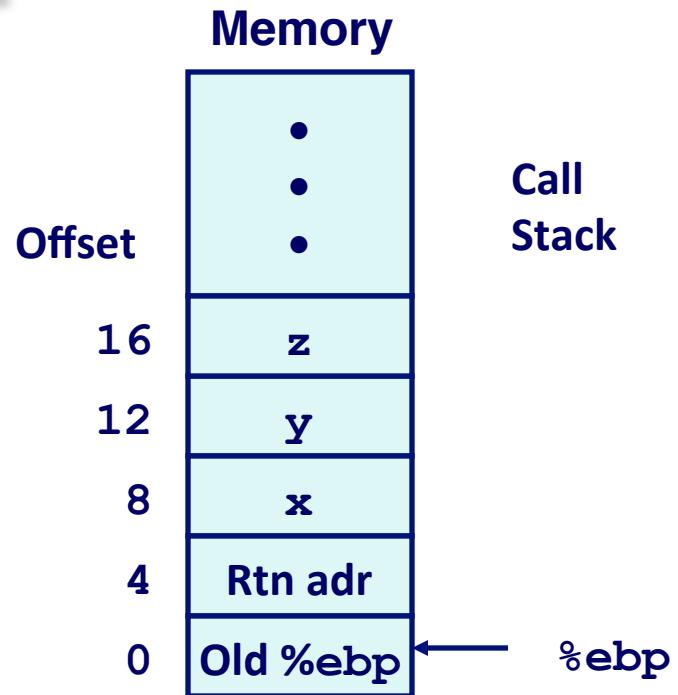
```
int arith  
    (int x, int y, int z)  
{  
    int t1 = x+y;  
    int t2 = z+t1;  
    int t3 = x+4;  
    int t4 = y * 48;  
    int t5 = t3 + t4;  
    int rval = t2 * t5;  
    return rval;  
}
```



```
movl 8(%ebp),%eax  
movl 12(%ebp),%edx  
leal (%edx,%eax),%ecx  
leal (%edx,%edx,2),%edx  
sall $4,%edx  
addl 16(%ebp),%ecx  
leal 4(%edx,%eax),%eax  
-36imull %ecx,%eax
```

Understanding arith

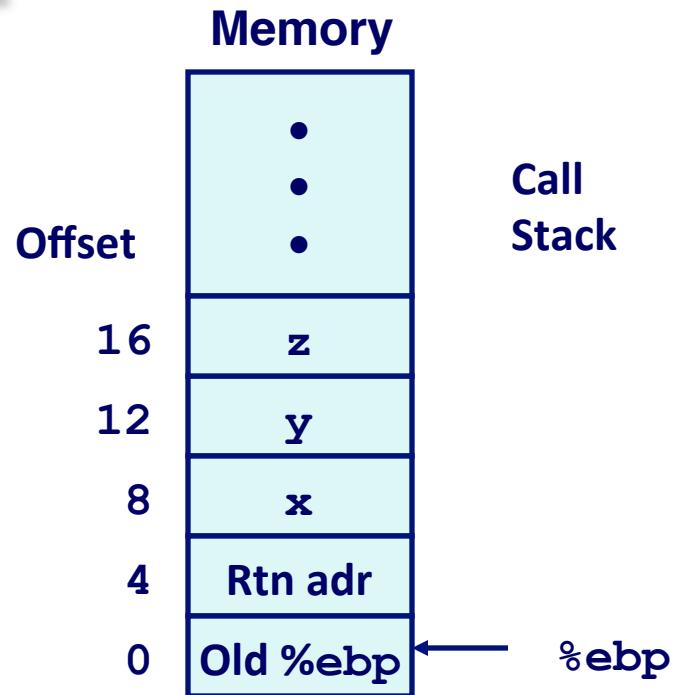
```
int arith  
    (int x, int y, int z)  
{  
    int t1 = x+y;  
    int t2 = z+t1;  
    int t3 = x+4;  
    int t4 = y * 48;  
    int t5 = t3 + t4;  
    int rval = t2 * t5;  
    return rval;  
}
```



movl 8(%ebp),%eax	# eax = x
movl 12(%ebp),%edx	# edx = y
leal (%edx,%eax),%ecx	# ecx = x+y (t1)
leal (%edx,%edx,2),%edx	# edx = 3*y
sall \$4,%edx	# edx = 48*y (t4)
addl 16(%ebp),%ecx	# ecx = z+t1 (t2)
leal 4(%edx,%eax),%eax	# eax = 4+t4+x (t5)
-37imull %ecx,%eax	# eax = t5*t2 (rval)

Understanding arith

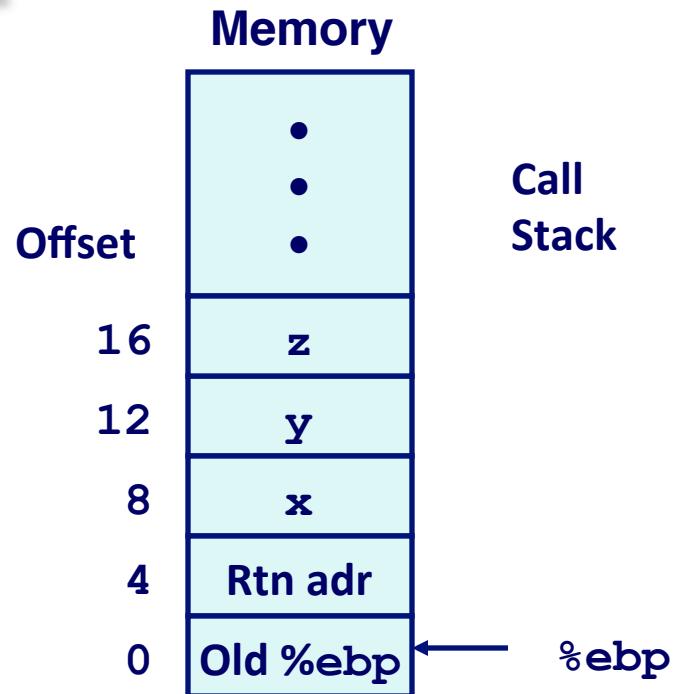
```
int arith
    (int x, int y, int z)
{
    int t1 = x+y;
    int t2 = z+t1;
    int t3 = x+4;
    int t4 = y * 48;
    int t5 = t3 + t4;
    int rval = t2 * t5;
    return rval;
}
```



```
movl 8(%ebp),%eax          # eax = x
movl 12(%ebp),%edx          # edx = y
leal (%edx,%eax),%ecx      # ecx = x+y (t1)
leal (%edx,%edx,2),%edx      # edx = 3*y
sall $4,%edx                # edx = 48*y (t4)
addl 16(%ebp),%ecx          # ecx = z+t1 (t2)
leal 4(%edx,%eax),%eax      # eax = 4+t4+x (t5)
-38imull %ecx,%eax          # eax = t5*t2 (rval)
```

Understanding arith

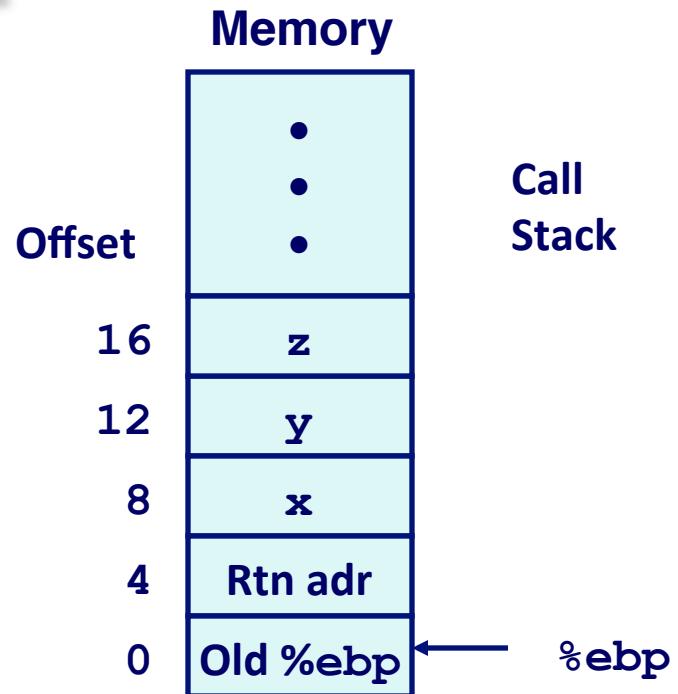
```
int arith
    (int x, int y, int z)
{
    int t1 = x+y;
    int t2 = z+t1;
    int t3 = x+4;
    int t4 = y * 48;
    int t5 = t3 + t4;
    int rval = t2 * t5;
    return rval;
}
```



movl 8(%ebp),%eax	# eax = x
movl 12(%ebp),%edx	# edx = y
leal (%edx,%eax),%ecx	# ecx = x+y (t1)
leal (%edx,%edx,2),%edx	# edx = 3*y
sall \$4,%edx	# edx = 48*y (t4)
addl 16(%ebp),%ecx	# ecx = z+t1 (t2)
leal 4(%edx,%eax),%eax	# eax = 4+t4+x (t5)
-39imull %ecx,%eax	# eax = t5*t2 (rval)

Understanding arith

```
int arith
    (int x, int y, int z)
{
    int t1 = x+y;
    int t2 = z+t1;
    int t3 = x+4;
    int t4 = y * 48;
    int t5 = t3 + t4;
    int rval = t2 * t5;
    return rval;
}
```



movl 8(%ebp),%eax	# eax = x
movl 12(%ebp),%edx	# edx = y
leal (%edx,%eax),%ecx	# ecx = x+y (t1)
leal (%edx,%edx,2),%edx	# edx = 3*y
sall \$4,%edx	# edx = 48*y (t4)
addl 16(%ebp),%ecx	# ecx = z+t1 (t2)
leal 4(%edx,%eax),%eax	# eax = 4+t4+x (t5)
-40imull %ecx,%eax	# eax = t5*t2 (rval)

Note that compiler can reorder the original code (t2 now computed after t4)

Supplementary Slides

Assembly Characteristics

Minimal Data Types

- “Integer” data of 1, 2, or 4 bytes
 - Data values
 - Addresses (untyped pointers)
- Floating point data of 4, 8, or 10 bytes
- No aggregate types such as arrays or structures
 - Just contiguously allocated bytes in memory

Primitive Operations

- Perform arithmetic function on register or memory data
- Transfer data between memory and register
 - Load data from memory into register
 - Store register data into memory
- Transfer control
 - Unconditional jumps to/from procedures
 - Conditional branches

Machine Instruction Example

```
int t = x+y;
```

```
addl 8(%ebp),%eax
```

Similar to
expression
 $x += y$

```
0x401046: 03 45 08
```

C Code

- Add two signed integers

Assembly

- Add 2 4-byte integers
 - “Long” words in GCC parlance
 - Same instruction whether signed or unsigned
- Operands:

x:	Register	%eax
y:	Memory	M[%ebp+8]
t:	Register	%eax

» Return function value in %eax

Object Code

- 3-byte instruction
- Stored at address 0x401046

Object Code

Code for sum

```
0x401040 <sum>:  
 0x55      • Total of 13  
 0x89      bytes  
 0xe5      • Each  
 0x8b      instruction 1,  
 0x45      2, or 3 bytes  
 0x0c      • Starts at  
 0x03      address  
 0x45      0x401040  
 0x08  
 0x89  
 0xec  
 0x5d  
 0xc3
```

Assembler

- Translates .s into .o
- Binary encoding of each instruction
- Nearly-complete image of executable code
- Missing linkages between code in different files

Linker

- Resolves references between files
- Combines with static run-time libraries
 - E.g., code for malloc, printf
- Some libraries are *dynamically linked*
 - Linking occurs when program begins execution

What Can be Disassembled?

```
% objdump -d WINWORD.EXE

WINWORD.EXE:      file format pei-i386

No symbols in "WINWORD.EXE".
Disassembly of section .text:

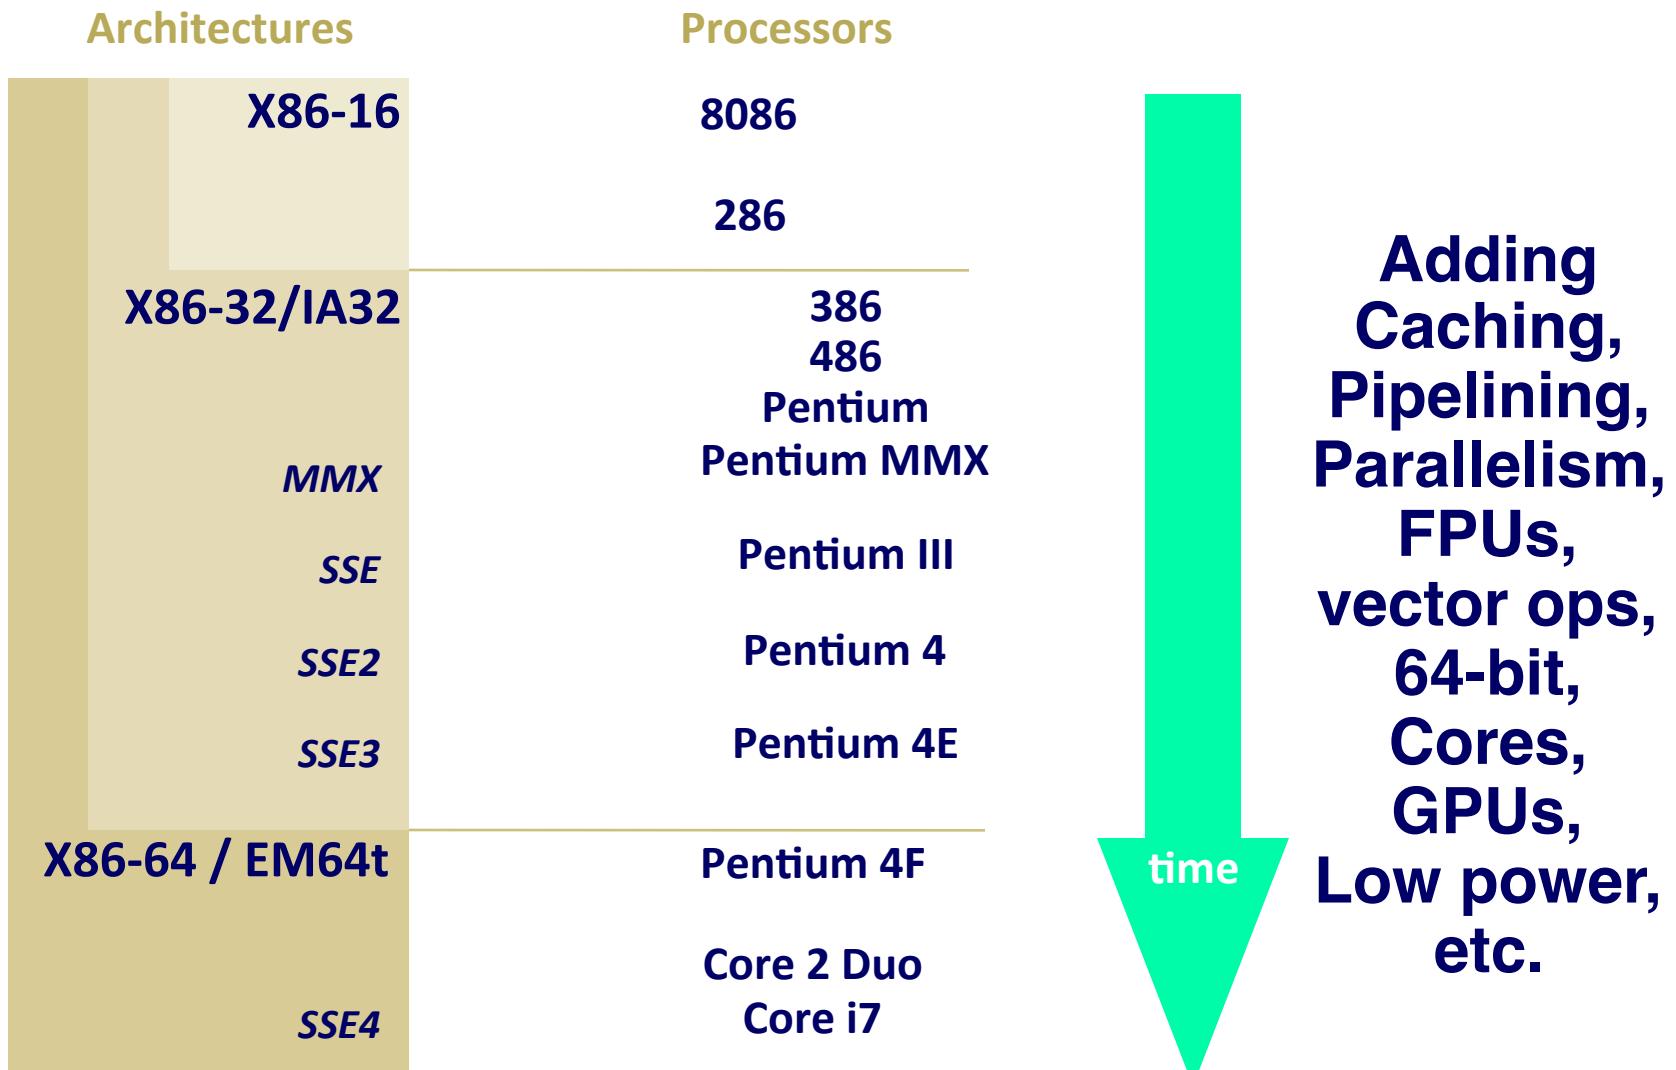
30001000 <.text>:
30001000: 55          push    %ebp
30001001: 8b ec        mov     %esp,%ebp
30001003: 6a ff        push    $0xffffffff
30001005: 68 90 10 00 30  push    $0x30001090
3000100a: 68 91 dc 4c 30  push    $0x304cdc91
```

- Anything that can be interpreted as executable code
- Disassembler examines bytes and reconstructs assembly source

Intel x86 Evolution: Milestones

<i>Name</i>	<i>Date</i>	<i>Transistors</i>	<i>MHz</i>
· 8086	1978	29K	5-10
		■ Intel's first 16-bit processor. Basis for IBM PC & DOS	
		■ 1MB address space	
· 386	1985	275K	16-33
		■ Intel's first 32 bit processor , referred to as IA32	
		■ Added “flat addressing” and support for MMU with paging	
		■ Capable of running Unix	
		■ 32-bit Linux/gcc uses no instructions introduced in later models	
· Pentium 4F	2005	230M	2800-3800
		■ A 64-bit processor. Intel's Itanium 64-bit CPU line debuting in 2001 was less successful due to performance & backward compatibility issues	
		■ Meanwhile, Pentium 4s (Netburst arch.) phased out in favor of “Core” line	

Intel x86 Processors: Overview

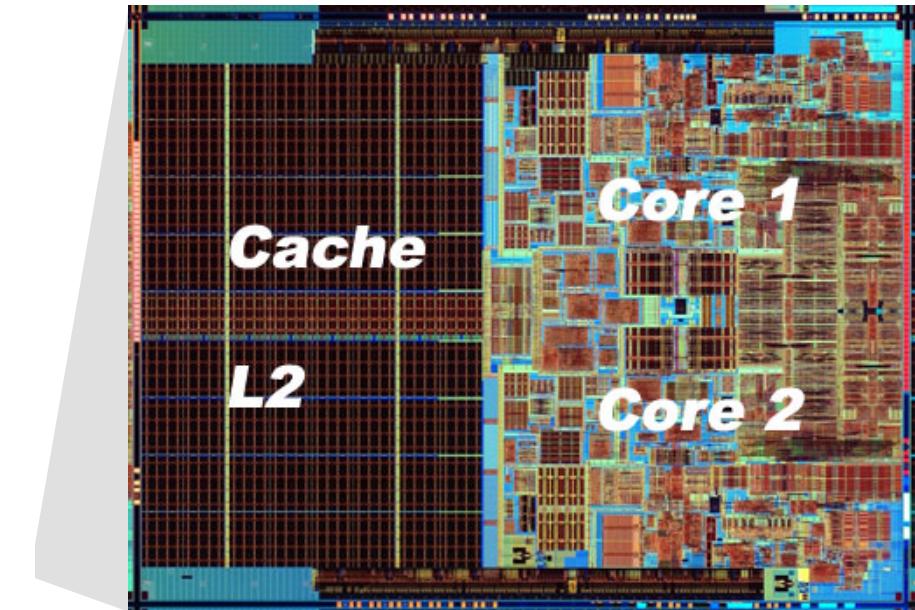


IA: often redefined as latest Intel architecture

Intel x86 Processors, contd.

▪ Machine Evolution

■ 486	1989	1.9M
■ Pentium	1993	3.1M
■ Pentium/MMX	1997	4.5M
■ PentiumPro	1995	6.5M
■ Pentium III	1999	8.2M
■ Pentium 4	2001	42M
■ Core 2 Duo	2006	291M



▪ Added Features

- Instructions to support multimedia operations
 - Parallel operations on 1, 2, and 4-byte data, both integer & FP
- Instructions to enable more efficient conditional operations

▪ Linux/GCC Evolution

- Very limited

Pentium Pro (P6)

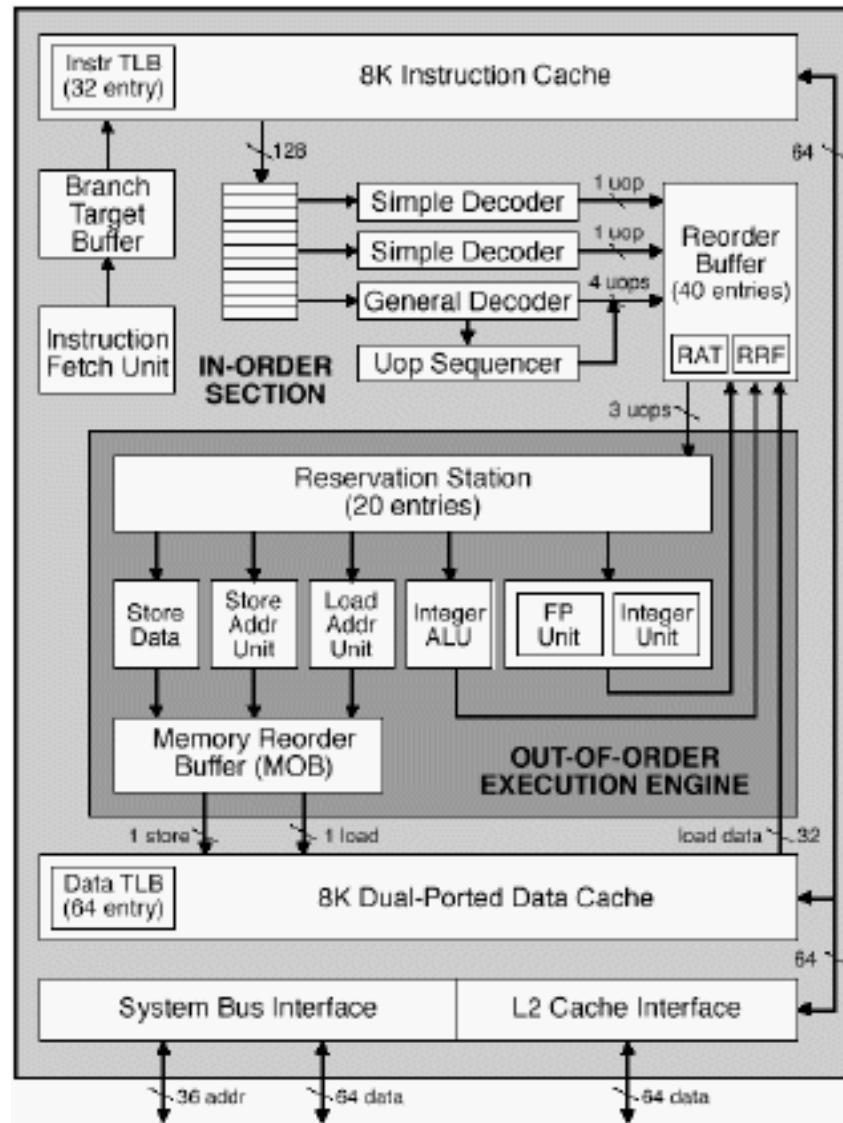
History

- Announced in Feb. '95
- Basis for Pentium II, Pentium III, and Celeron processors
- Pentium 4 similar idea, but different details

Features

- Dynamically translates instructions to more regular format
 - Very wide, but simple instructions
- Executes operations in parallel
 - Up to 5 at once
- Very deep pipeline
 - 12–18 cycle latency

Pentium Pro Block Diagram



Microprocessor Report
2/16/95

PentiumPro Operation

Translates instructions dynamically into “Uops”

- 118 bits wide
- Holds operation, two sources, and destination

Executes Uops with “Out of Order” engine

- Uop executed when
 - Operands available
 - Functional unit available
- Execution controlled by “Reservation Stations”
 - Keeps track of data dependencies between uops
 - Allocates resources

Consequences

- Indirect relationship between IA32 code & what actually gets executed
- Tricky to predict / optimize performance at assembly level

New Species: IA64, then IPF, then Itanium,...

<i>Name</i>	<i>Date</i>	<i>Transistors</i>
· Itanium	2001	10M
		<ul style="list-style-type: none">■ First shot at 64-bit architecture: first called IA64<ul style="list-style-type: none">■ DEC's Alpha CPU was 64-bit in 1992, Sun SPARC had 64-bit in 1995■ Radically new VLIW instruction set designed for high performance■ Can run existing IA32 programs in compatibility mode<ul style="list-style-type: none">● On-board “x86 engine” – not great performance■ Joint project with Hewlett-Packard
· Itanium 2	2002	221M
		<ul style="list-style-type: none">■ Big performance boost
· Itanium 2 Dual-Core	2006	1.7B
· Itanium has not taken off in marketplace		<ul style="list-style-type: none">■ Lack of backward compatibility, no good compiler support,– Pentium 4 got too good

x86 Clones: Advanced Micro Devices (AMD)

- **Historically**

- AMD has followed just behind Intel
 - A little bit slower, a lot cheaper

- **Then**

- Recruited top circuit designers from Digital Equipment Corp. and other downward trending companies
 - Built Opteron: tough competitor to Pentium 4
 - In 2003, developed x86-64, their own extension to 64 bits

- **Recently**

- Intel much quicker with dual core design
 - Intel currently far ahead in performance
 - em64t backwards compatible to x86-64

Intel's 64-Bit

- Intel Attempted Radical Shift from IA32 to IA64
 - Totally different architecture (Itanium)
 - Executes IA32 code only as legacy
 - Performance disappointing
- AMD Stepped in with Evolutionary Solution
 - x86-64 (now called “AMD64”)
- Intel Felt Obligated to Focus on IA64
 - Hard to admit mistake or that AMD is better
- 2004: Intel Announces EM64T extension to IA32
 - Extended Memory 64-bit Technology
 - Almost identical to x86-64!
- Meanwhile: EM64t well introduced, wide adoption now
- Our focus will be on 32-bit assembly language programming
 - But we will point out 64-bit equivalents as we go

Other Processors

ARM-based RISC processors

- Large market in mobile smartphones, e.g. Android phones and tablets as well as iPhone/iPad
- low power, low cost, reasonably high performance for mobile space

CISC vs. RISC – not evident there's a clear winner

- Market forces affect outcome, not just technical issues
- For more, see RISC vs. CISC and RISC.

CISC Properties

Instruction can reference different operand types

- Immediate, register, memory

Arithmetic operations can read/write memory

Memory reference can involve complex computation

- $R_b + S^*R_i + D$
- Can be used for arithmetic!

Instructions can have varying lengths

- IA32 instructions 1-15 bytes

RISC Properties

Only load/store instructions can access and move to/from memory

Data instructions only manipulate registers, not memory

Each instruction takes one clock cycle

Complex computations are performed by many instructions

Deep pipelining

CISC vs RISC

“A common misunderstanding of the phrase "reduced instruction set computer" is the mistaken idea that instructions are simply eliminated, resulting in a smaller set of instructions. In fact, over the years, RISC instruction sets have grown in size, and today many of them have a larger set of instructions than many CISC CPUs.”

and

“The term "reduced" in that phrase was intended to describe the fact that the amount of work any single instruction accomplishes is reduced—at most a single data memory cycle—compared to the "complex instructions" of CISC CPUs that may require dozens of data memory cycles in order to execute a single instruction.”

-- Wikipedia

Other Processors (2)

Intel's Atom targeting small/embedded low power space too, e.g. netbooks

- IA32-compatible

Graphics Processing Units (GPUs)

- Designed for highly parallel data processing, such as computing the 3D rendering of a graphics scene
- General-purpose CPUs are sequential with some parallelism

Data Representations: IA32 + x86-64

- Sizes of C Objects (in Bytes)

<i>C Data Type</i>	<i>Typical 32-bit</i>	<i>Intel IA32</i>	<i>x86-64</i>
• unsigned	4	4	4
• int	4	4	4
• long int	4	4	8
• char	1	1	1
• short	2	2	2
• float	4	4	4
• double	8	8	8
• long double	8	10/12	16
• char *	4	4	8

Or any other pointer

Whose Assembler?

Intel/Microsoft Format

```
lea    eax, [ecx+ecx*2]
sub   esp, 8
cmp   dword ptr [ebp-8], 0
mov   eax, dword ptr [eax*4+100h]
```

GAS/Gnu Format

```
leal  (%ecx,%ecx,2),%eax
subl $8,%esp
cmpl $0,-8(%ebp)
movl $0x100(,%eax,4),%eax
```

Intel/Microsoft Differs from GAS in syntax

- Operands listed in opposite order

mov Dest, Src

movl Src, Dest

- Constants not preceded by '\$', Denote hex with 'h' at end

100h

\$0x100

- Operand size indicated by operands rather than operator suffix

sub

subl

- Addressing format shows effective address computation

[eax*4+100h]

\$0x100(,%eax,4)

C operators – Assembly Equivalents?

Operators

() [] -> .
! ~ ++ -- + - * & (type) sizeof
* / %
+ -
<< >>
< <= > >=
== !=
&
^
|
&&
||
?:
= += -= *= /= %= &= ^= != <<= >>=
,

Associativity

left to right
right to left
left to right
right to left
right to left
left to right

Note: Unary +, -, and * have higher precedence than binary forms