

# **Chapter 9: Virtual Memory**

## **Topics**

- **Motivations for virtual memory VM (not to be confused with virtual machines VM)**
- **Address translation**
- **Accelerating translation with TLBs**

# Announcements

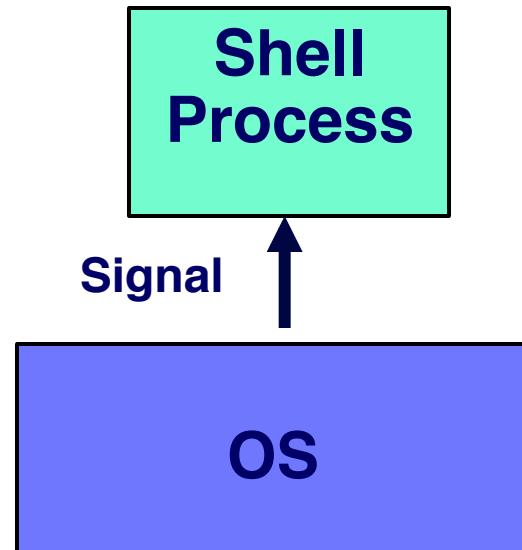
- **Shell lab is due Monday Dec 8 by 8 am**
  - Interview grading time slots for next week available later this week
- **Last Recitation Exercise #5 due Friday Dec 12 by 5 pm**
  - Upload to moodle or hand into TA at TA office hours
- **Reading:**
  - Read Chapter 9, except 9.6 and 9.7 (no case study and no memory mapping, can also skip multi-level page tables)
- **Need a volunteer for FCQs on Thursday**
- **Final exam is Thursday Dec 18, 4:30-7 pm, more next week**

# Recap

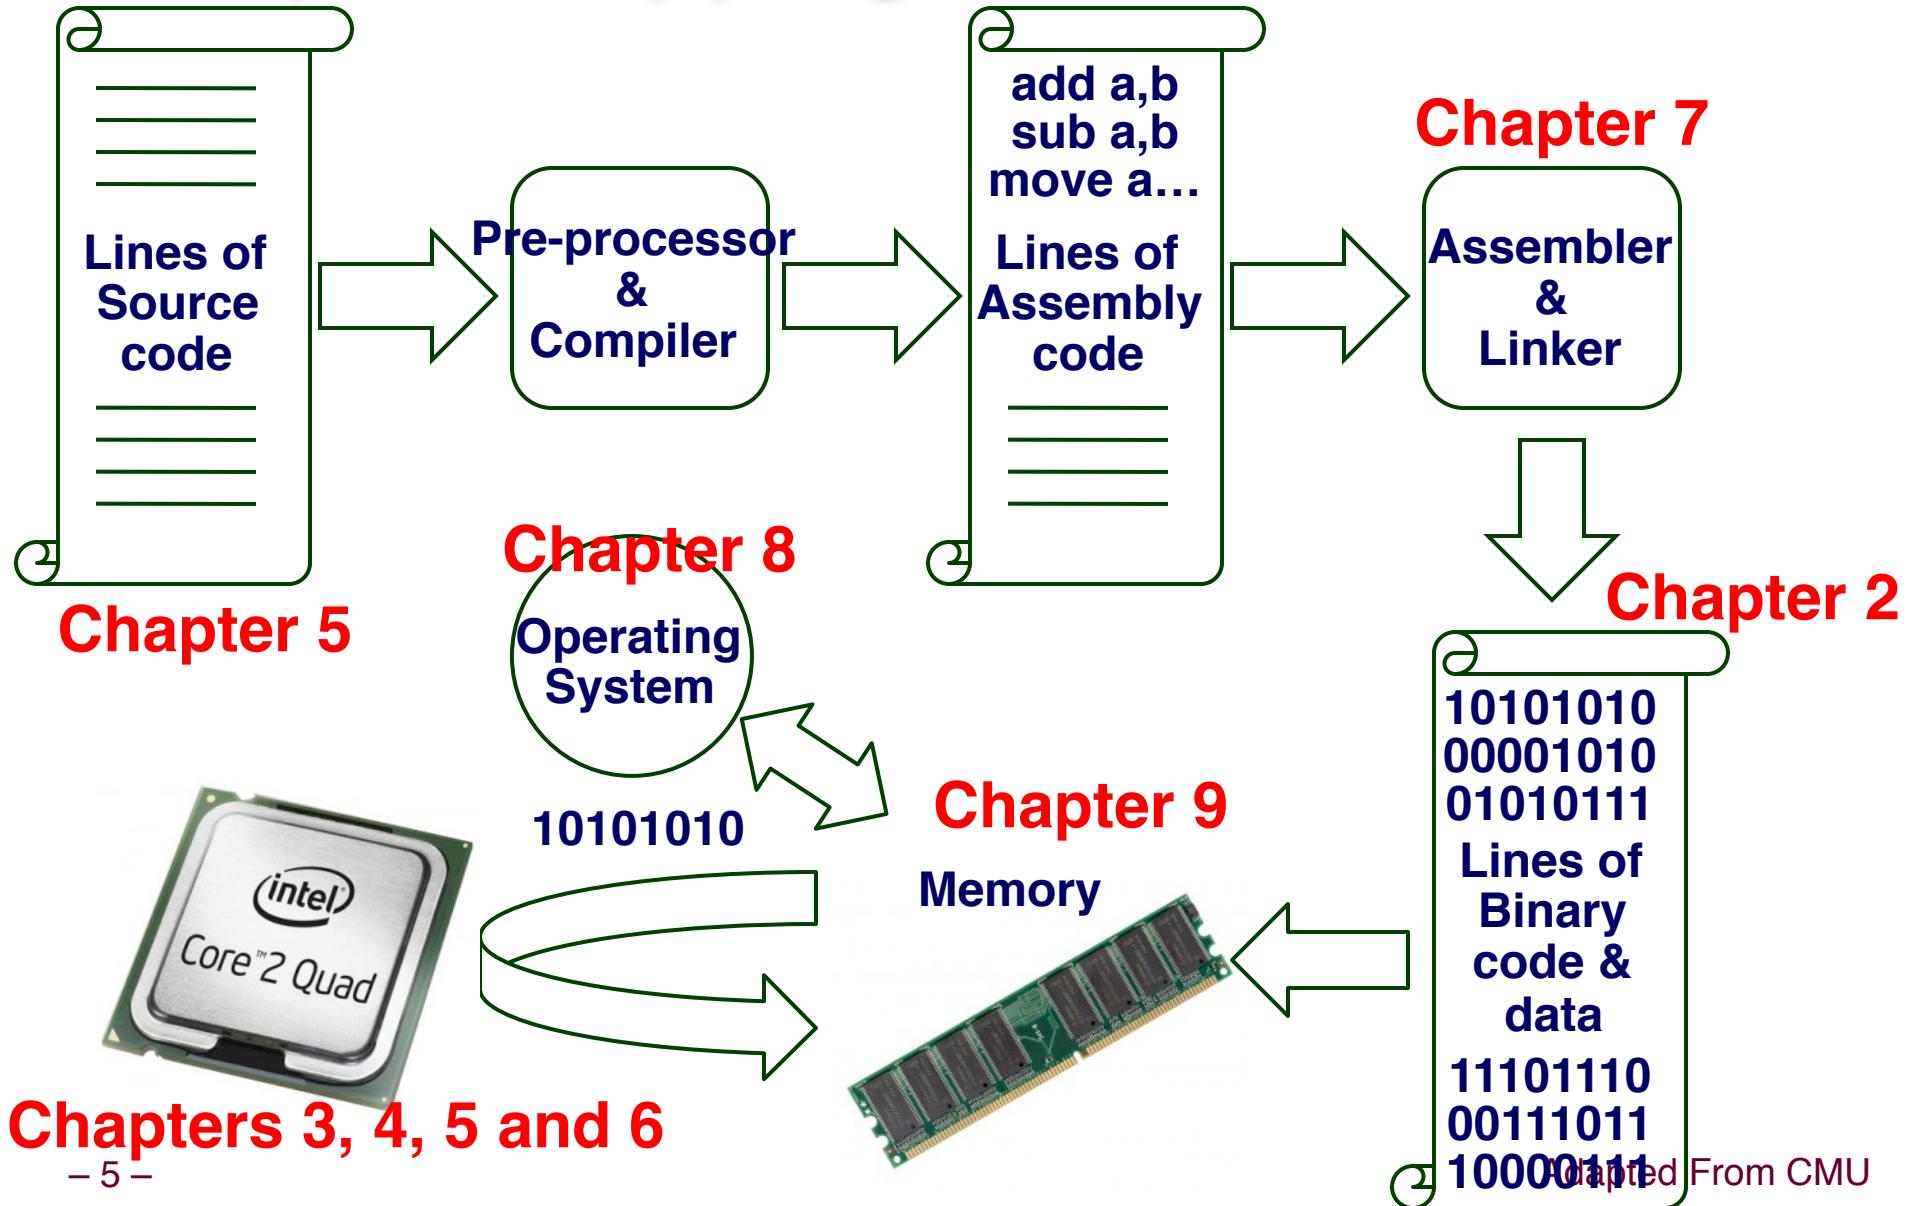
- How does software execute on modern computers?
    - It interacts with OS and hardware via HW/SW interrupts, faults, ...
  - A Process is an actively executing program
    - Use `fork()` to create a child process in Unix/Linux
    - Use `exit()` to finish executing a process
    - Parent calls `wait()`/`waitpid()` to free child process' resources
  - A Shell process
    - Forked child can run in foreground or background (type “`&`” at end of command line)
    - Use signals to properly reap child processes and respond to CTL-C/CTL-Z
- ```
int main() {
    char cmdline[MAXLINE];
    ...
    while (1) {
        ...
        /* get string */
        Fgets(cmdline, MAXLINE, stdin);
        ...
        /* execute string */
        if ((pid = Fork()) == 0) {
            /* child runs user job */
            execve(args from cmdline);
        }
        ...
    }
}
```

# Recap

- Your shell will have to handle signals
  - CTL-C sends a SIGINT
  - CTL-Z sends a SIGTSTP
  - Child process exiting sends a SIGCHLD
- Register your signal handler with  
`signal()` or `sigaction()`
- Signals don't queue
  - pending bit indicates at least one signal of that type has occurred – so add a while loop to handler
- Race conditions may occur
  - If signal handler modifies global state while normal control flow also modifies global state – see Ch 8 slides' 2 examples
  - Use masking of signals (or synchronization) to mitigate race conditions



# Chapter Mapping



Chapters 3, 4, 5 and 6

- 5 -

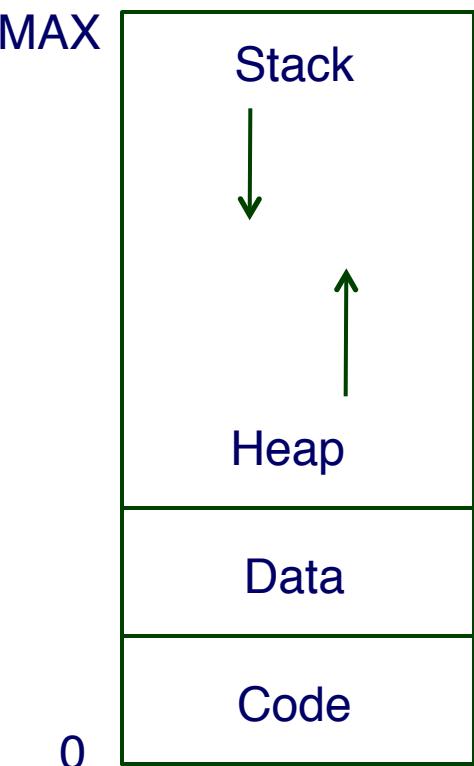
Adapted From CMU

# Motivations for Virtual Memory

Provides the Abstraction that a Process is executing in its own Address Space of memory from address 0 to address MAX

For a Single Process:

- Compiler and linker can generate code targeting a [0..MAX] memory paradigm
- Can be placed anywhere in memory, and virtual addresses are *translated* to physical addresses at run time as each instruction executes
- frees a process and the compiler from having to worry where in physical memory the process is or will be located



# Virtual Addresses

- Assembly code is compiled assuming virtual addresses
  - address of each instruction is a virtual address, not a physical address
  - data referenced in memory is referred to by its virtual address, not its actual physical address
  - executing program is completely unaware of where it is actually located in physical memory

| addresses | code instructions                                                                                                                |
|-----------|----------------------------------------------------------------------------------------------------------------------------------|
| 0xFFC0    | ...<br>pushl %ebp<br>movl %esp,%ebp<br>movl 12(%ebp),%eax<br>movl 8(%ebp),%edx<br>movl %ebp,%esp<br>popl %ebp<br>jmp *eax<br>... |
| 0xFFD8    |                                                                                                                                  |
| .         |                                                                                                                                  |
| .         |                                                                                                                                  |
| .         |                                                                                                                                  |

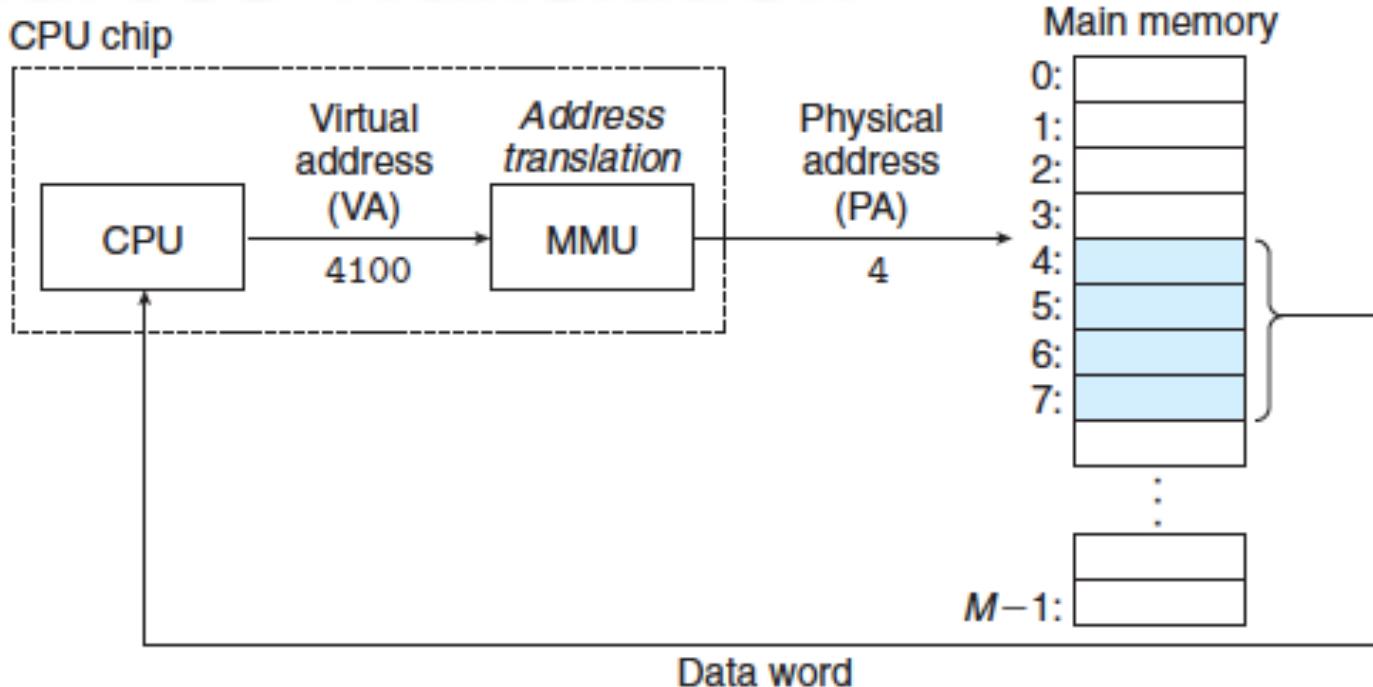
# MMU Translates Virtual Addresses to Physical Addresses

As the program executes, a Memory Management Unit (MMU) in the CPU does the following:

- When fetching the next instruction, the MMU translates the virtual address of the instruction into a physical address where the instruction is actually located and fetches it
- When an instruction references a data location in memory, the MMU translates the data's virtual address into its actual physical address, and fetches the data

| addresses | code instructions                                                                                             |
|-----------|---------------------------------------------------------------------------------------------------------------|
| 0xFFC0    | ...<br>pushl %ebp<br>movl %esp,%ebp<br>movl 12(%ebp),%eax<br>movl 8(%ebp),%edx<br>movl %ebp,%esp<br>popl %ebp |
| .         |                                                                                                               |
| .         |                                                                                                               |
| .         |                                                                                                               |
| 0xFFD8    | jmp *eax<br>...                                                                                               |

# MMU: Virtual $\rightarrow$ Physical Address Translation



**The MMU is a unit in the CPU that operates invisibly, underneath the program, translating its virtual addresses to physical addresses**

- The OS keeps track for each process of how virtual addresses are mapped to physical addresses and informs MMU

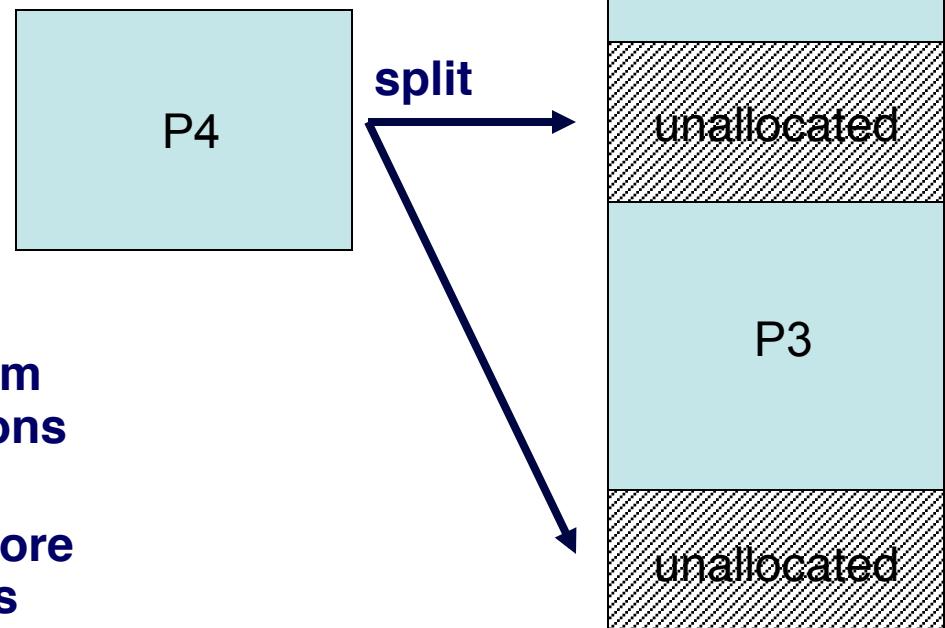
# Fragmentation Problem

Virtual addresses allow programs to be placed anywhere in main memory

- P1, P2, and P3 are placed contiguously, with some remaining unallocated memory

But over time, main memory becomes fragmented as processes finish

- P2 finishes, and its memory is deallocated
- P4 can't fit into any of the unallocated regions



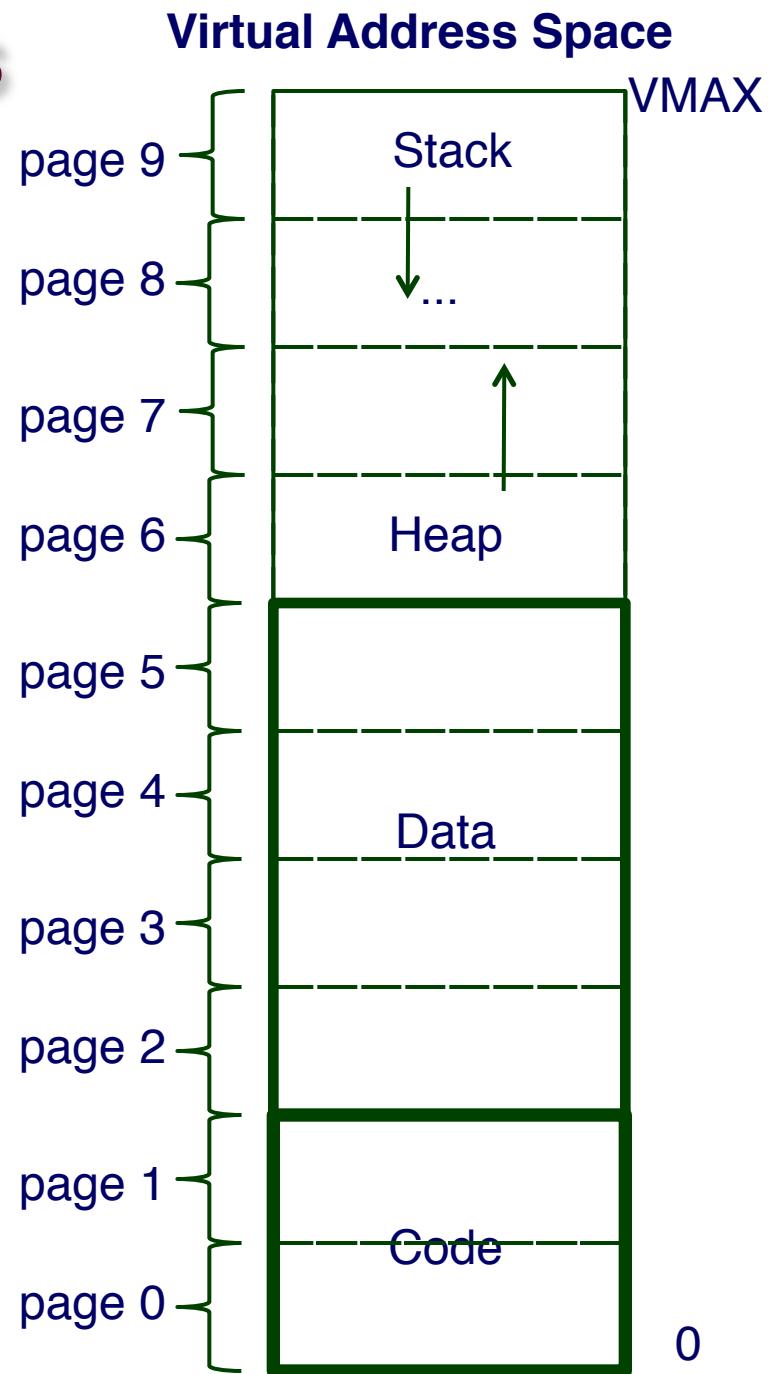
Better approach:

- Split P4 into pieces, and fit them each into the unallocated regions of memory
- Advantage: memory is used more efficiently, and more processes can run

# Dividing the Address Space into Pages

## Divide the Address Space into Fixed-size Pages

- Can put pages anywhere in physical memory – don't need contiguous space in physical memory to layout entire virtual address space of each process
- Solves fragmentation problem
- Typical page size is 4 KB – depends on OS
- virtual page #  
= virtual address/`sizeof(page)`  
= the most significant bits of the virtual address!



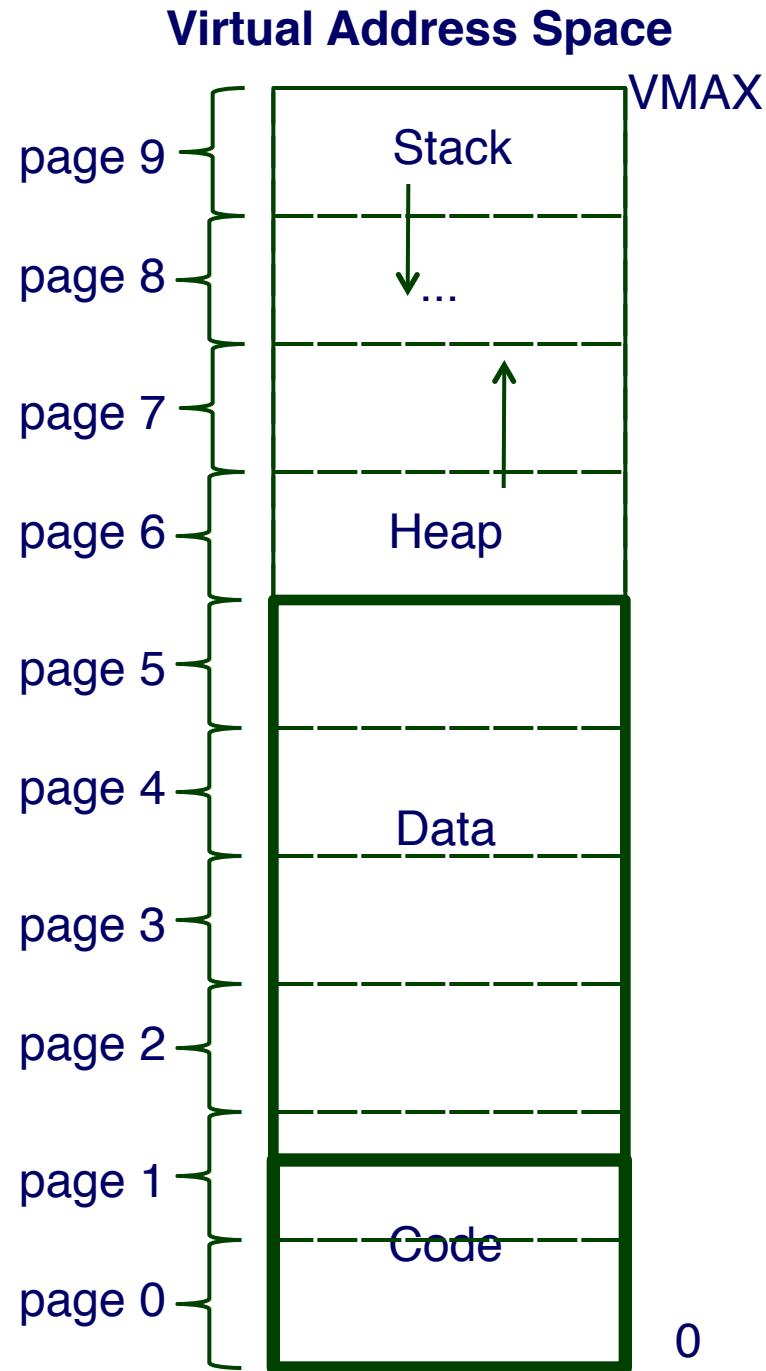
# Page Tables

**Divide Main Memory into Fixed-size pages as well**

**Must keep track of which virtual pages map to which physical pages**

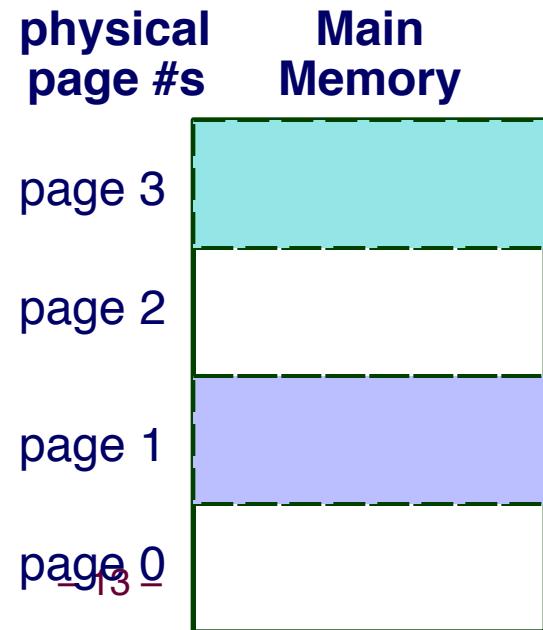
**Introduce a data structure per process called a **Page Table** that:**

- records virtual page -> physical page mapping
- The page table differs for different processes

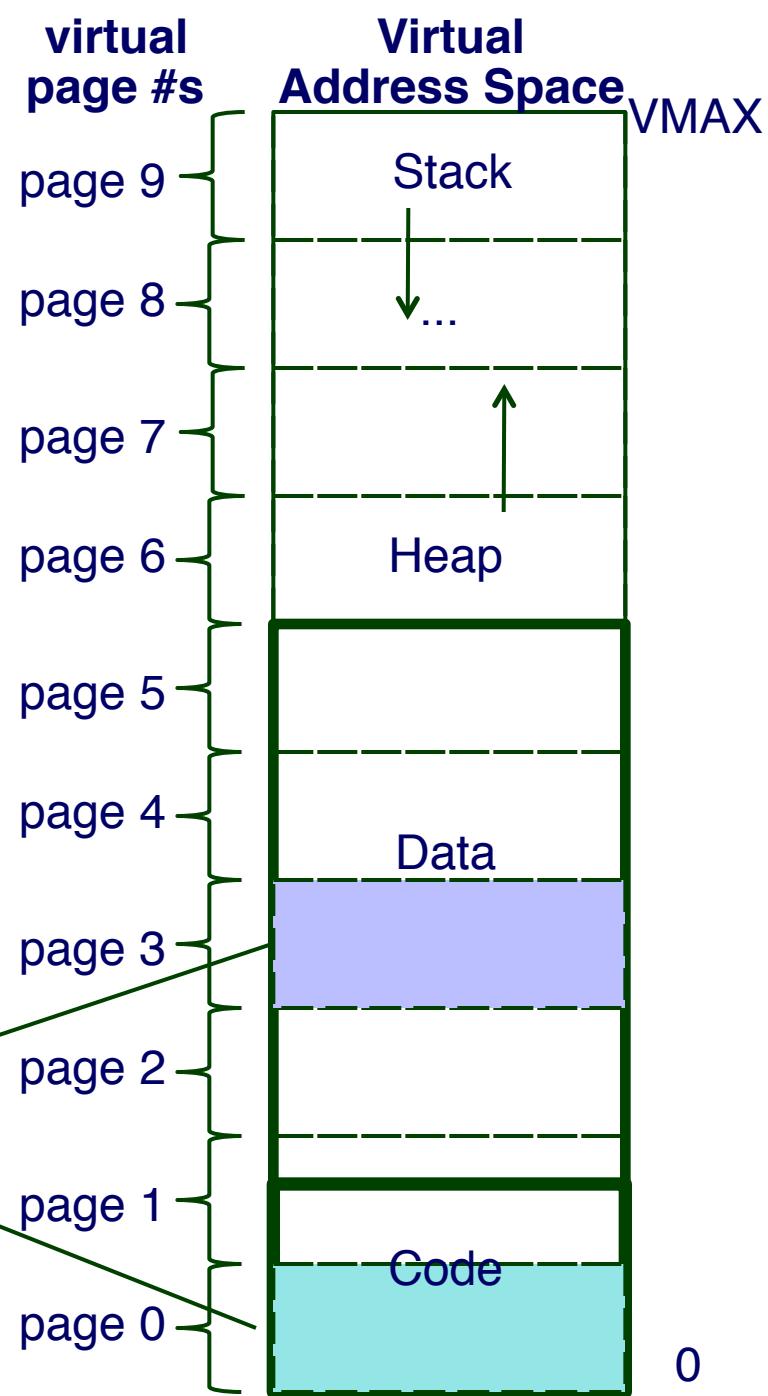


# Page Tables

valid bit indicates if page is in memory (more on this later)

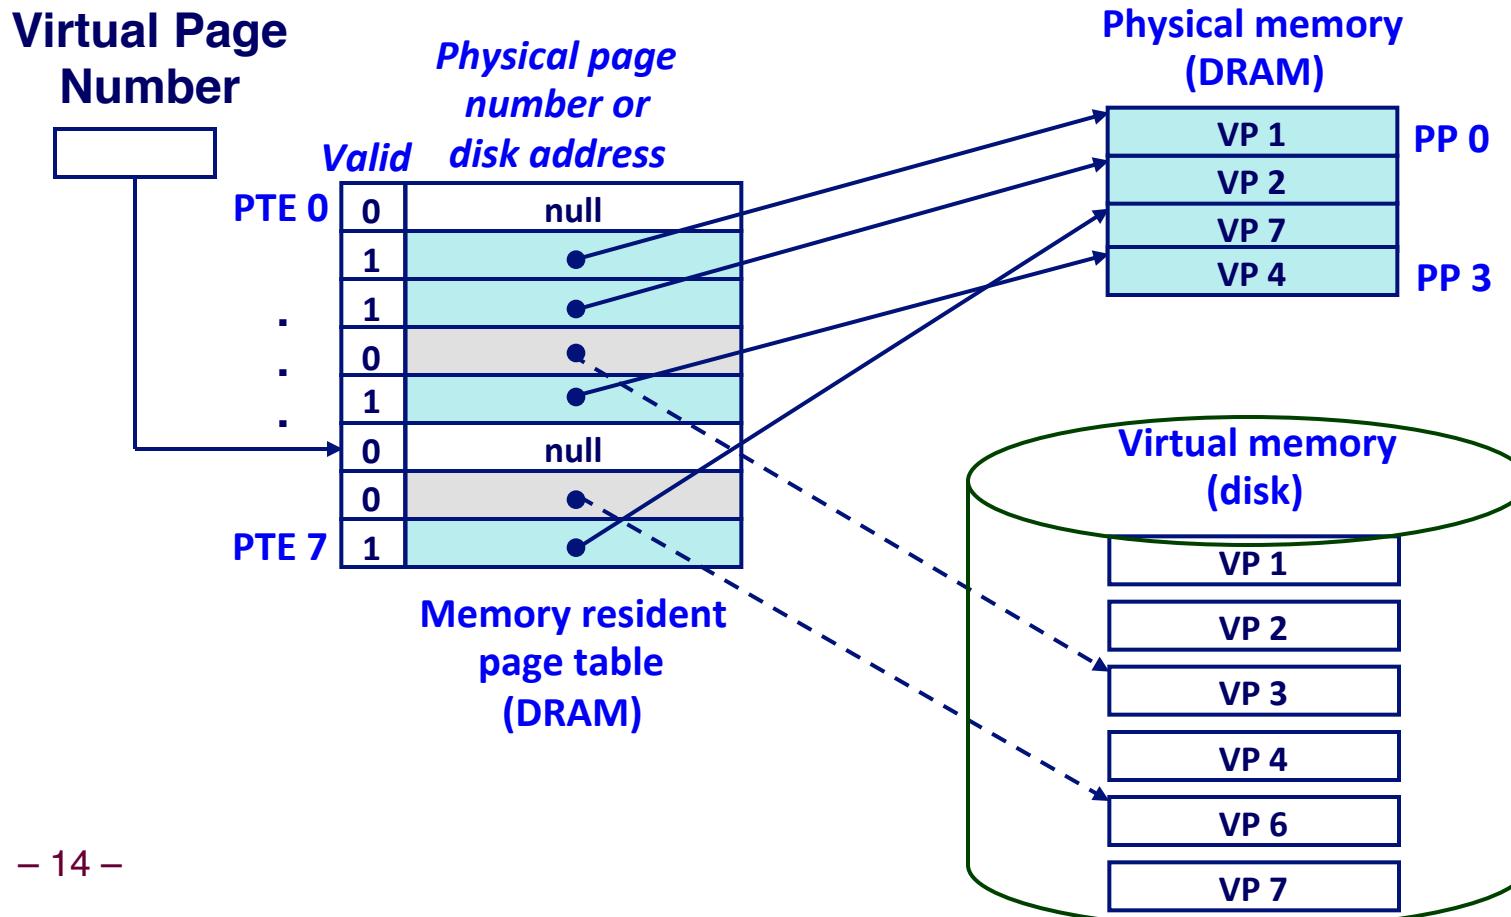


| Virtual Page # | Physical Page # | va-lid ? |
|----------------|-----------------|----------|
| 0              | 3               | 1        |
| 1              |                 | 0        |
| 2              |                 | 0        |
| 3              | 1               | 1        |
| 4              |                 | 0        |
| 5              |                 | 0        |
| 6              |                 | 0        |
| 7              |                 | 0        |
| 8              |                 | 0        |
| 9              |                 | 0        |



# Address Translation: Page Tables

- A **page table** is an array of page table entries (PTEs) that maps virtual pages to physical pages. Here: 8 VPs
  - Per-process kernel data structure in DRAM

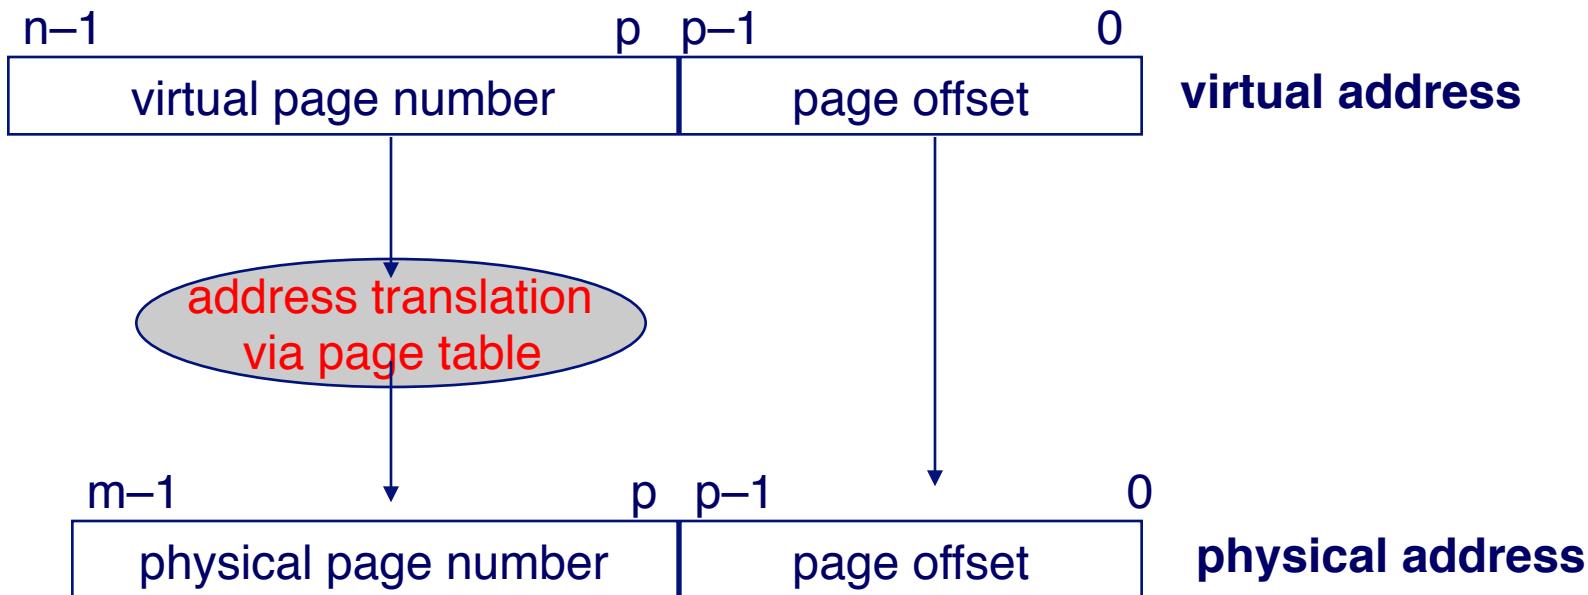


# VM Address Translation

## Parameters

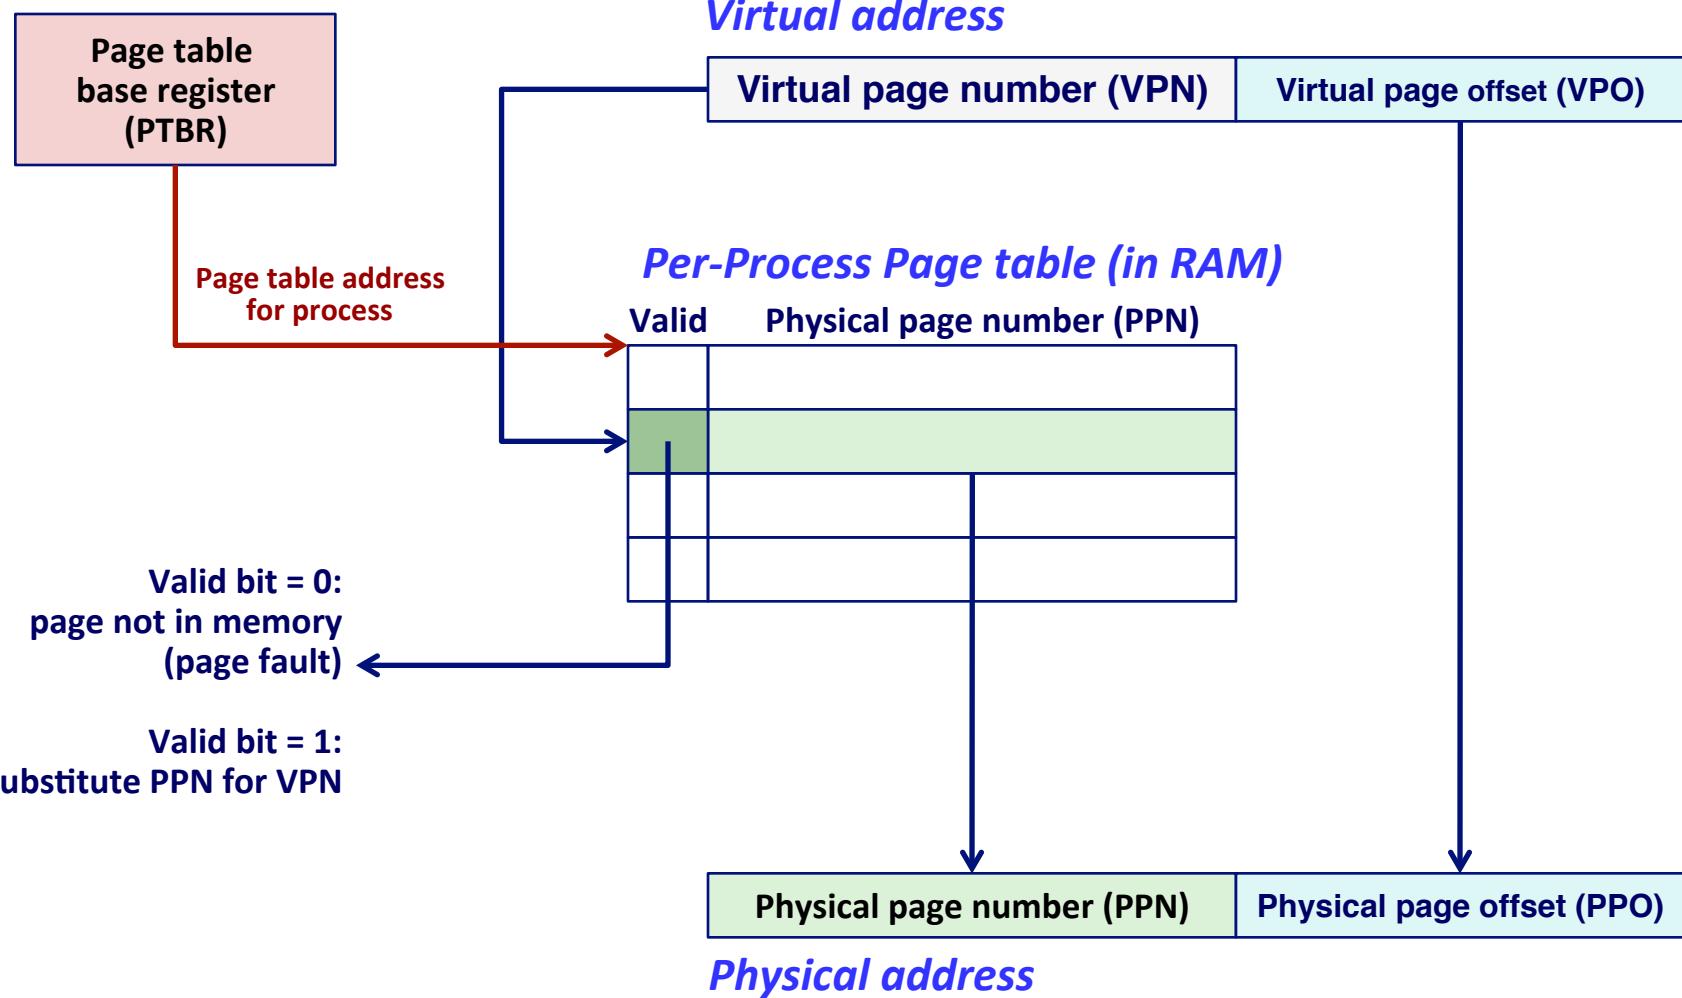
- $P = 2^p$  = page size (bytes).
- $N = 2^n$  = Virtual address limit
- $M = 2^m$  = Physical address limit

Example:  
page size = 4 KB =  $2^{12}$ ,  
for 32-bit addresses  
then page # =  $32-12$   
= 20 most significant bits



Page offset bits don't change as a result of translation

# Address Translation With a Page Table

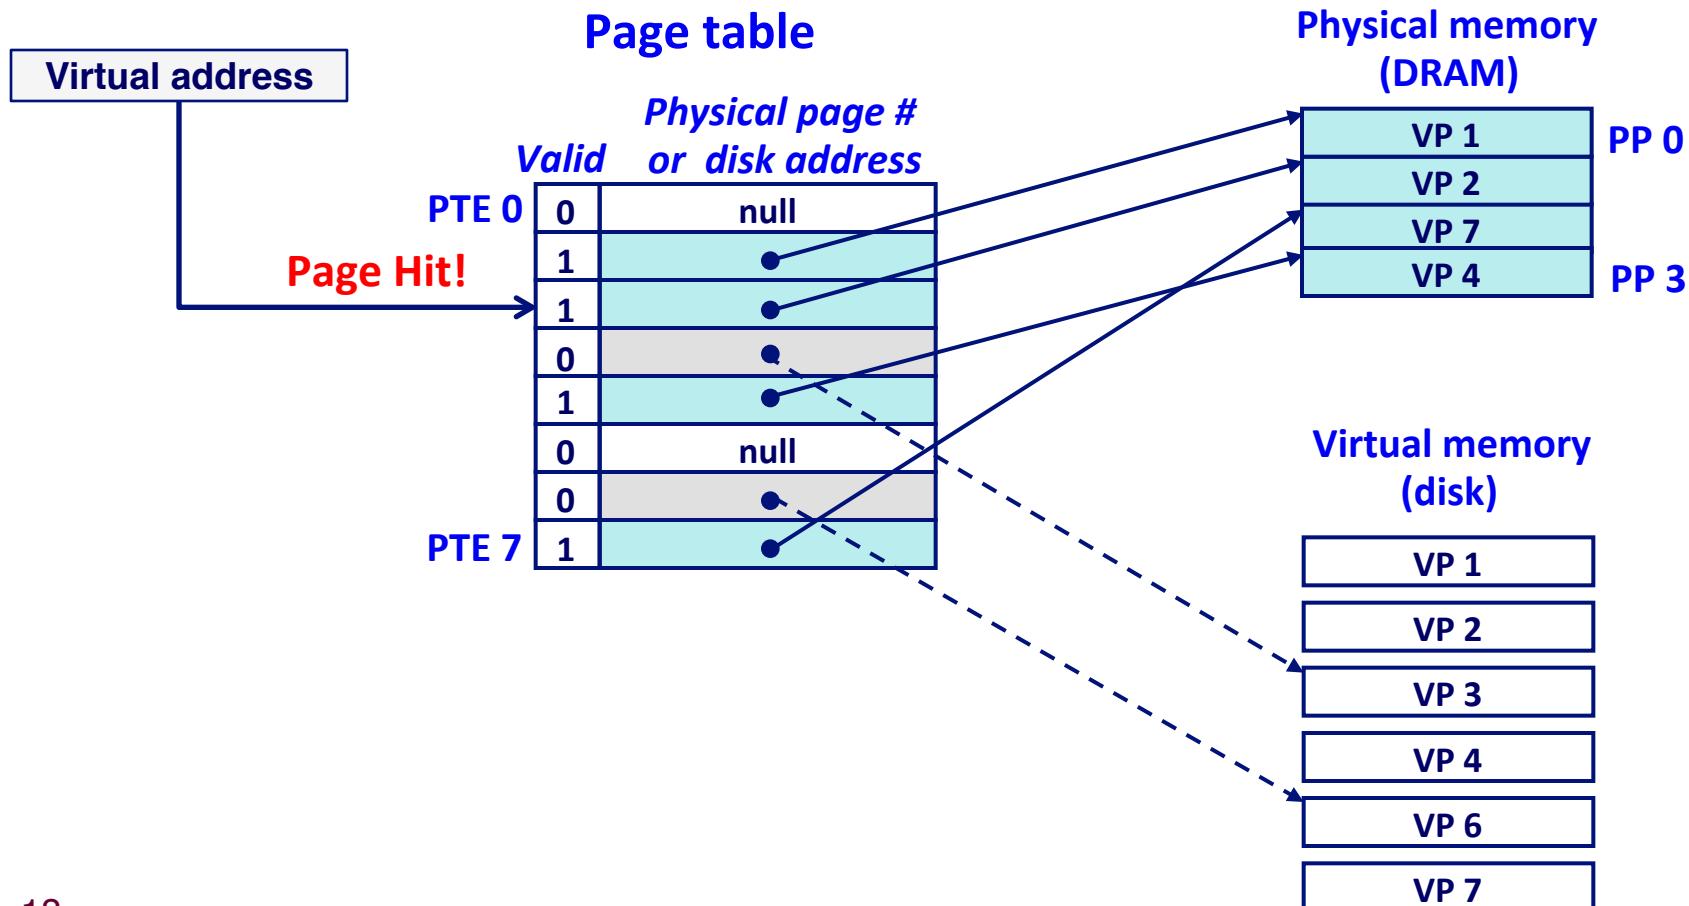


# Keep only a subset of pages in memory = Demand Paging

- Not all pages of address space (code, data, etc.) have to be kept in memory
  - Exploit locality, e.g. loops, to keep only the pages that were recently accessed (the “*working set*”) in memory
  - A small working set is best
  - Thrashing (performance meltdown) occurs when working set is too large for allocated memory
  - Main memory acts as a cache for disk
- Demand paging = load pages on demand as needed
  - Advantage: can now fit more processes into memory, since only a subset of pages are kept in memory at any one time
- Valid bit in page table keeps track of whether a page is in memory or on disk

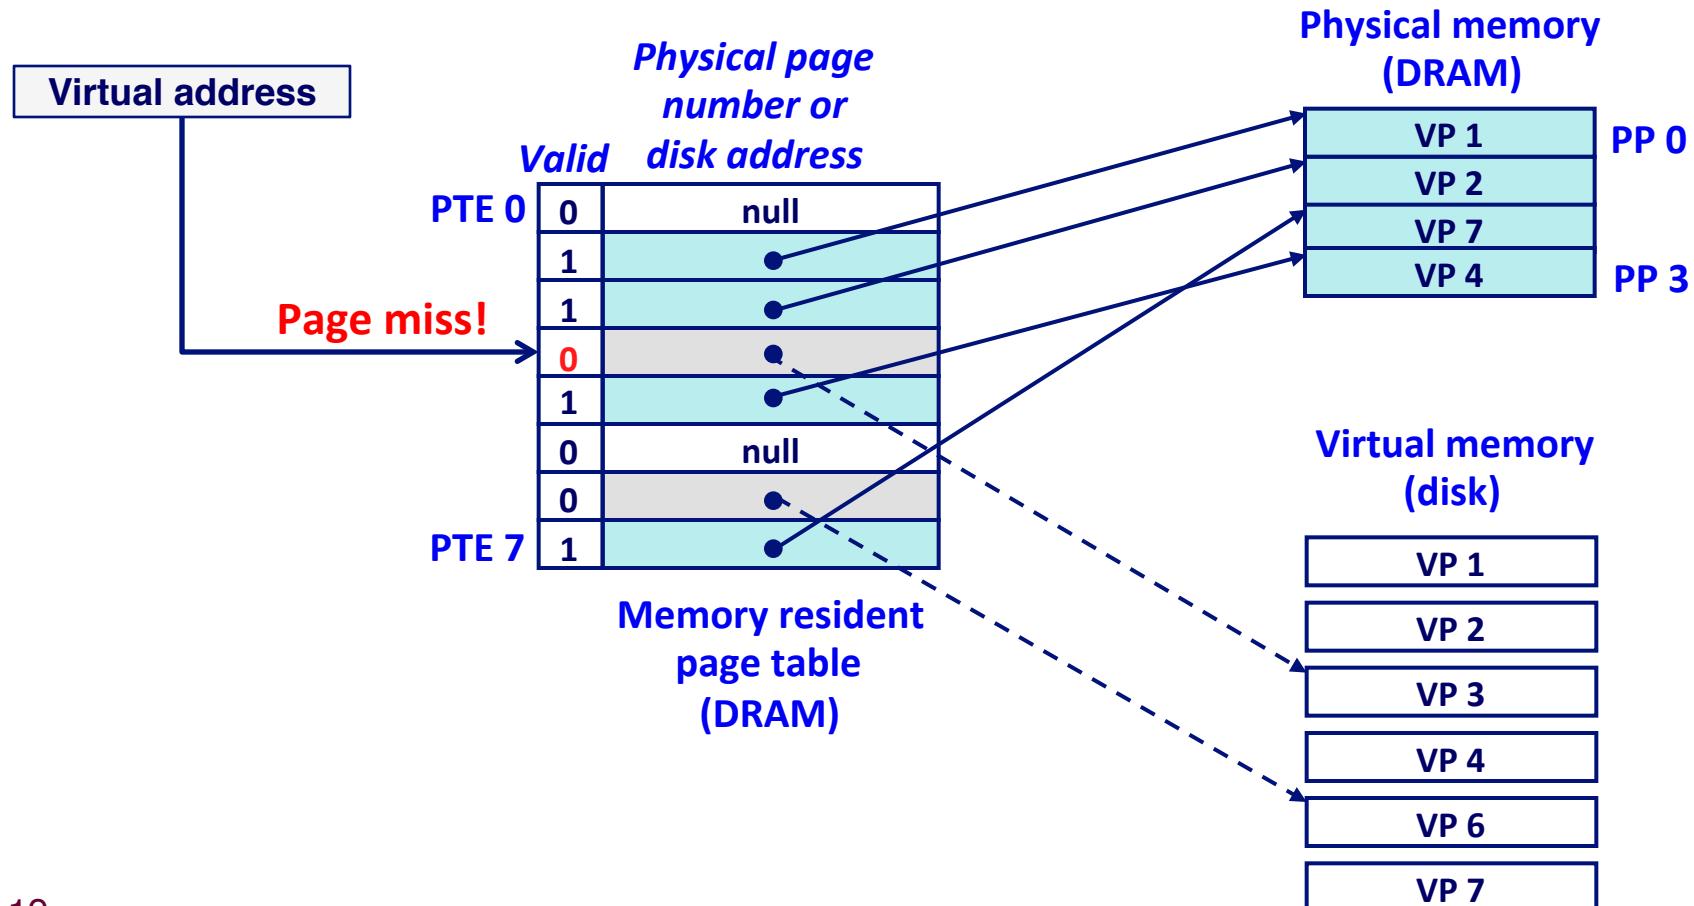
# Page Hit

- **Page hit:** reference to a virtual memory page that is in physical memory



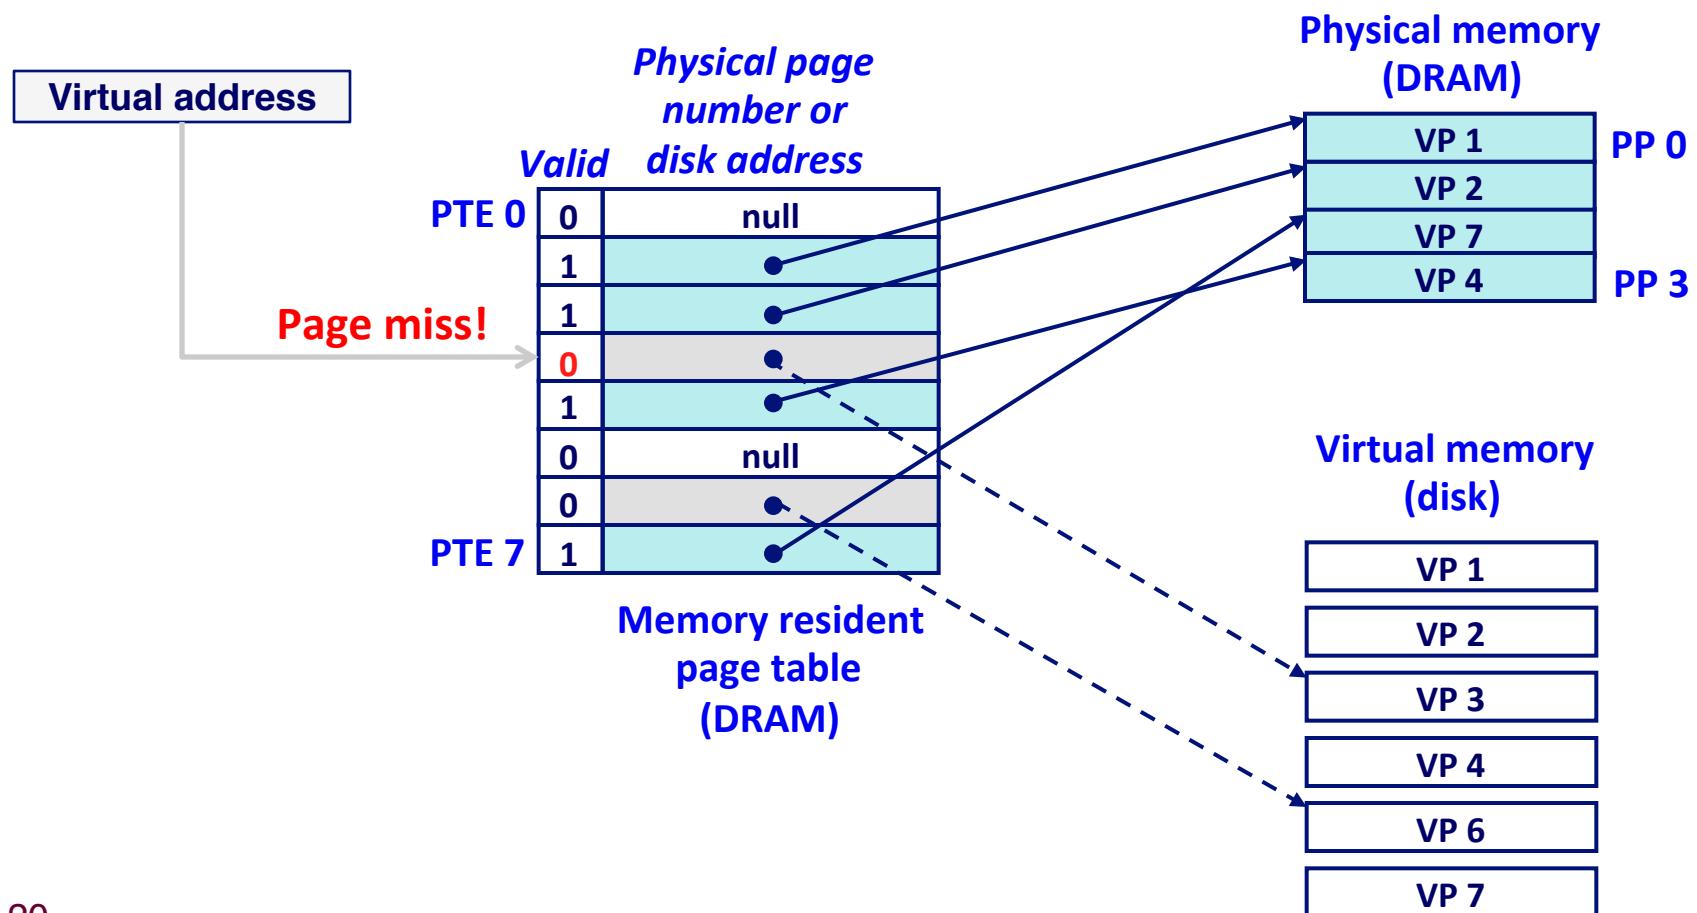
# Page Miss

- **Page miss:** reference to VM word that is not in physical memory



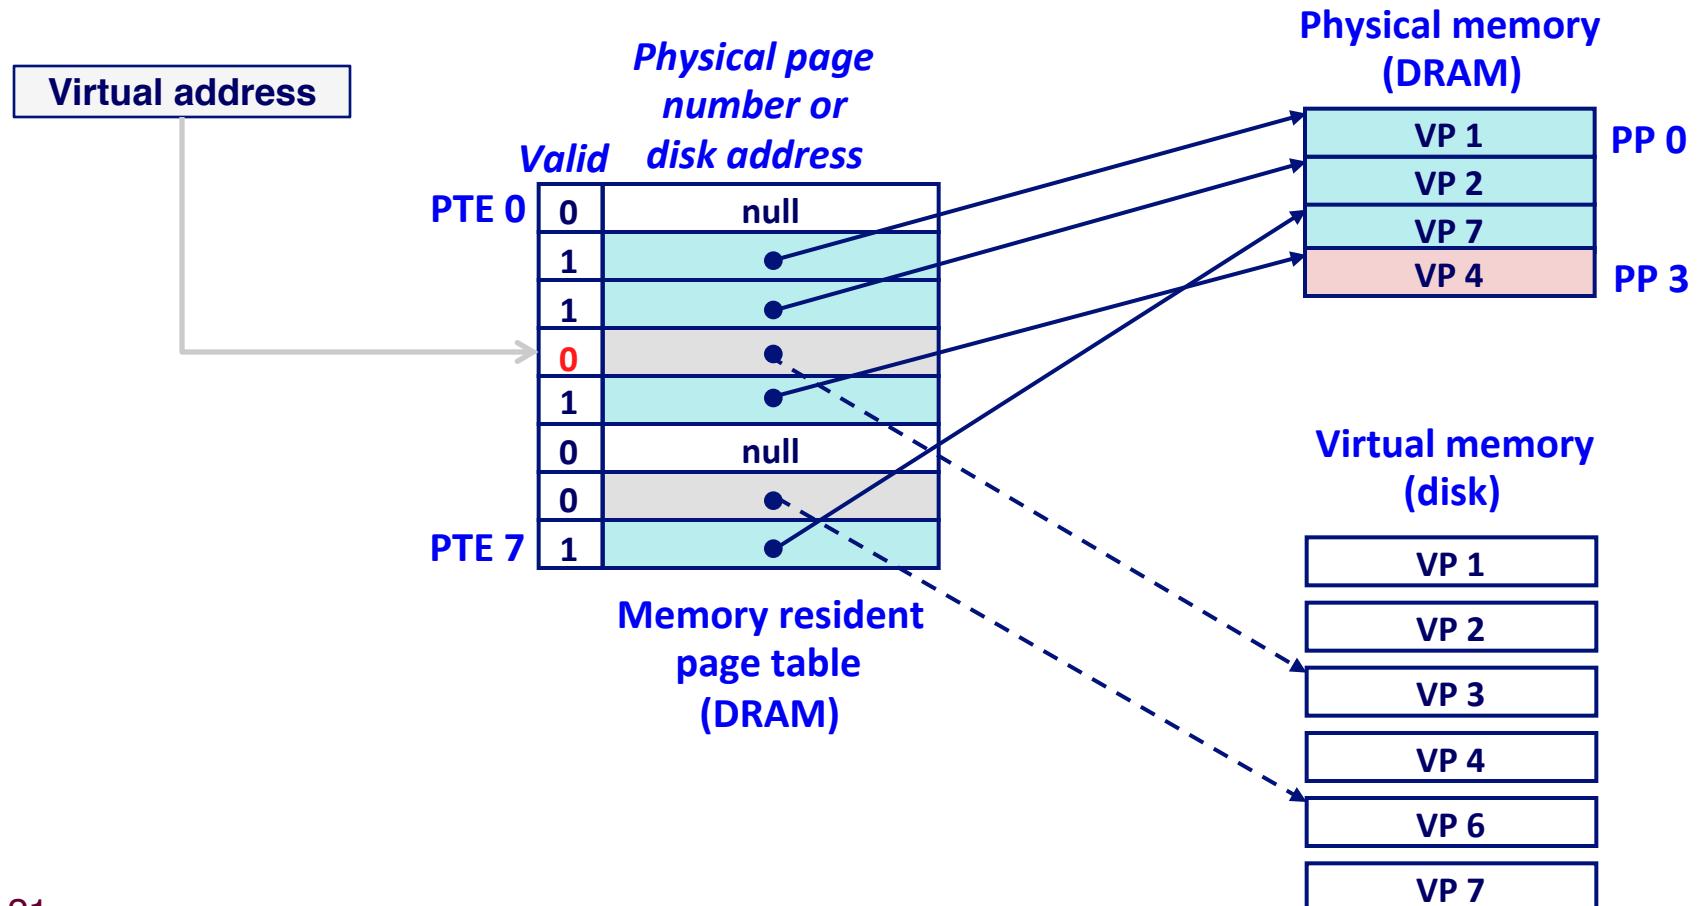
# Handling Page Fault

- Page miss causes page fault (an exception)



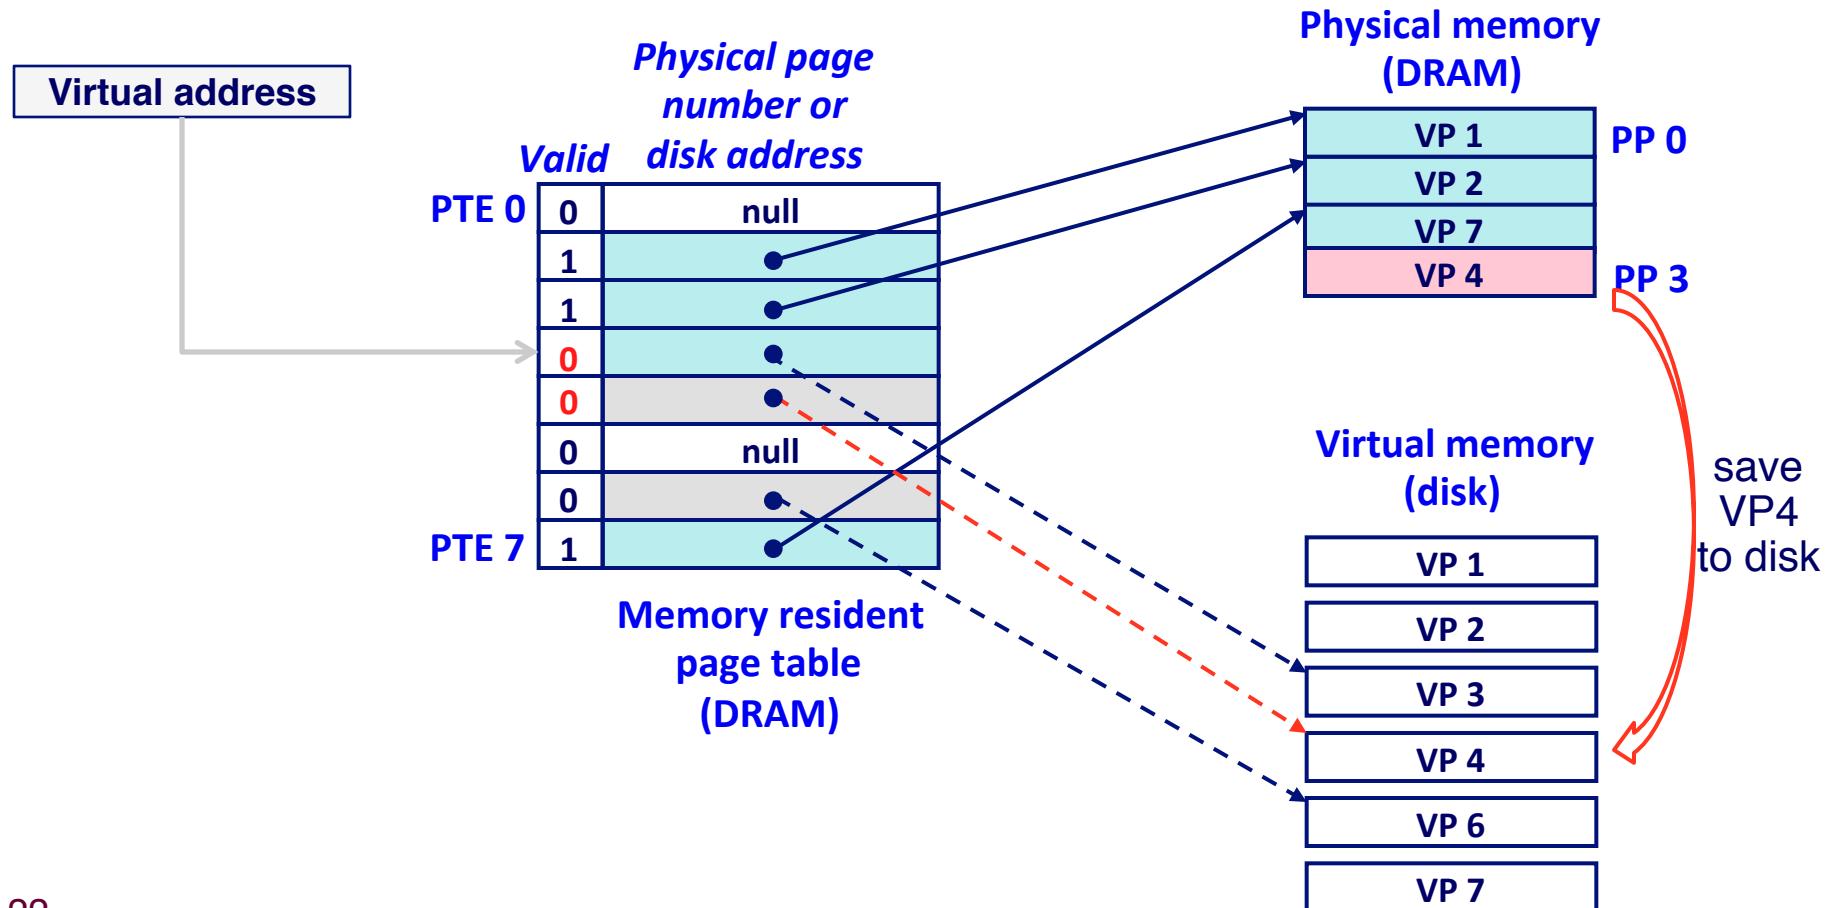
# Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)



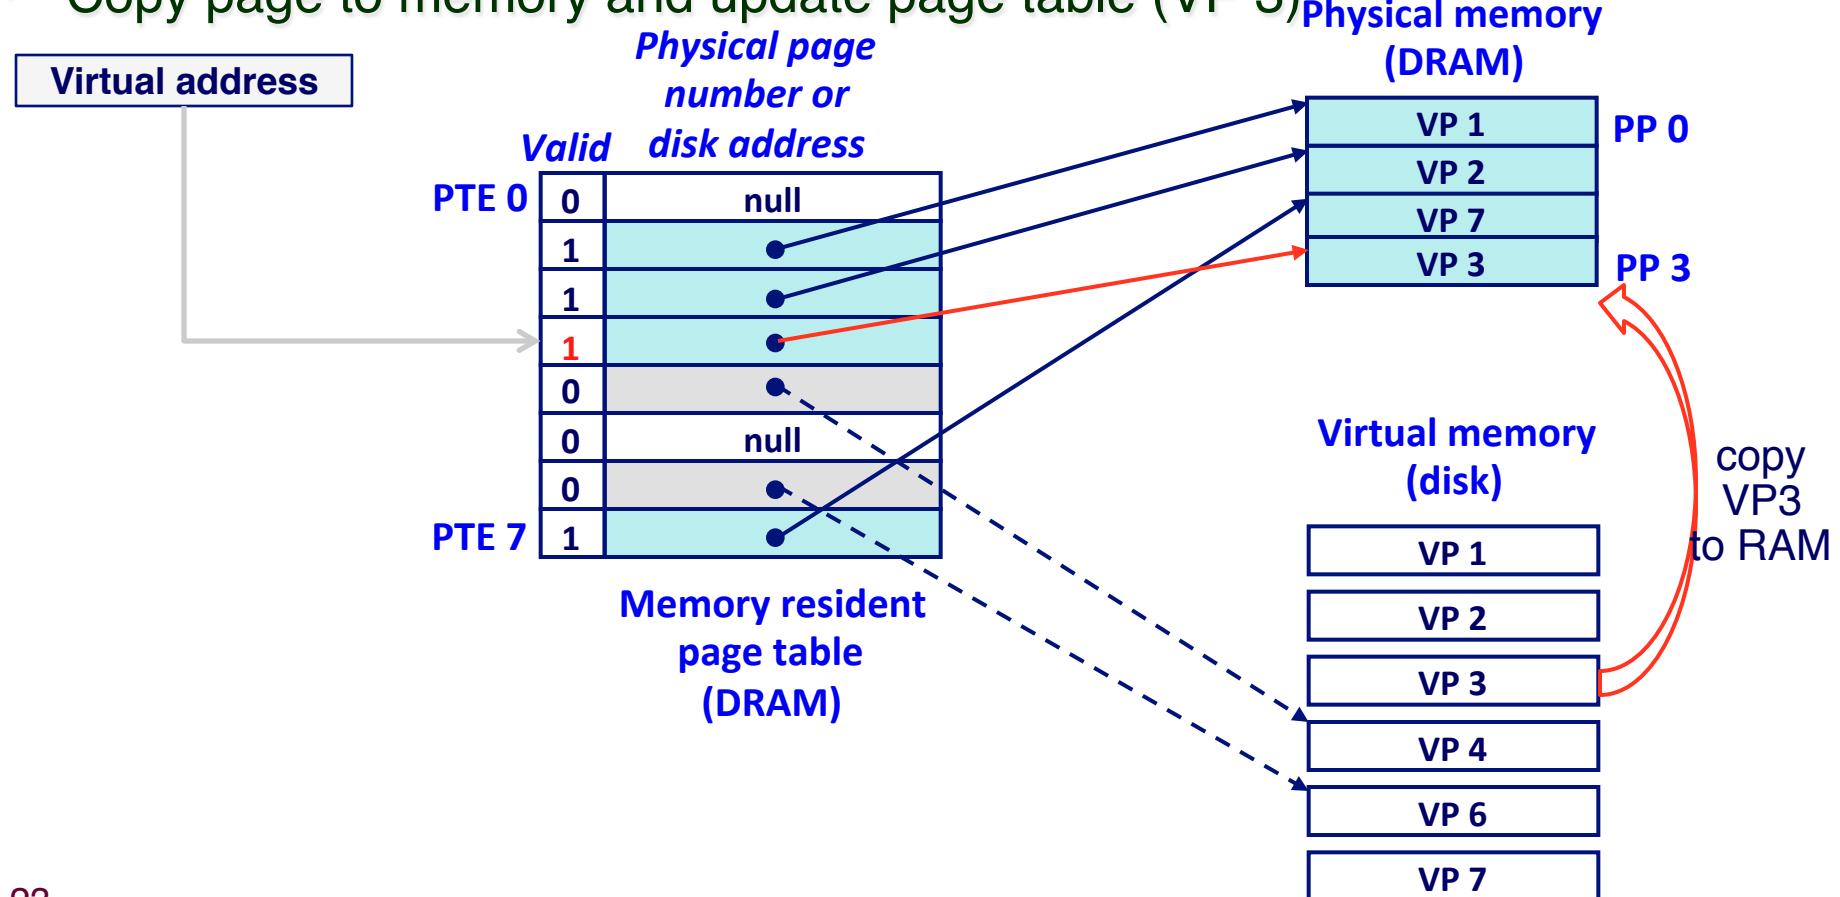
# Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)



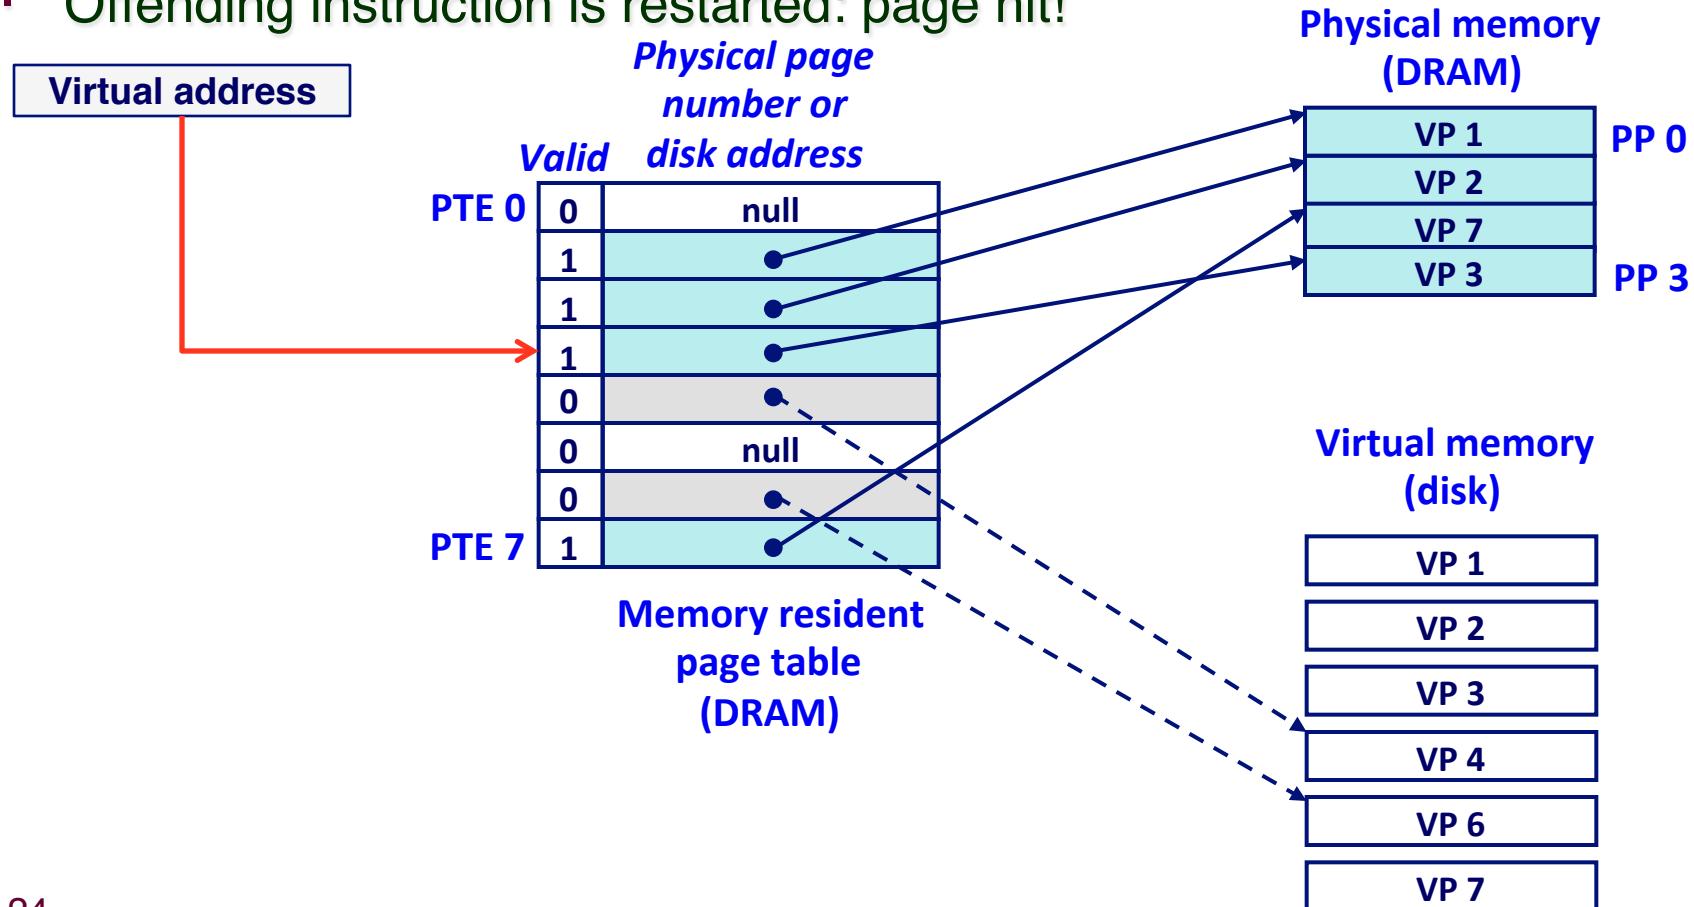
# Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)
- Copy page to memory and update page table (VP 3)



# Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)
- Offending instruction is restarted: page hit!



# Advantages of Page-based Virtual Memory (so far)

1. Simplifies memory management by compiler, program, and OS – all view the program as executing in an abstract [0..VMAX] address space
2. Can place processes anywhere in physical memory
3. Can fragment process into fixed-size pages that fit into unallocated memory fragments and thus use memory more efficiently
  - Processes no longer have to be contiguous.
  - Page table keeps track of where pages are placed physically in memory.
4. Demand paging keeps only a subset of recently accessed pages in memory – uses memory even more efficiently by allowing even more processes to fit into main memory

# More Advantages of Page-based Virtual Memory

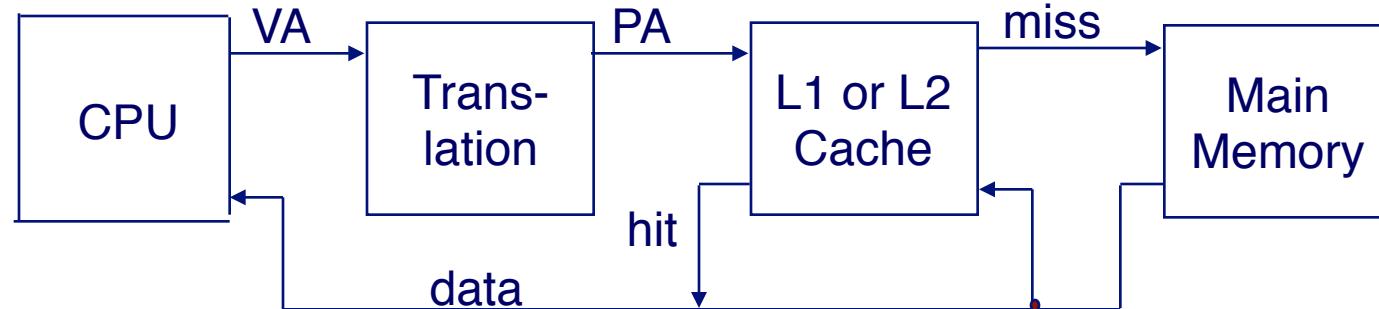
**5. Size of virtual address space VMAX is no longer limited to the amount of physical memory allocated to the process, but can equal or exceed the physical memory!**

- i.e. VMAX can be  $\geq$  PMAX
- Heap, stack, code and data segments can all be much larger now and not be constrained by physical memory – the set of addressable memory addresses is much larger

**6. Page tables help protect address spaces**

- A process X cannot easily write into the pages of another process Y's address space
  - OS typically assigns non-overlapping physical pages to fill in the entries of each process's page table
  - any virtual address X used will be translated by X's page table to a physical page of memory that is typically not shared with any other process like Y

# Integrating Virtual Memory and Cache



## Most Caches “Physically Addressed”

- Accessed by physical addresses
- Allows multiple processes to have blocks in cache at same time
- Allows multiple processes to share pages
- Cache doesn’t need to be concerned with protection issues
  - Access rights checked as part of address translation

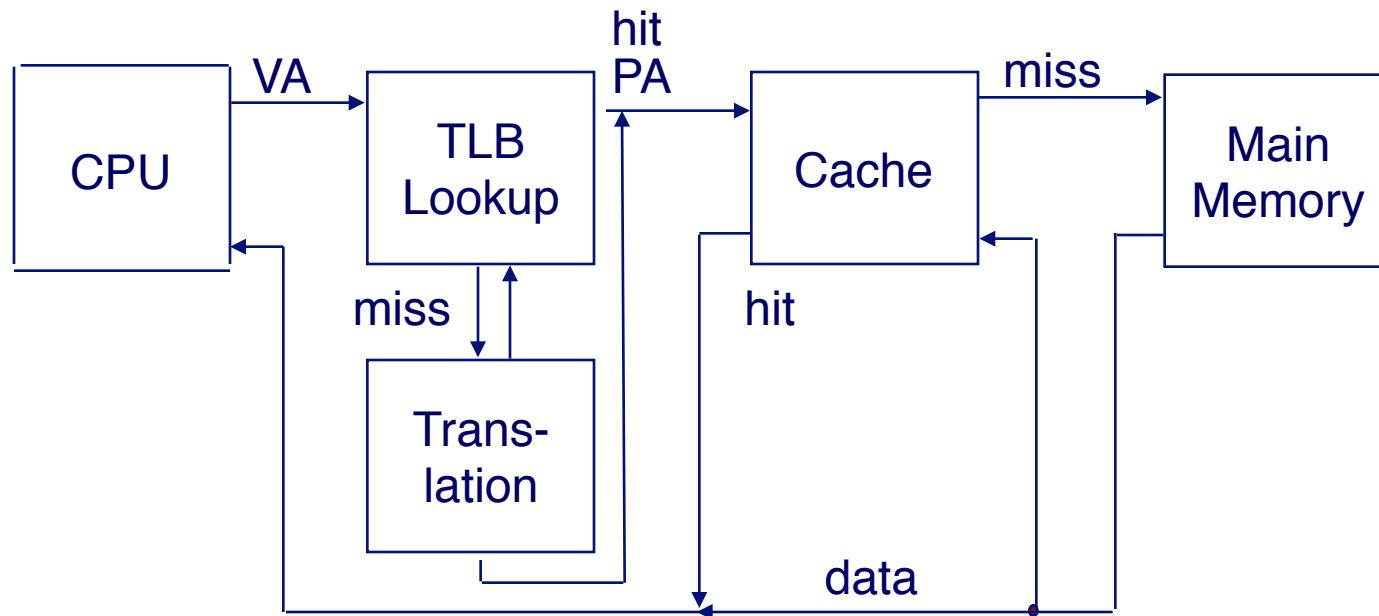
## Perform Address Translation Before Cache Lookup

- But this could involve a memory access itself (of the PTE) – Big Slowdown!
- Therefore, cache the page table entries too!

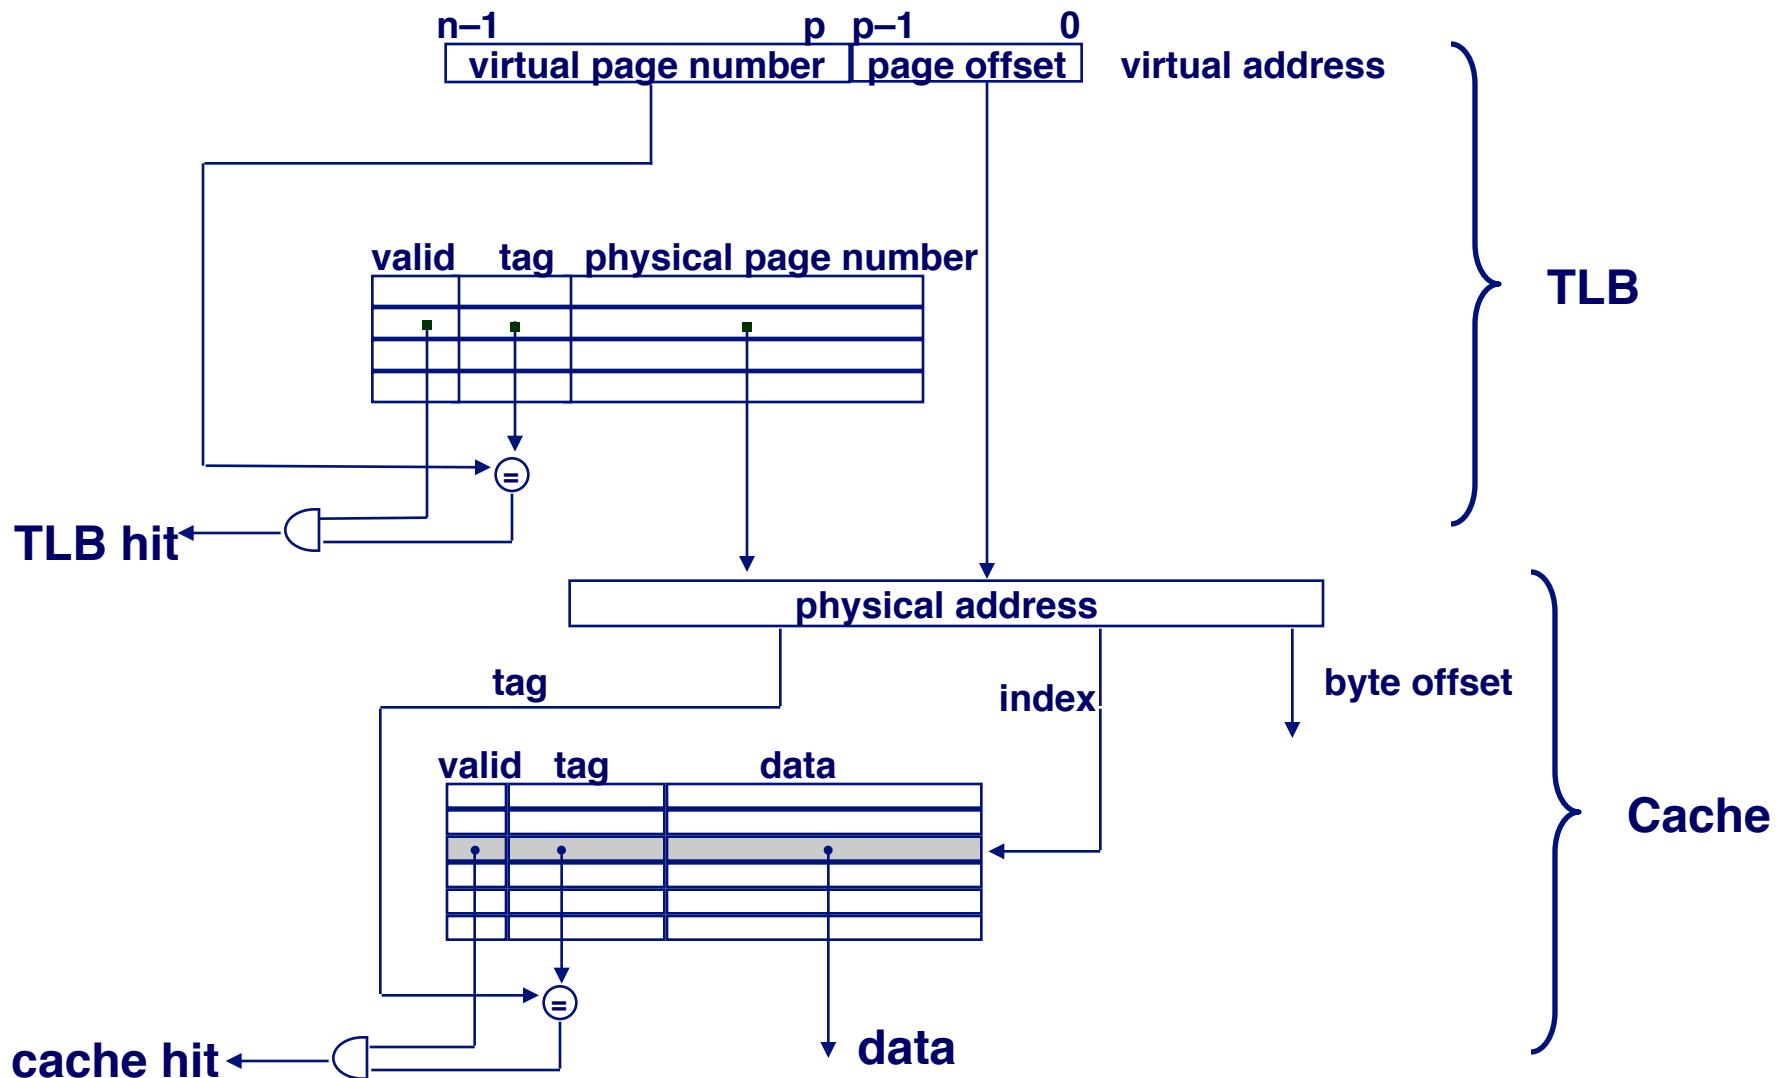
# Speeding up Translation with a TLB

## “Translation Lookaside Buffer” (TLB)

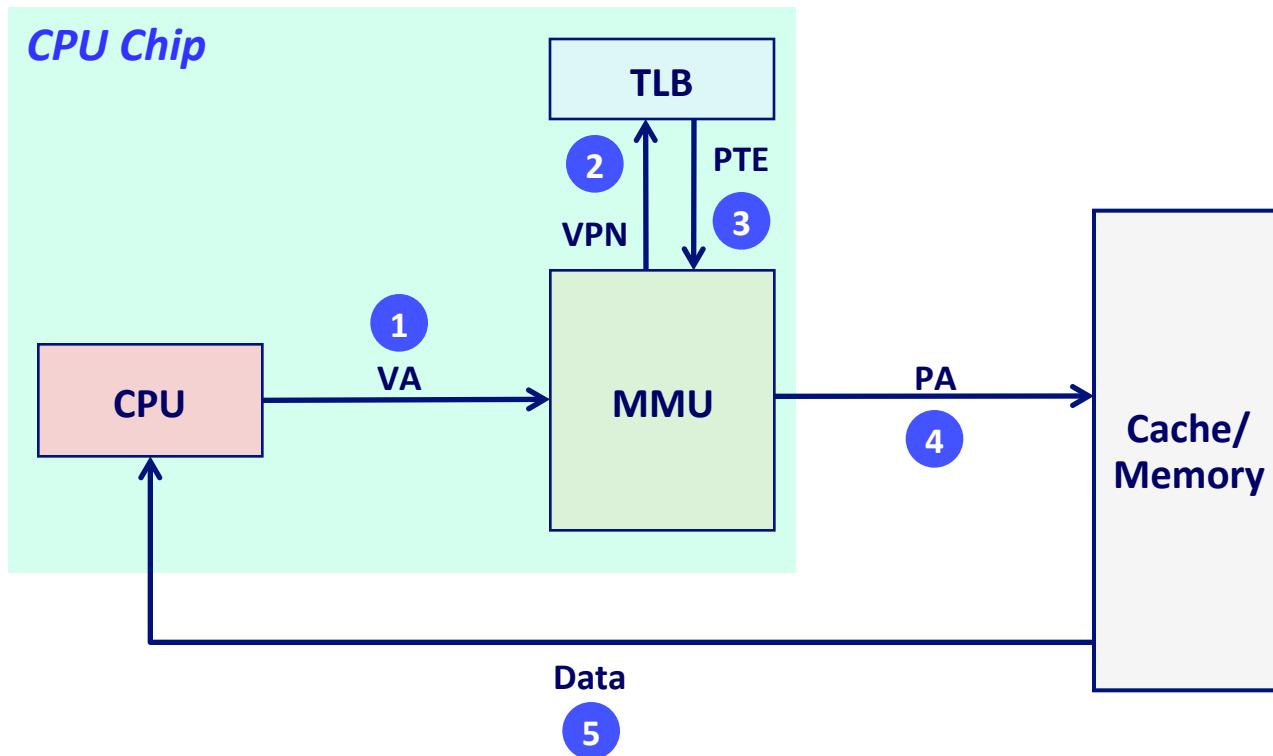
- Small hardware cache in MMU
- Maps virtual page numbers to physical page numbers
- Contains complete page table entries for small number of pages



# Address Translation with a TLB

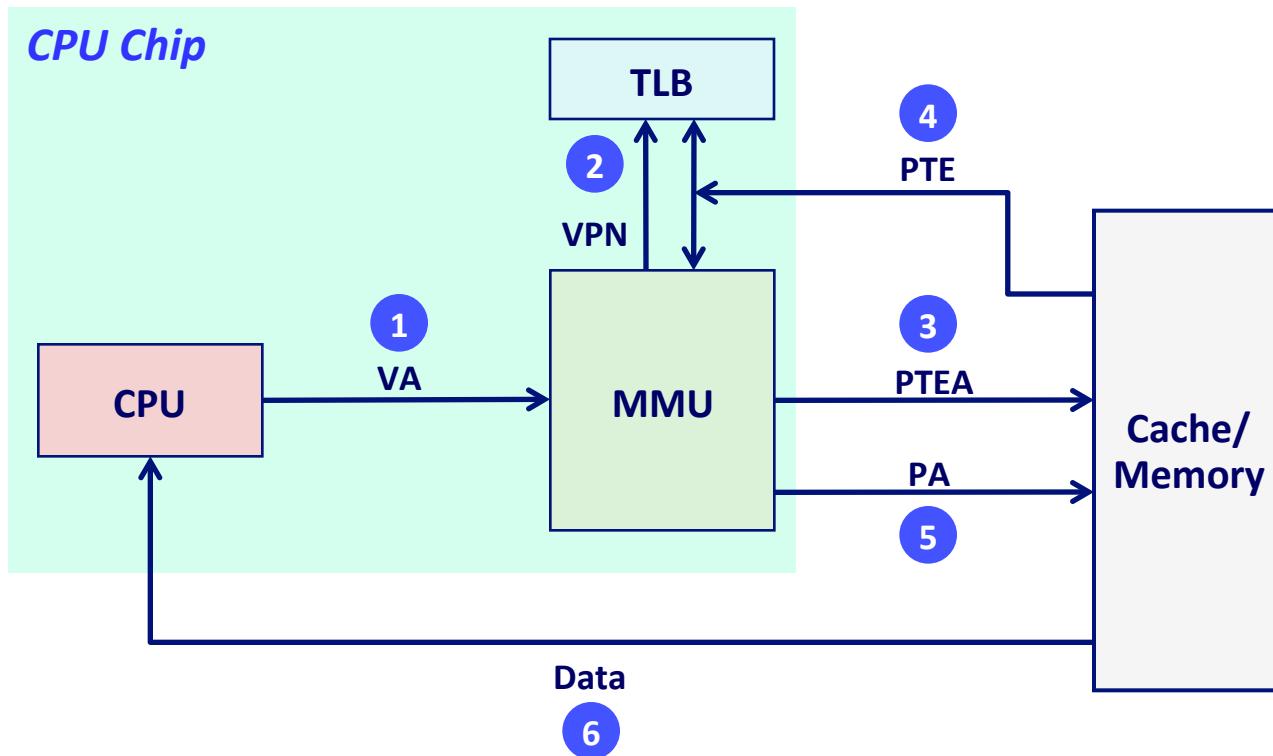


# TLB Hit



A TLB hit eliminates a memory access

# TLB Miss



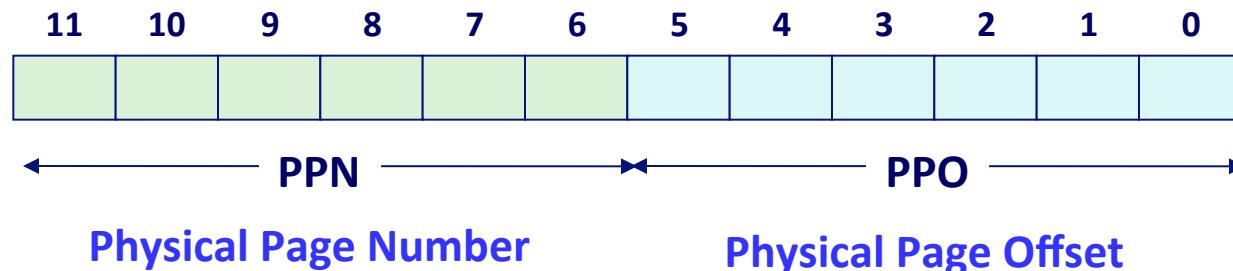
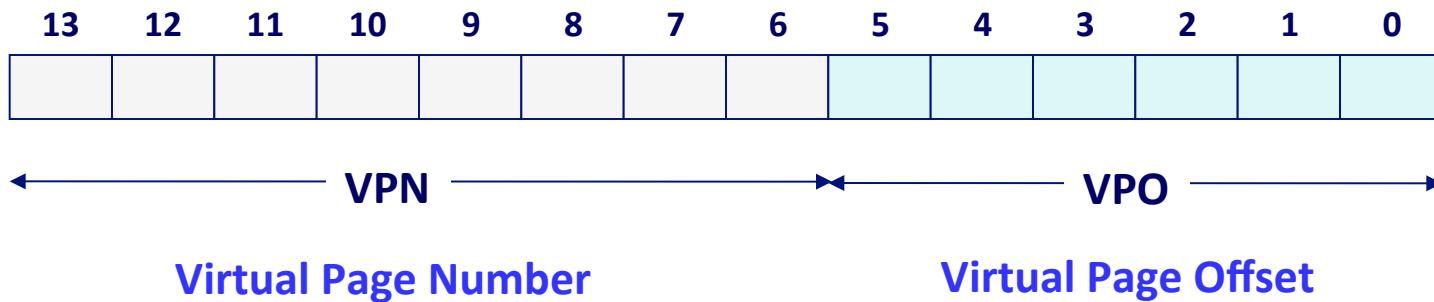
A TLB miss incurs an add'l memory access (the PTE)

Fortunately, TLB misses are rare

# Simple Memory System Example

## Addressing

- 14-bit virtual addresses
- 12-bit physical address
- Page size = 64 bytes



# Simple Memory System Page Table

Only show first 16 entries (out of 256)

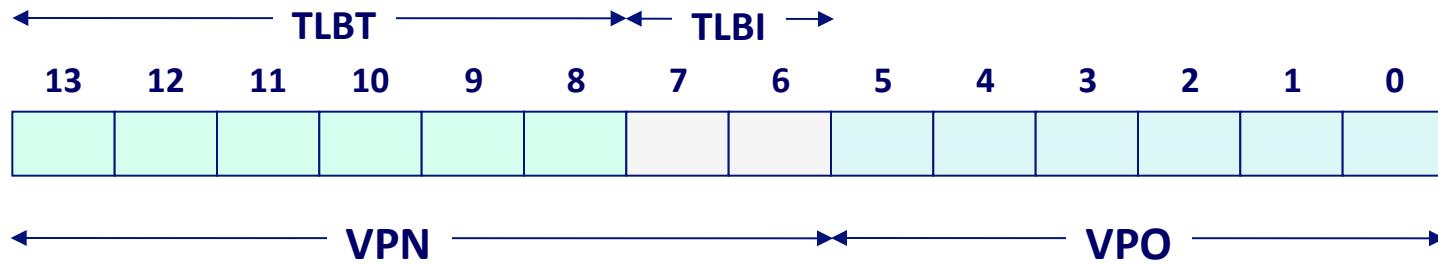
| <i>VPN</i> | <i>PPN</i> | <i>Valid</i> |
|------------|------------|--------------|
| 00         | 28         | 1            |
| 01         | –          | 0            |
| 02         | 33         | 1            |
| 03         | 02         | 1            |
| 04         | –          | 0            |
| 05         | 16         | 1            |
| 06         | –          | 0            |
| 07         | –          | 0            |

| <i>VPN</i> | <i>PPN</i> | <i>Valid</i> |
|------------|------------|--------------|
| 08         | 13         | 1            |
| 09         | 17         | 1            |
| 0A         | 09         | 1            |
| 0B         | –          | 0            |
| 0C         | –          | 0            |
| 0D         | 2D         | 1            |
| 0E         | 11         | 1            |
| 0F         | 0D         | 1            |

For a later example, let's assume VPN page 0x2E is invalid,  
i.e. valid bit =0, so it is not in memory

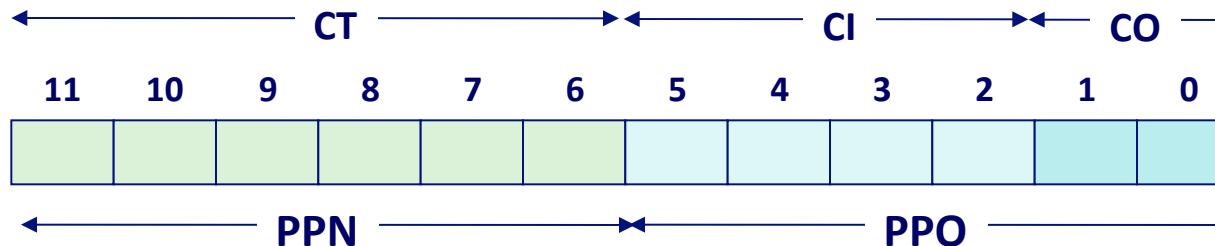
# Simple Memory System TLB

- 16 entries
- Virtually addressed
- 4-way associative



| <i>Set</i> | <i>Tag</i> | <i>PPN</i> | <i>Valid</i> | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 0          | 03         | -          | 0            | 09         | 0D         | 1            | 00         | -          | 0            | 07         | 02         | 1            |
| 1          | 03         | 2D         | 1            | 02         | -          | 0            | 04         | -          | 0            | 0A         | -          | 0            |
| 2          | 02         | -          | 0            | 08         | -          | 0            | 06         | -          | 0            | 03         | -          | 0            |
| 3          | 07         | -          | 0            | 03         | 0D         | 1            | 0A         | 34         | 1            | 02         | -          | 0            |

# Simple Memory System Cache

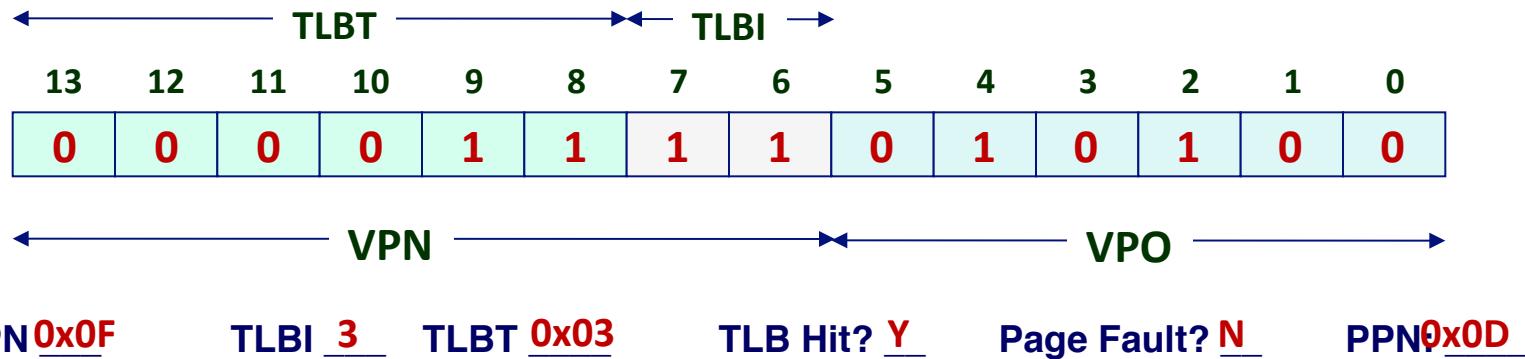


- **16 lines, 4-byte block size**
- **Physically addressed**
- **Direct mapped**

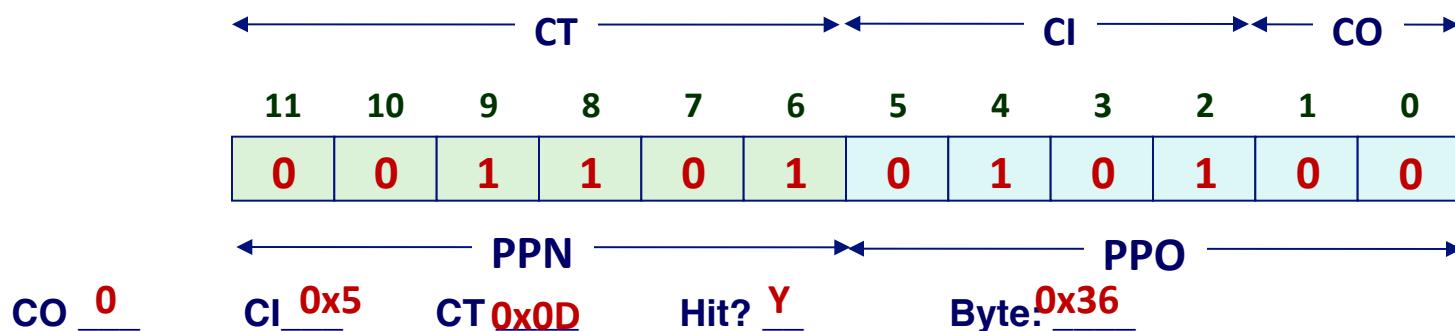
| <i>Idx</i> | <i>Tag</i> | <i>Valid</i> | <i>B0</i> | <i>B1</i> | <i>B2</i> | <i>B3</i> |
|------------|------------|--------------|-----------|-----------|-----------|-----------|
| 0          | 19         | 1            | 99        | 11        | 23        | 11        |
| 1          | 15         | 0            | -         | -         | -         | -         |
| 2          | 1B         | 1            | 00        | 02        | 04        | 08        |
| 3          | 36         | 0            | -         | -         | -         | -         |
| 4          | 32         | 1            | 43        | 6D        | 8F        | 09        |
| 5          | 0D         | 1            | 36        | 72        | F0        | 1D        |
| 6          | 31         | 0            | -         | -         | -         | -         |
| 7          | 16         | 1            | 11        | C2        | DF        | 03        |
| 8          | 24         | 1            | 3A        | 00        | 51        | 89        |
| 9          | 2D         | 0            | -         | -         | -         | -         |
| A          | 2D         | 1            | 93        | 15        | DA        | 3B        |
| B          | 0B         | 0            | -         | -         | -         | -         |
| C          | 12         | 0            | -         | -         | -         | -         |
| D          | 16         | 1            | 04        | 96        | 34        | 15        |
| E          | 13         | 1            | 83        | 77        | 1B        | D3        |
| F          | 14         | 0            | -         | -         | -         | -         |

# Address Translation Example #1

Virtual Address: 0x03D4

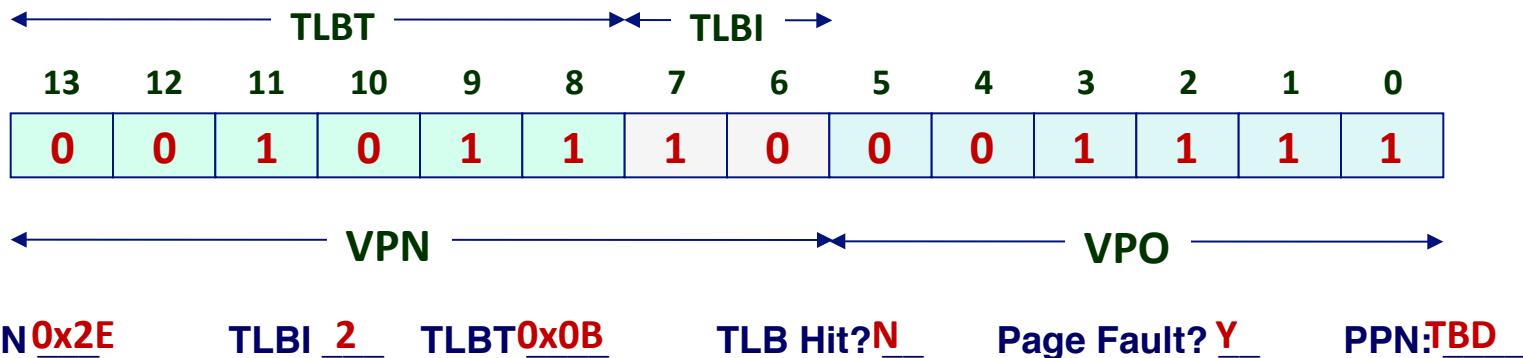


Physical Address



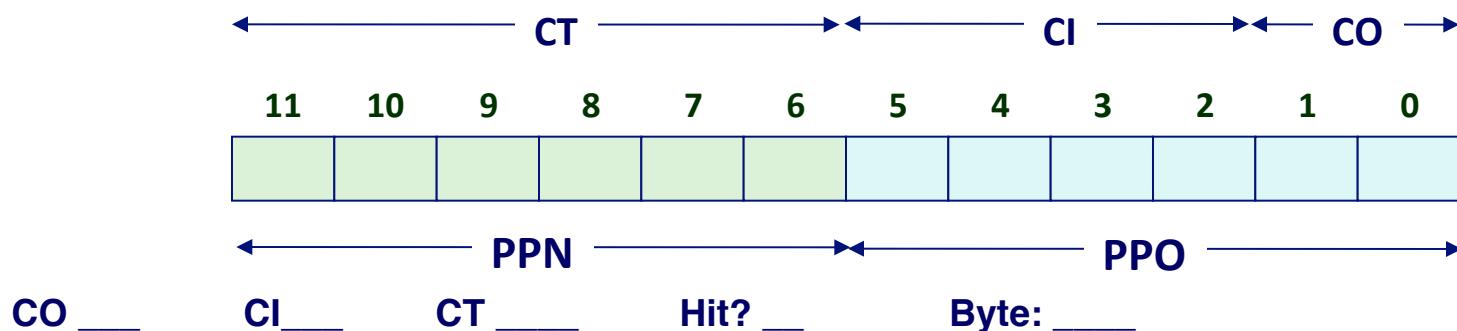
# Address Translation Example #2

Virtual Address: 0x0B8F



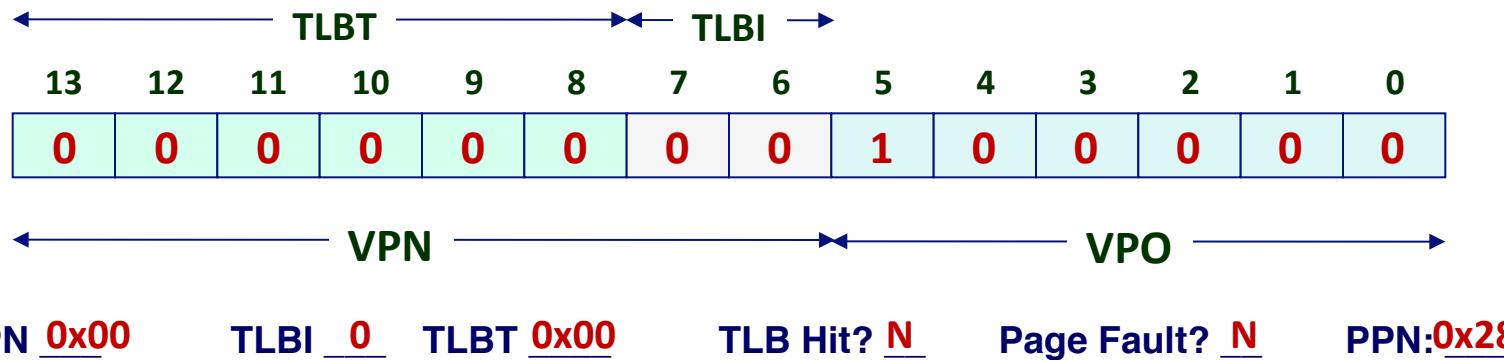
Physical Address

here, the page is neither in TLB nor memory,  
so have to go out to disk – MMU waits, i.e.  
valid bits of both TLB and page table =0

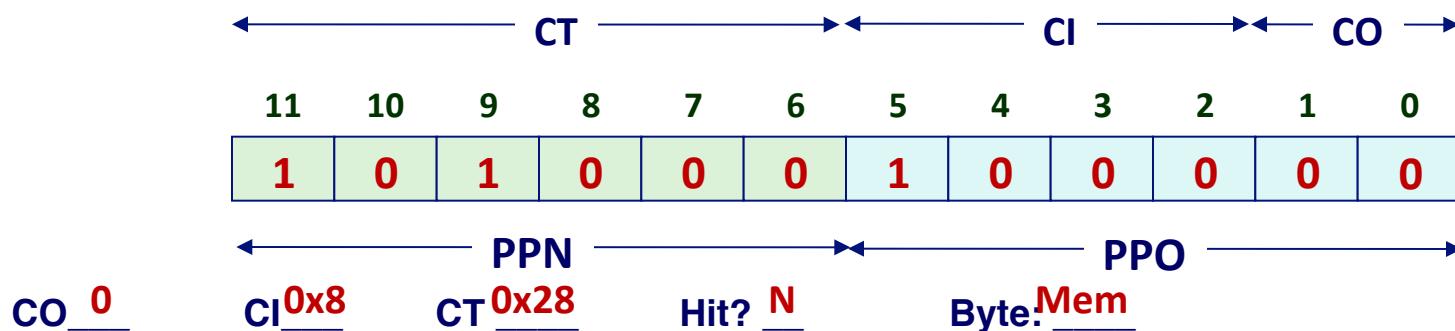


# Address Translation Example #3

Virtual Address: 0x0020



Physical Address



# Cache Hit Combinations

Is Virtual->Physical page mapping entry cached in Page Table?

- Hit on TLB
- Miss on TLB, Hit on Main Memory/RAM
- Miss on TLB, Miss on Main Memory/RAM = Page Fault

Is data/instruction associated with a Physical address cached in L1/L2 Cache?

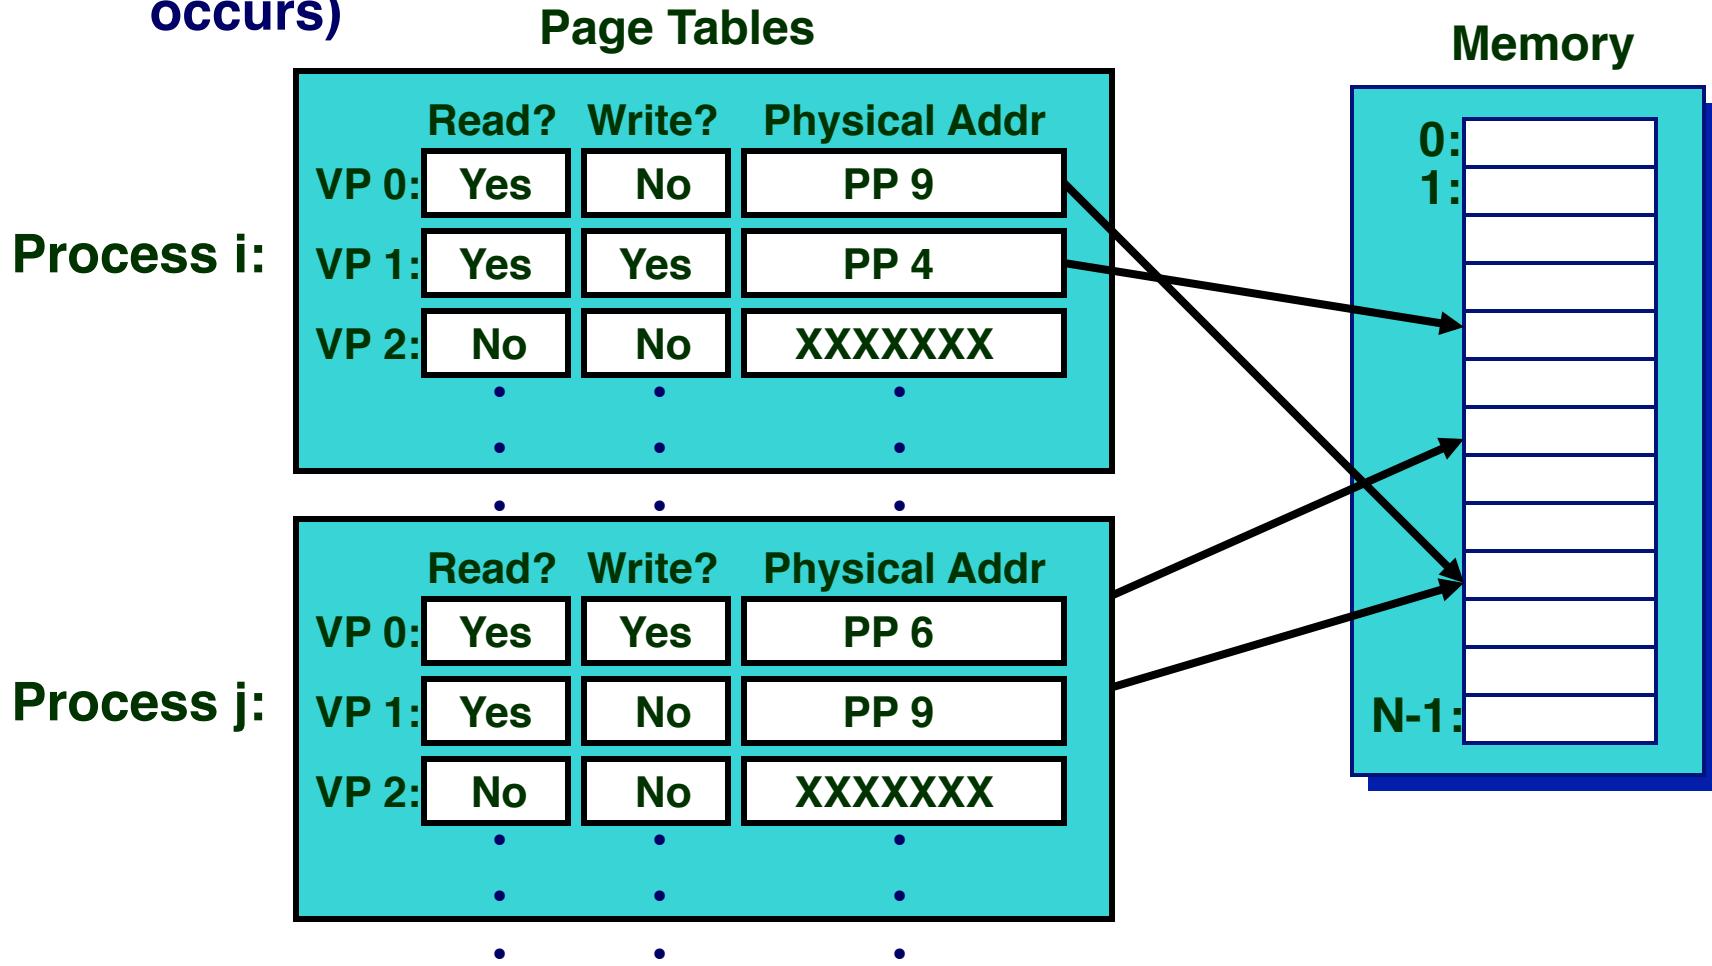
- Hit on L1
- Miss on L1, Hit on L2
- Miss on L1 and L2, go to Main Memory/RAM

# Supplementary Slides

# Memory Protection

Page table entry contains access rights information

- hardware enforces this protection (trap into OS if violation occurs)



# Motivations for Virtual Memory

For *multiple processes*, virtual memory brings benefits:

## Use Physical DRAM as a Cache for the Disk

- Only “active” code and data is actually cached in memory for each process
  - Allocate more memory to process as needed.
- Physical memory can support many more processes

## Simplify Memory Management

- Each process has its own virtual address space
- OS keeps track of the virtual->physical memory mapping and informs MMU whenever a new process is context-switched in to execute on the CPU

## Provide Protection

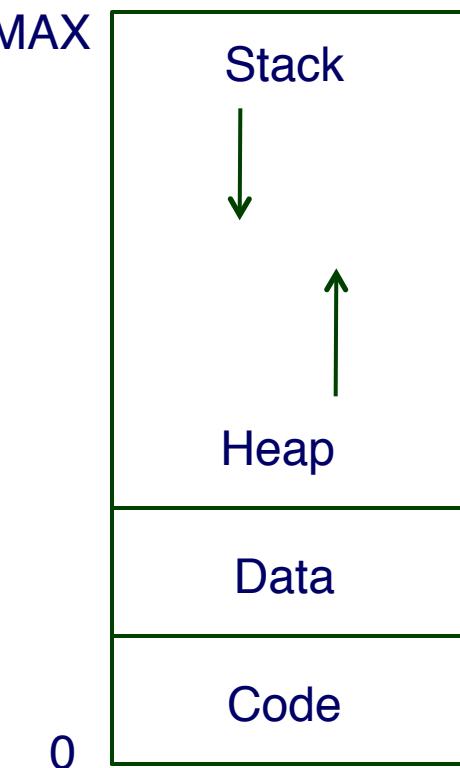
- One process can't interfere with another, because each one operates in its own address space & is mapped to different physical memory – can't overwrite same physical memory
  - different sections of address spaces have different permissions.

# Motivations for Virtual Memory (2)

Provides the Abstraction that a Process is executing in its own Address Space of memory from address 0 to address MAX

For a Single Process:

- Set of virtual addresses = Address Space of a process, which can *exceed physical memory size*
  - Only “active” code and data is actually in memory
  - frees a process from having to worry about physical limitations of main memory
    - for example, could allocate very large arrays and data structures, only a portion of which are kept in memory
  - Still limited by maximum size of virtual memory: for a 32-bit system, maximum addressable virtual memory is  $2^{32} = 4\text{GB}$



# Motivation #1: DRAM a “Cache” for Disk

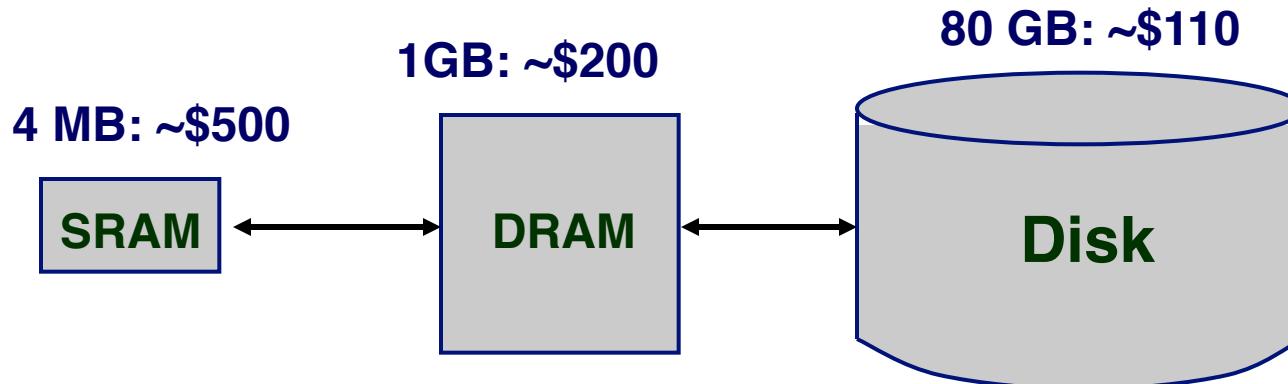
Full address space is quite large:

- 32-bit addresses: ~4,000,000,000 (4 billion) bytes
- 64-bit addresses: ~16,000,000,000,000,000,000 (16 quintillion) bytes

Disk storage is ~300X cheaper than DRAM storage

- 80 GB of DRAM: ~ \$33,000
- 80 GB of disk: ~ \$110

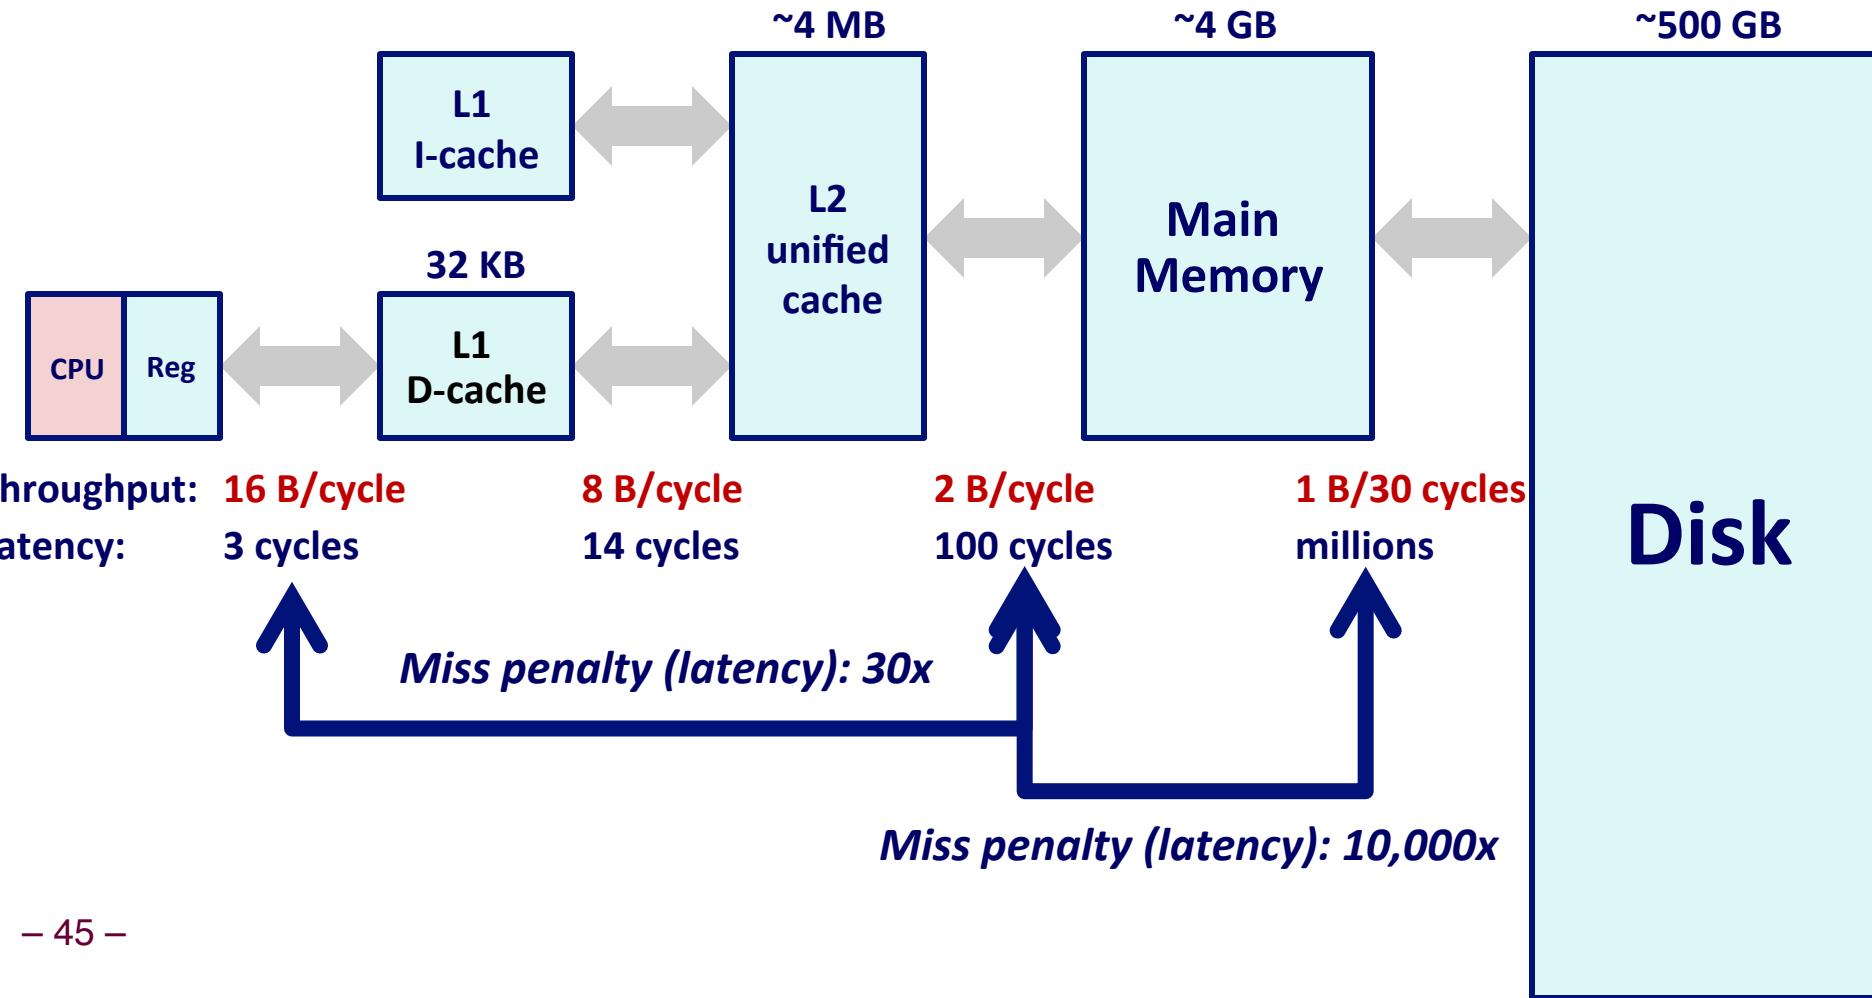
To access large amounts of data in a cost-effective manner, the bulk of the data must be stored on disk



# Memory Hierarchy: Core 2 Duo

*Not drawn to scale*

L1/L2 cache: 64 B blocks



# DRAM vs. SRAM as a “Cache”

**DRAM vs. disk is more extreme than SRAM vs. DRAM**

- Access latencies:

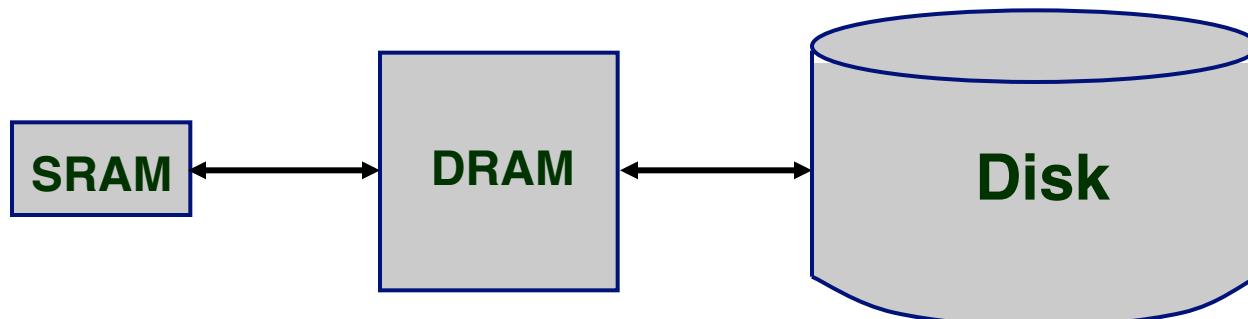
- DRAM ~10X slower than SRAM
- Disk ~100,000X slower than DRAM

- Importance of exploiting spatial locality:

- First byte is ~100,000X slower than successive bytes on disk
  - » vs. ~4X improvement for page-mode vs. regular accesses to DRAM

- Bottom line:

- Design decisions made for DRAM caches driven by enormous cost of misses



# Impact of Properties on Design

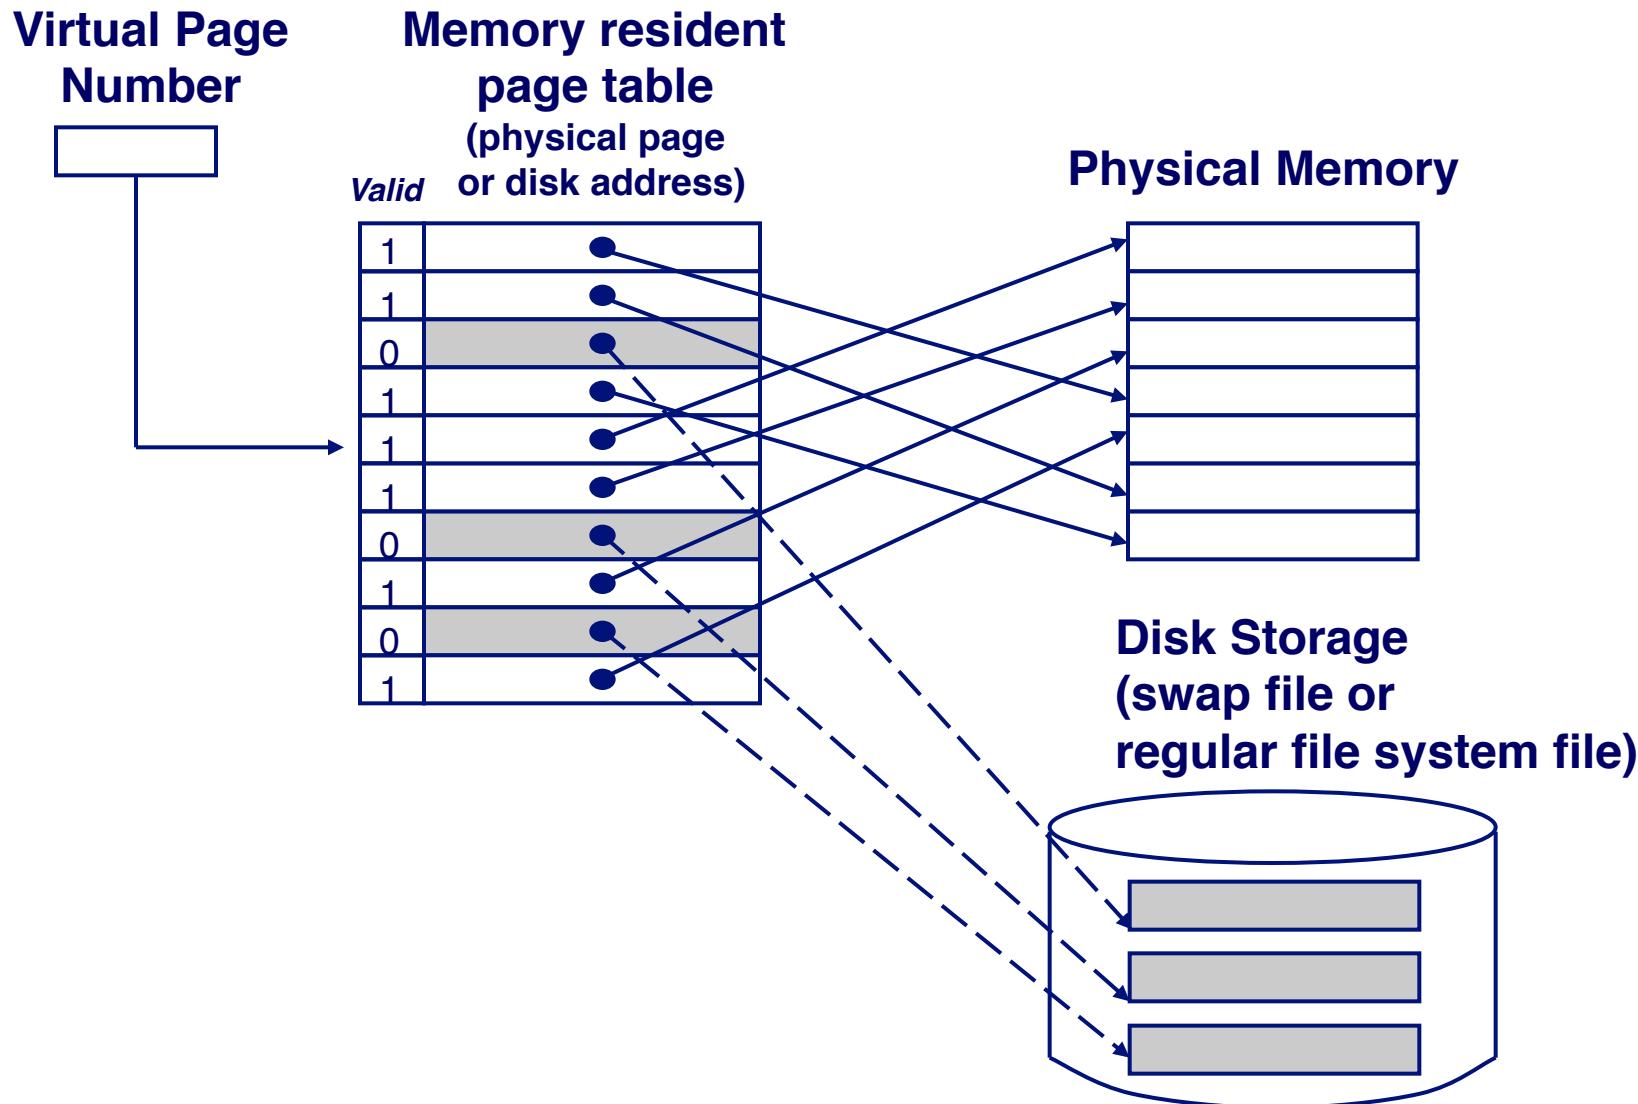
If DRAM was to be organized similar to an SRAM cache, how would we set the following design parameters?

- Line size?
  - Large, since disk better at transferring large blocks
- Associativity?
  - High, to minimize miss rate
- Write through or write back?
  - Write back, since can't afford to perform small writes to disk

What would the impact of these choices be on:

- miss rate
  - Extremely low. << 1%
- hit time
  - Must match cache/DRAM performance
- miss latency
  - Very high. ~20ms
- tag storage overhead
  - Low, relative to block size

# Page Tables



# VM Address Translation

## Virtual Address Space

- $V = \{0, 1, \dots, N-1\}$

## Physical Address Space

- $P = \{0, 1, \dots, M-1\}$
- $M < N$

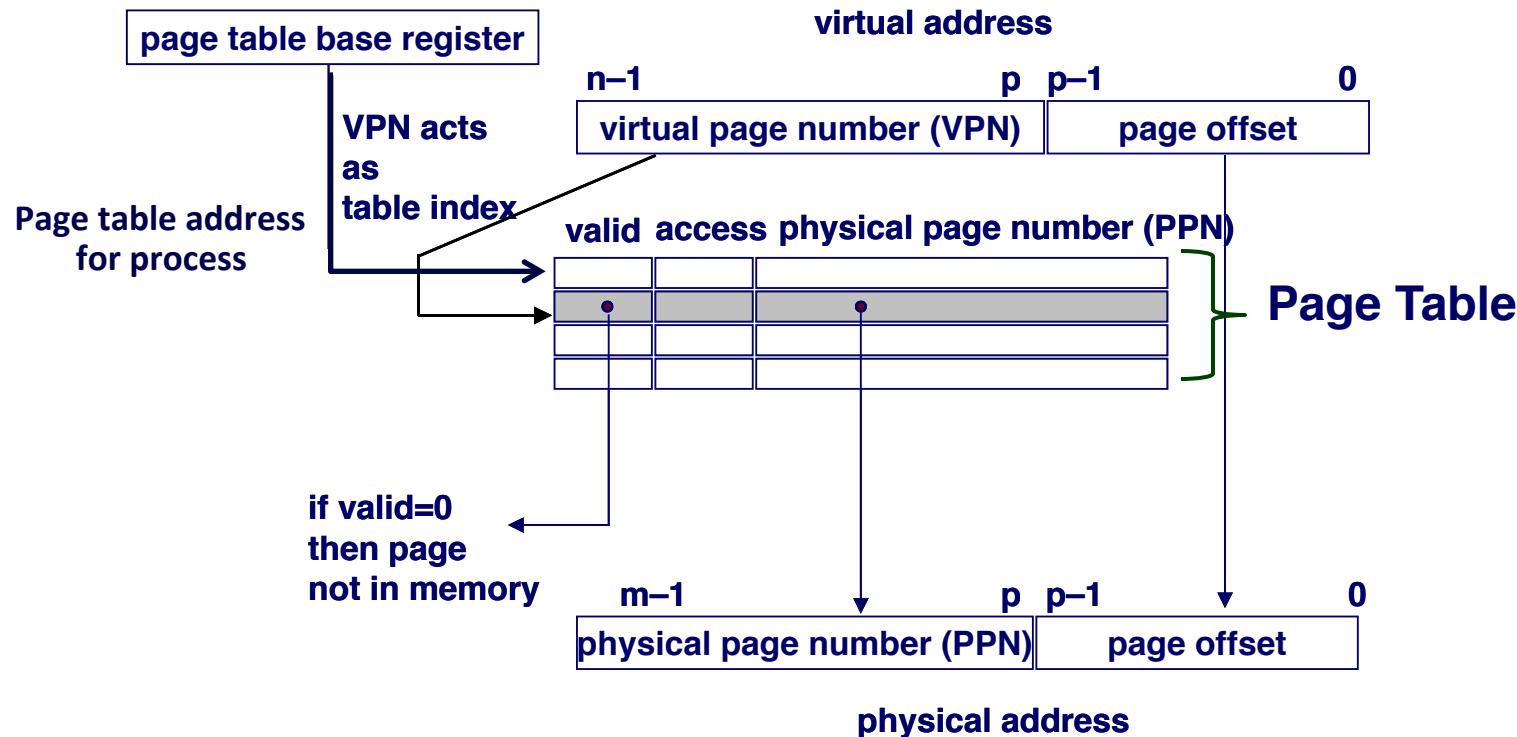
## Address Translation

- MAP:  $V \rightarrow P \cup \{\emptyset\}$
- For virtual address  $a$ :
  - $\text{MAP}(a) = a'$  if data at virtual address  $a$  at physical address  $a'$  in  $P$
  - $\text{MAP}(a) = \emptyset$  if data at virtual address  $a$  not in physical memory
    - » Either invalid or stored on disk

# Page Table Operation

## Translation

- Separate (set of) page table(s) per process
- VPN forms index into page table (points to a page table entry)



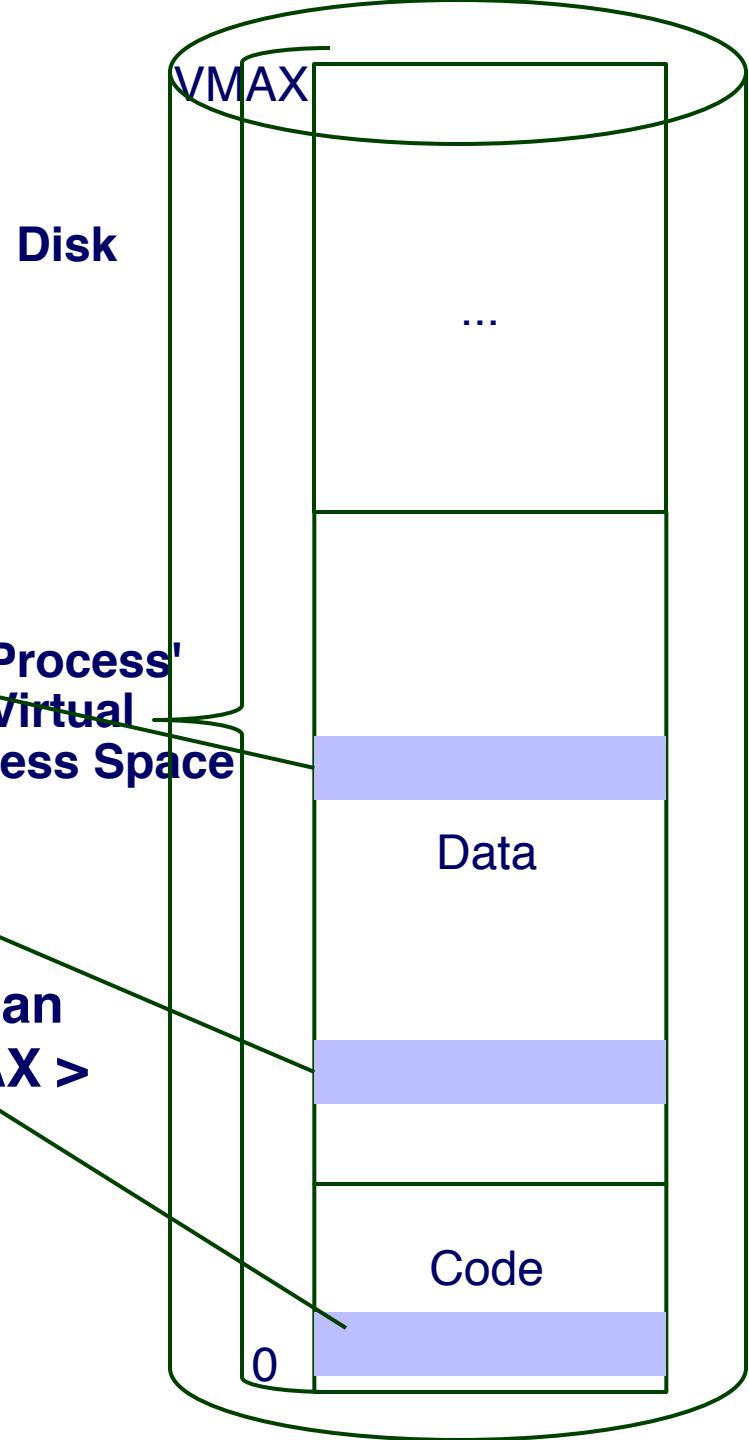
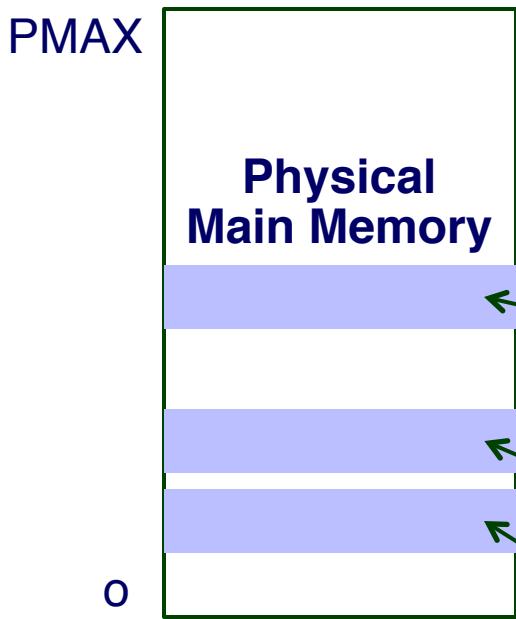
# Demand Paging uses Main Memory as a Cache for Disk

- The most recently accessed pages are kept in memory for fast access, compared to the disk, where the rest of the pages are stored
- Similar to caches, which keep the most recently accessed instructions and data in L1, L2, and L3 caches for fast access, compared to memory, where the rest of the instructions & data are stored

# Why does Caching Pages work? Locality

- Virtual memory works because of locality
- At any point in time, programs tend to access a set of active virtual pages called the **working set**
  - Programs with better temporal locality will have smaller working sets
- If (working set size < amount of physical main memory allocated to a process)
  - Good performance for one process after compulsory misses
- If (  $\text{SUM}(\text{working set sizes}) > \text{main memory size}$  )
  - **Thrashing:** Performance meltdown where pages are swapped (copied) in and out continuously

# Main Memory as a Cache for Disk



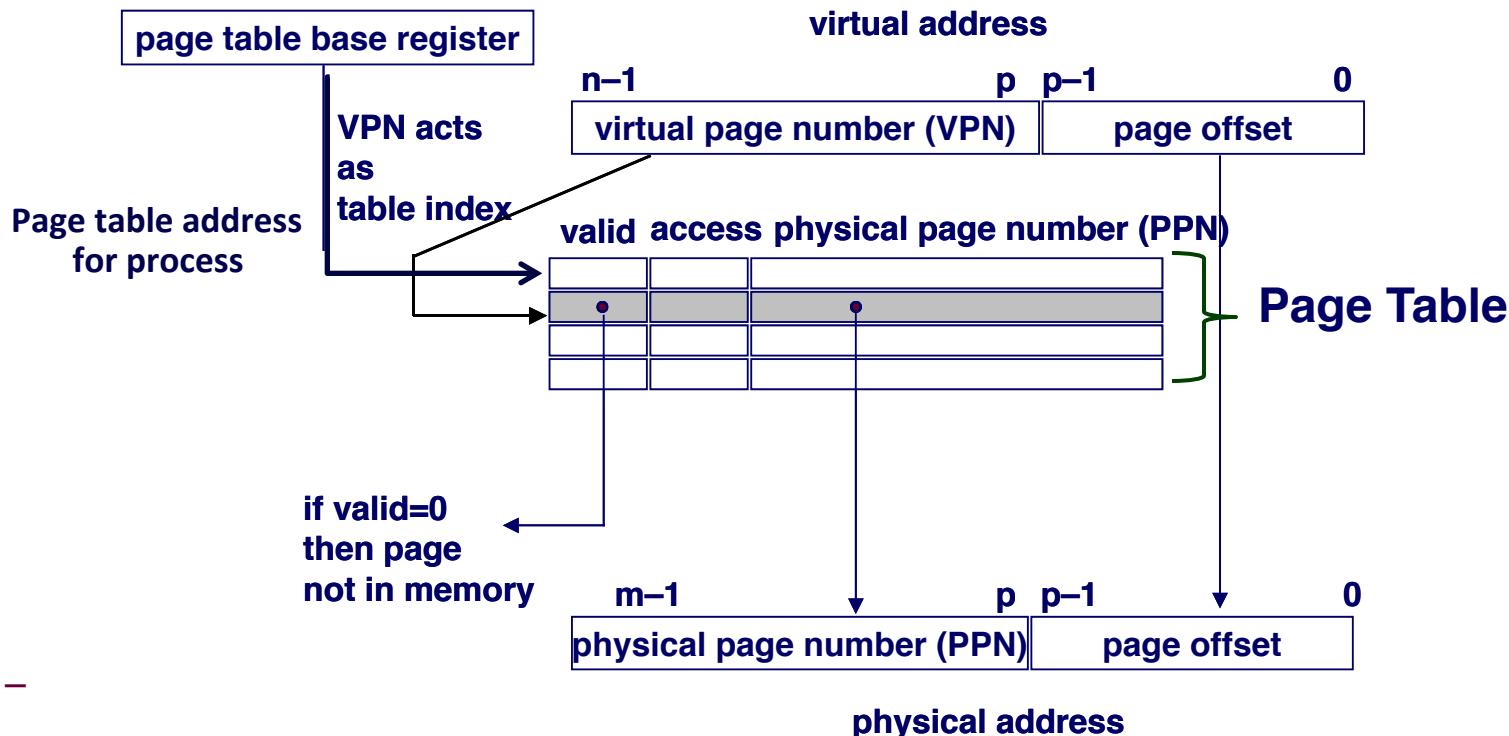
- Virtual address space of each process can be much larger than main memory ( $VMAX > PMAX$ )

- Most of this is stored on disk
- Only “active” code and data is actually **cached** in memory for fast execution and access

# Page Table Operation with Demand Paging

## Computing Physical Address

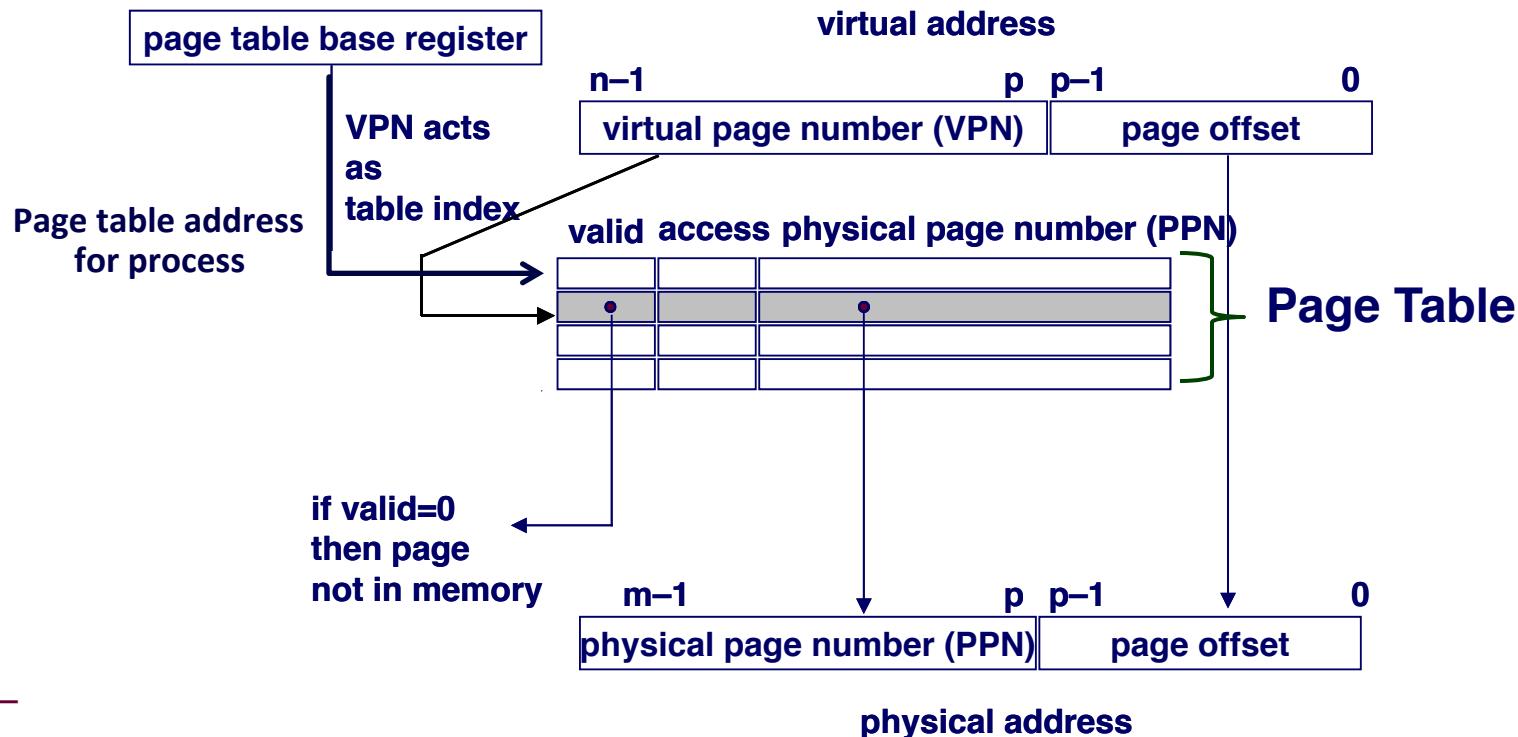
- if (valid bit = 1) then the page is in memory.
  - » Use physical page number (PPN) to construct address
- if (valid bit = 0) then the page is on disk
  - » Page fault – must load page from disk into a free page in memory and update page table



# Page Table Operation

## Checking Protection

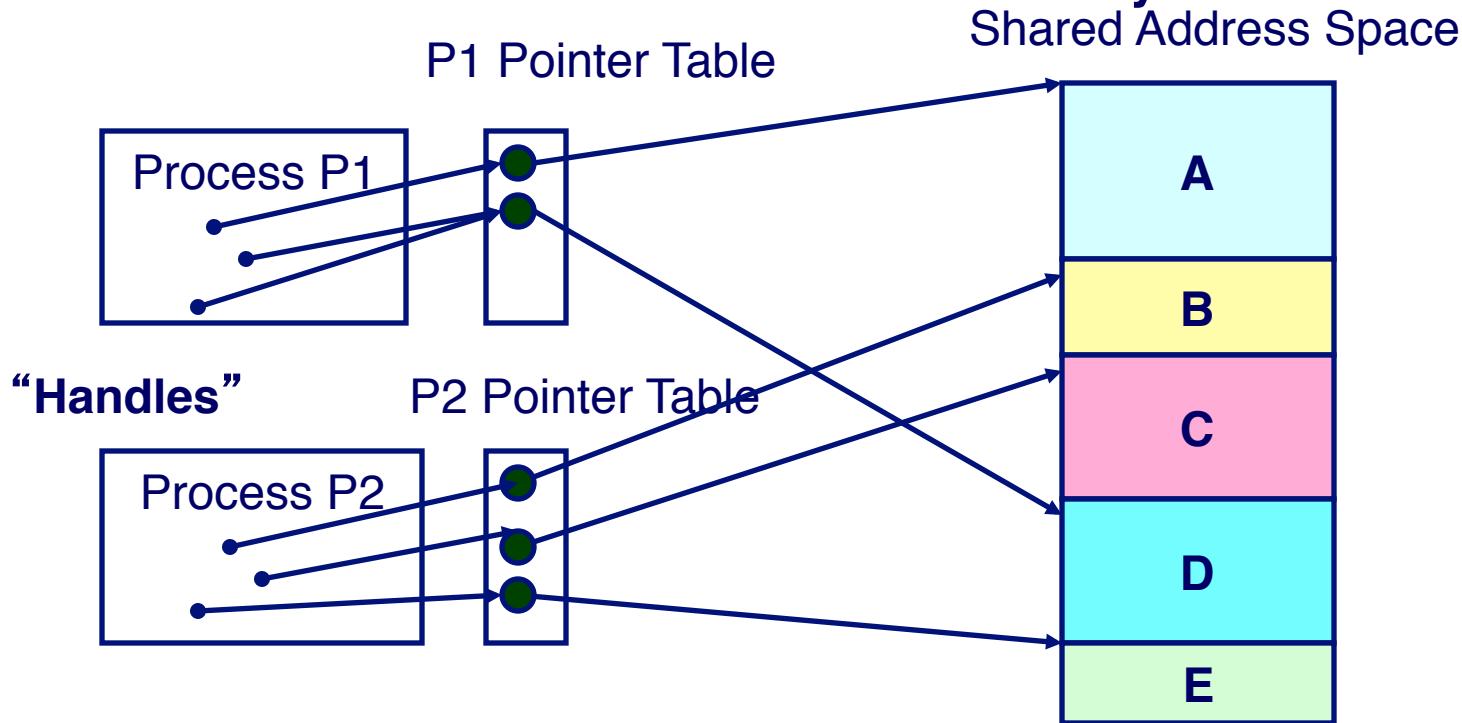
- Access rights field indicate allowable access
  - e.g., read-only, read-write, execute-only
  - typically support multiple protection modes (e.g., kernel vs. user)
- Protection violation fault if user doesn't have necessary permission



# Contrast: Macintosh Memory Model

## MAC OS 1–9

- Does not use traditional virtual memory



All program objects accessed through “handles”

- Indirect reference through pointer table
- Objects stored in shared global address space

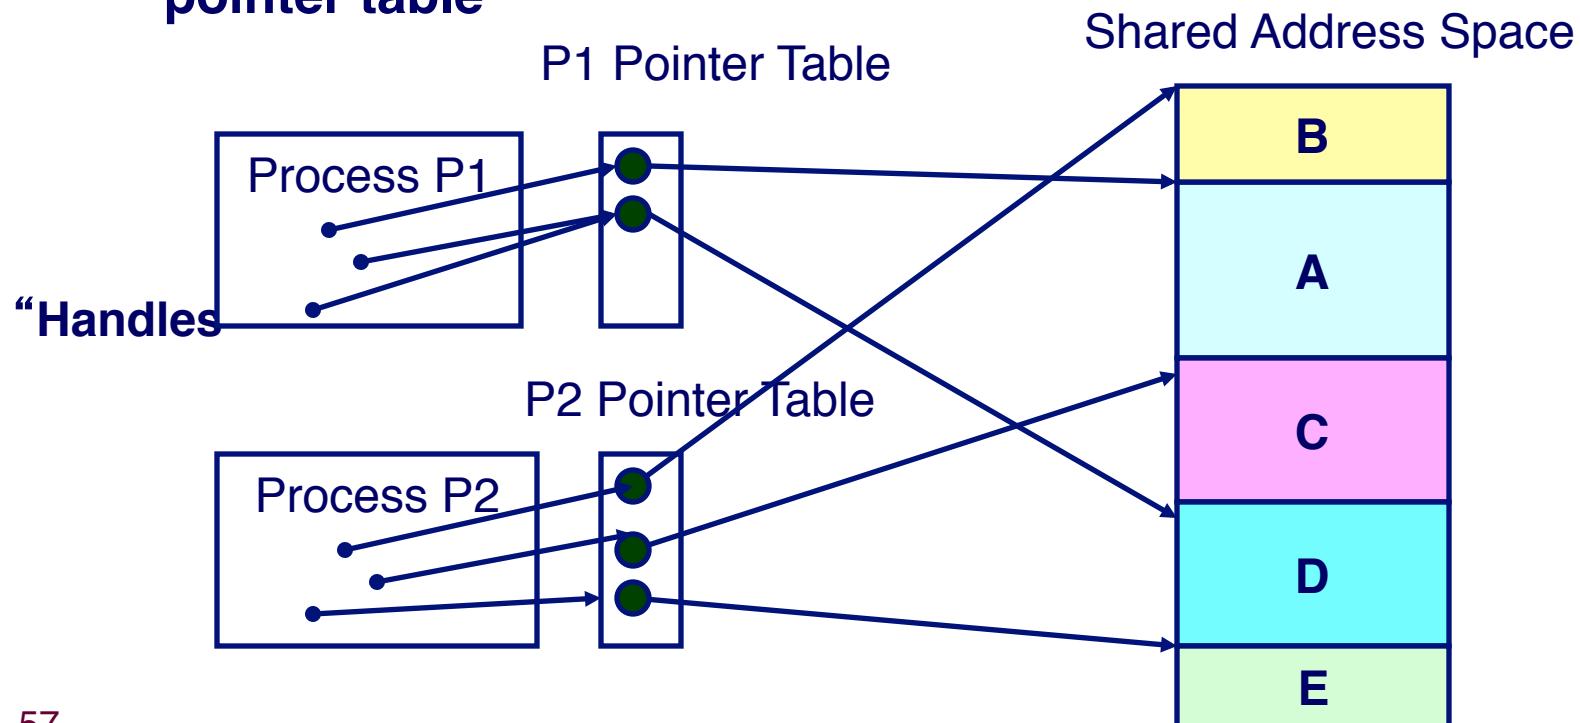
# Macintosh Memory Management

## Allocation / Deallocation

- Similar to free-list management of malloc/free

## Compaction

- Can move any object and just update the (unique) pointer in pointer table



# Mac vs. VM-Based Memory Mgmt

## Allocating, deallocating, and moving memory:

- can be accomplished by both techniques

## Block sizes:

- Mac: variable-sized
  - may be very small or very large
- VM: fixed-size
  - size is equal to *one page* (4KB on x86 Linux systems)

## Allocating contiguous chunks of memory:

- Mac: contiguous allocation is *required*
- VM: can map contiguous range of virtual addresses to disjoint ranges of physical addresses

## Protection

- Mac: “wild write” by one process can corrupt another’s data

# Modern O/S

## “Modern” Operating System

- Virtual memory with protection
- *Preemptive multitasking*

## Windows/NT: Based on VMS from Digital Equipment

- Primary designer (Cutler) from Digital Equipment
- “Heavy” processes, event queues & messages

## Mac OS/X: Based on FreeBSD & MACH OS

- FreeBSD derived from “Berkeley Standard Distribution”
  - Contracted by DARPA to build a more robust UNIX
- MACH OS Developed at CMU in late 1980’s
  - Key idea was “micro-kernel” - limited functionality implemented by kernel