

# Chapter 9: Memory Management

## Topics

- Virtual Memory
- Heap Management

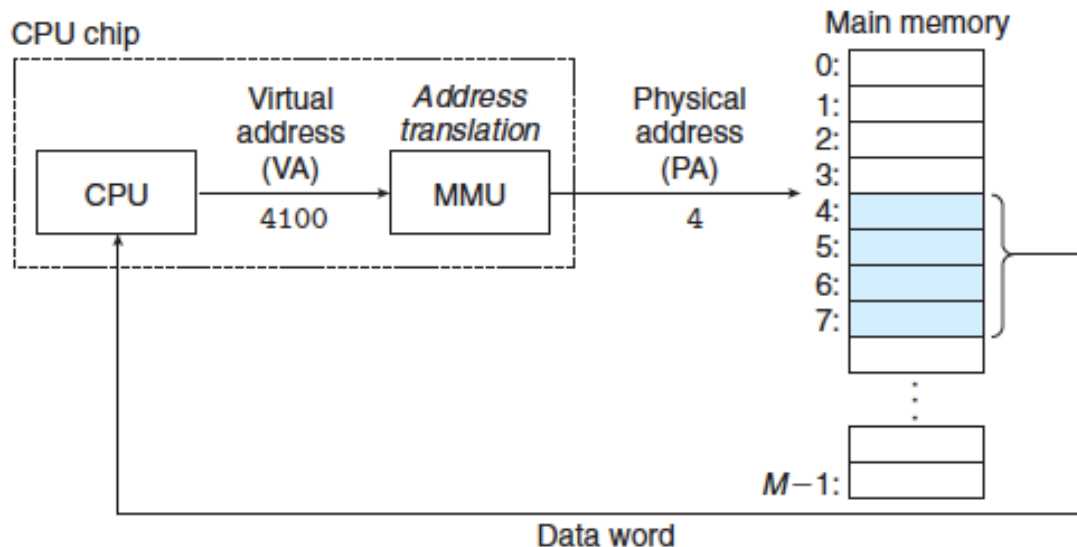
# Announcements

- **Shell lab is due Monday Dec 8 by 8 am**
  - Interview grading time slots for next week available later this week
  - TA office hours Thursday and Friday
- **Last Recitation Exercise #5 due Friday Dec 12 by 5 pm**
  - Upload to moodle or hand into TA at TA office hours
- **Final exam is Thursday Dec 18, 4:30-7 pm, more next week**
- **Reading:**
  - Read Chapter 9, except 9.6 and 9.7 (no case study and no memory mapping, can also skip multi-level page tables)
- **FCQs at end of class today**

# Recap

## Virtual Memory



- executables are compiled as if they would execute in their own virtual address space of memory addresses [0..VMAX]
  - code & data addresses are virtual
  - Many advantages to this approach
- a Memory Management Unit (MMU) translates each virtual addresses reference into a physical address, and access memory with that physical address



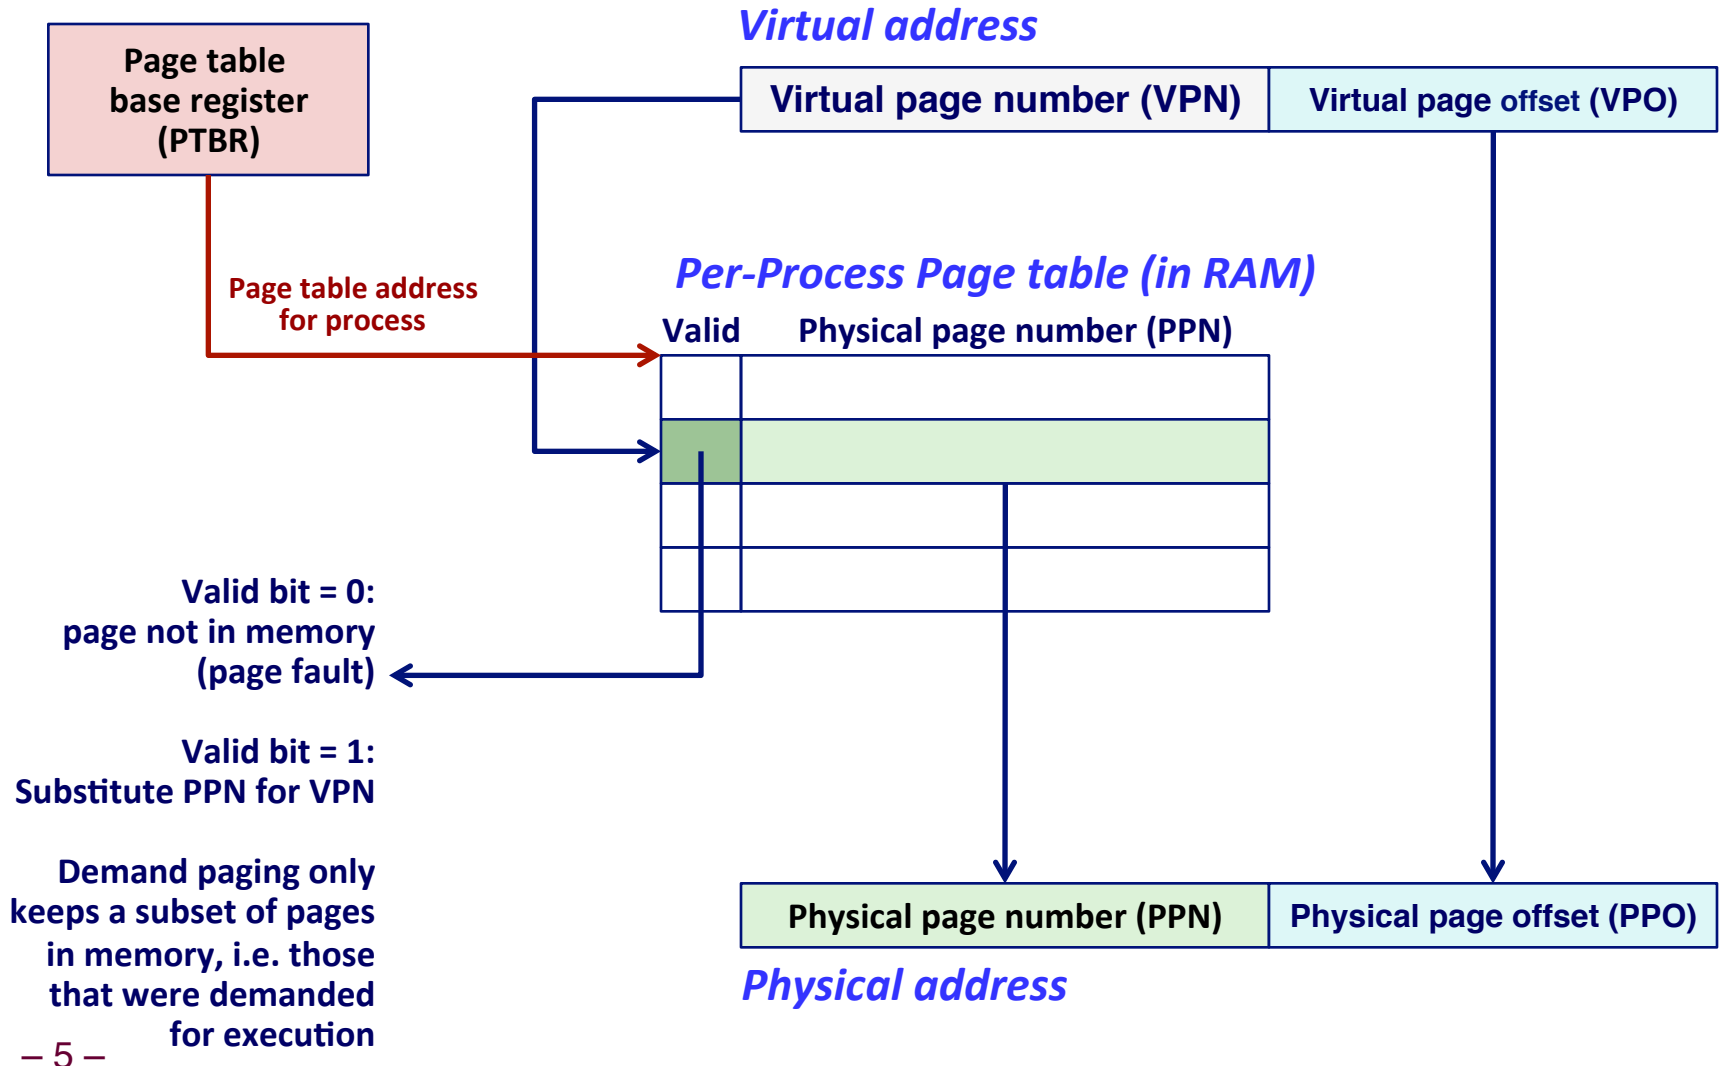
# Recap: Page Tables

valid bit indicates if page is in memory (more on this later)

Virtual Page #	Physical Page #	valid ?
0	3	1
1		0
2		0
3	1	1
4		0
5		0
6		0
7		0
8		0
9		0



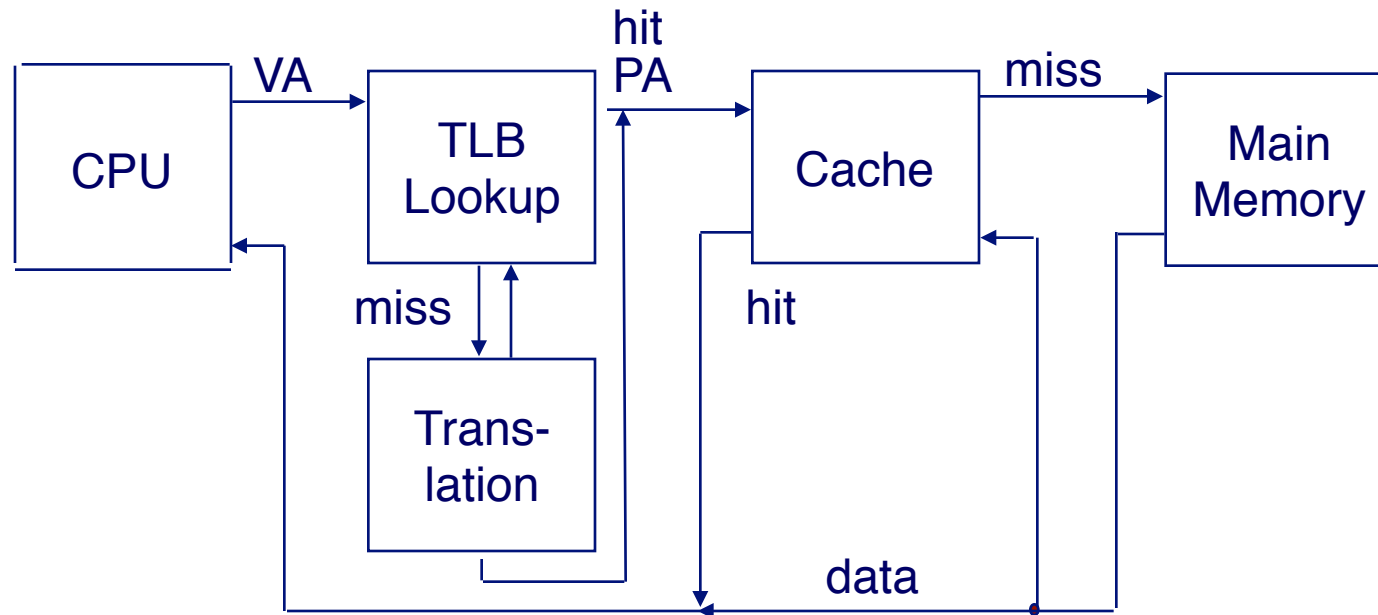
# Recap: Address Translation With a Page Table



# Recap: Speeding up Translation with a TLB

## “Translation Lookaside Buffer” (TLB)

- Small hardware cache in MMU
- Maps virtual page numbers to physical page numbers
- Contains complete page table entries for small number of pages




# Complete example of virtual memory

- See previous lecture's slides

# Dynamic Memory Allocation

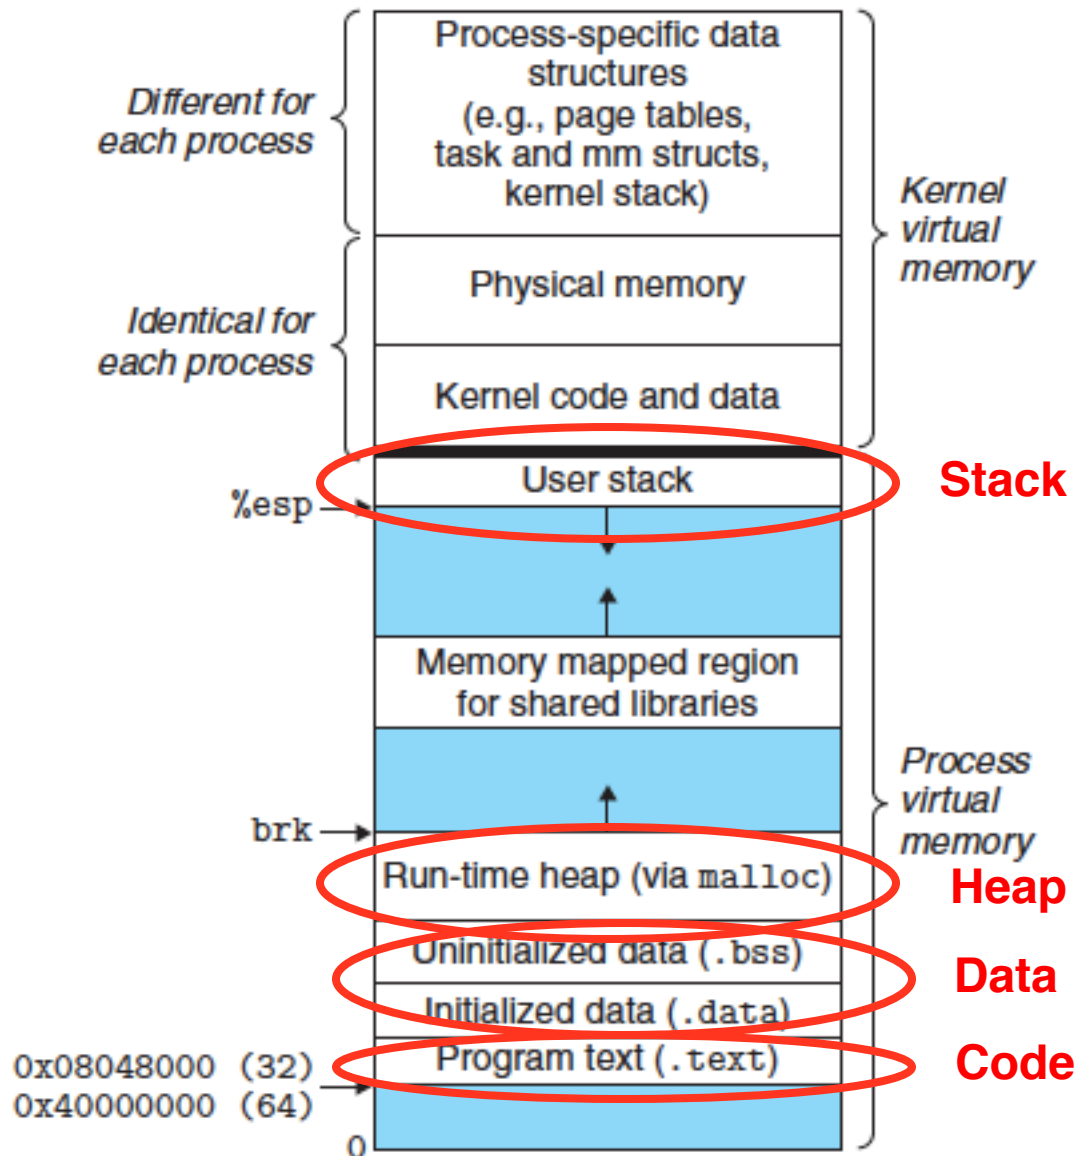
**Allocate variables dynamically at run time from the “Heap”**

- Grows upwards in terms of memory addresses
- May not know until run time how large of an array, linked list, or data structure to allocate
- Useful to dynamically expand the size of a data structure, e.g. a linked list or a binary tree, by allocating more memory as needed.





# Example: Linux Address Space



# Heap Allocation Example


```
int x=0,y[1000];  
char *p;  
  
main (int argc, char *argv[]) {  
  
    p = (char*) malloc(256);  
  
    function1(p);  
  
    free(p); /* return p to  
available memory pool */  
}  
  
function1 (char *ptr) {  
  
    int i,j=50;  
  
    for (i=0; i<100; i++) {  
        *(ptr+i) = i*j;  
    }  
  
}
```

- Global variables allocated in data section of address space


- Local variables allocated on the stack

- Dynamic variables allocated on the heap

- In C++, the “new” command is equivalent to malloc



# Dynamic Memory Allocation



## Explicit vs. Implicit Memory Allocator

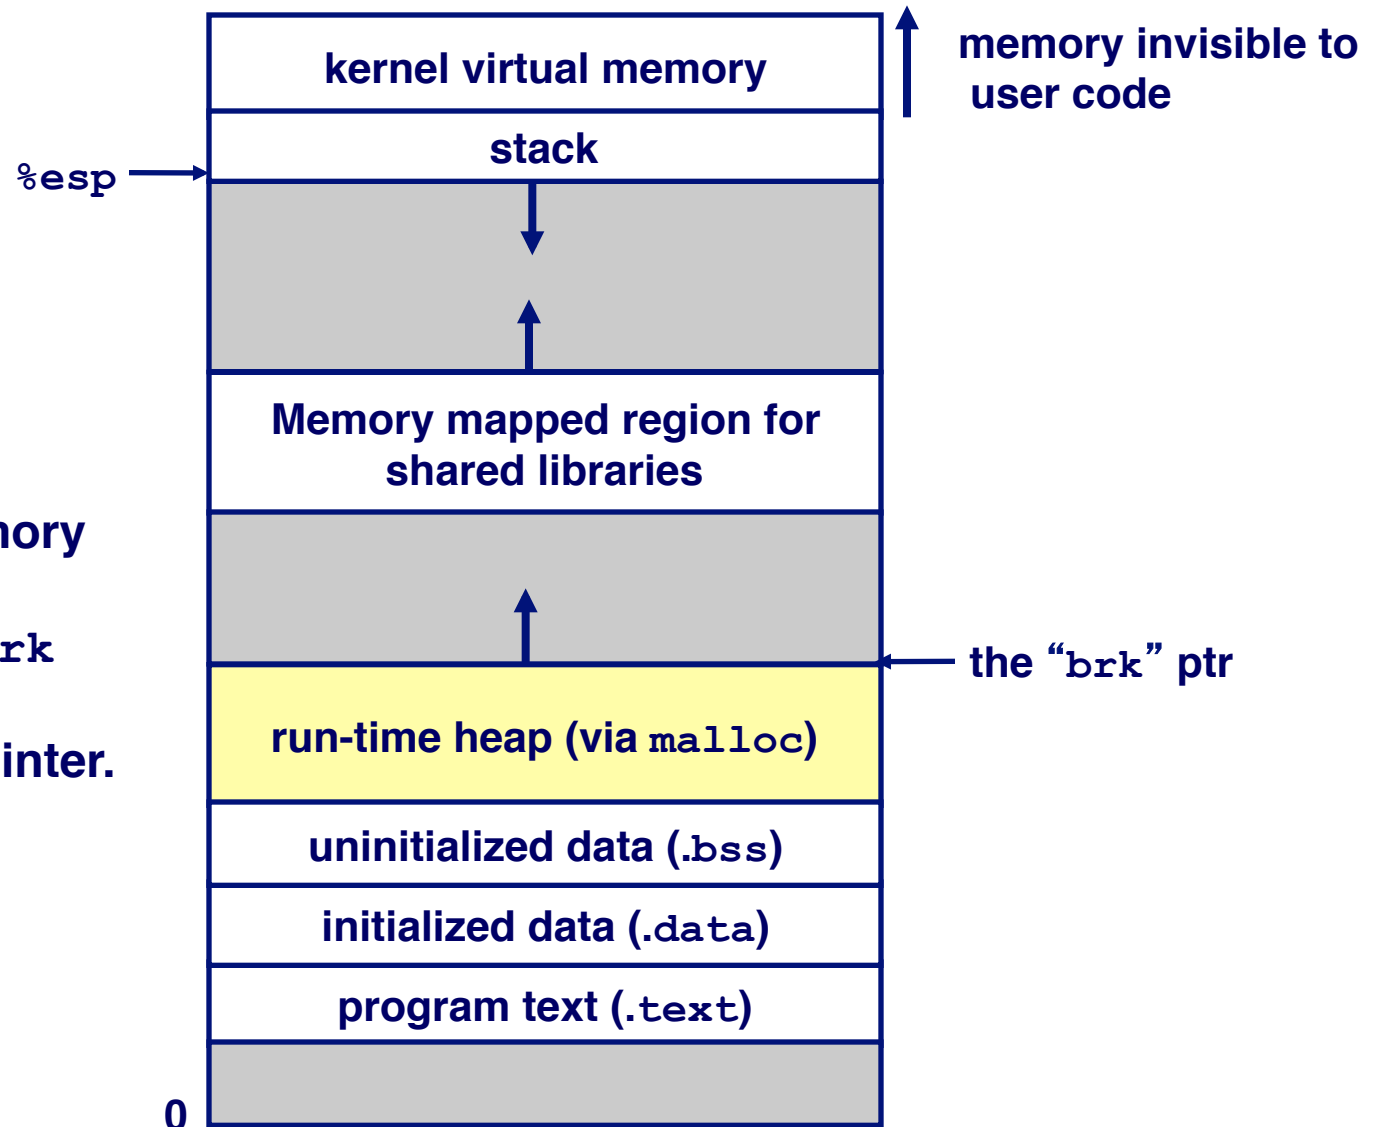
- **Explicit:** application allocates and frees space
  - E.g., `malloc` and `free` in C
  - In C++, the `new` command is equivalent to a `malloc`, and `delete` is equivalent to a `free`
- **Implicit:** application allocates, but does not free space
  - E.g. garbage collection in Java, ML or Lisp

## Allocation

- In both cases the memory allocator provides an abstraction of memory as a set of blocks
- Doles out free memory blocks to application

– 1 Will discuss explicit memory allocation first

# Process Memory Image



Allocators request additional heap memory from the operating system using the `sbrk` function, which increases the `brk` pointer. (can also shrink the heap)

# Malloc Package

```
#include <stdlib.h>
```

```
void *malloc(size_t size)
```

- If successful:

- Returns a pointer to a memory block of at least `size` bytes, (typically) aligned to 8-byte boundary.
- If `size == 0`, returns NULL

- If unsuccessful: returns NULL (0) and sets `errno`.

```
void free(void *p)
```

- Returns the block pointed at by `p` to pool of available memory
- `p` must come from a previous call to `malloc` or `realloc`.

```
void *realloc(void *p, size_t size)
```

- increases or decreases the size of the specified block of memory `p` and returns pointer to new block.
- Reallocates block if needed. Contents of new block unchanged up to min of old and new size.

```
void *calloc(size_t nelem, size_t elsize)
```

- 13 – ■ Similar to `malloc` except initializes block of memory to zero

# Malloc Example

```
void foo(int n, int m) {
    int i, *p;


    /* allocate a block of n ints */
    if ((p = (int *) malloc(n * sizeof(int))) == NULL) {
        perror("malloc");
        exit(0);
    }
    for (i=0; i<n; i++)
        p[i] = i;

    /* add m bytes to end of p block */
    if ((p = (int *) realloc(p, (n+m) * sizeof(int))) == NULL) {
        perror("realloc");
        exit(0);
    }
    for (i=n; i < n+m; i++)
        p[i] = i;

    /* print new array */
    for (i=0; i<n+m; i++)
        printf("%d\n", p[i]);

    free(p); /* return p to available memory pool */
}
```

# Allocation Examples




# External Fragmentation

Occurs when there is enough aggregate heap memory, but no single free block is large enough


```
p1 = malloc(4)
```



```
p2 = malloc(5)
```



```
p3 = malloc(6)
```



```
free(p2)
```



```
p4 = malloc(6)
```

**oops!**

External fragmentation depends on the pattern of future requests, and thus is difficult to measure.




# Internal Fragmentation

Poor memory utilization caused by *fragmentation*.

- Comes in two forms: internal and external fragmentation

## Internal fragmentation

- For some block, internal fragmentation is the difference between the block size and the payload size.



- Caused by overhead of maintaining heap data structures, padding for alignment purposes, or explicit policy decisions (e.g., not to split the block).
- Depends only on the pattern of previous requests, and thus is easy to measure.

# Goals of Good malloc/free

## Primary goals


1. **Efficient memory utilization**
    - User allocated structures should be large fraction of the heap.
    - Want to minimize “fragmentation”.
  2. **Low latency/fast throughput for `malloc` and `free`**
    - Ideally should take constant time (not always possible)
    - Should certainly not take linear time in the number of blocks
- Goals 1 and 2 are often conflicting!

## Some other goals

- **Good locality properties**
- Structures allocated close in time should be close in space
  - “Similar” objects should be allocated close in space
- **Robust**
- Can check that `free(p1)` is on a valid allocated object `p1`
  - Can check that memory references are to allocated space

# Implementation Issues

- How do we know how much memory to free given only a pointer?
- How do we keep track of the free blocks?
- What do we do with the extra space when allocating a structure that is smaller than the free block it is placed in?
- How do we pick a block to use for allocation -- many might fit?
- How do we reinsert a freed block?




`free(p0)`

`p1 = malloc(1)`

# Knowing How Much to Free

## Standard method


- Keep the length of a block in the word preceding the block.
  - This word is often called the *header field* or *header*
- Requires an extra word for every allocated block
- Let's assume each square in figure is large enough (say 4 bytes) to contain a word, e.g. an int or pointer




- 20 - Note if p0 in free(p0) is not originally malloc'ed, this may cause an error

# Keeping Track of Free Blocks

**Method 1: *Implicit free list* using lengths -- links all blocks**



**Method 2: *Explicit free list* among the free blocks using pointers within the free blocks**



**Method 3: *Segregated free list***

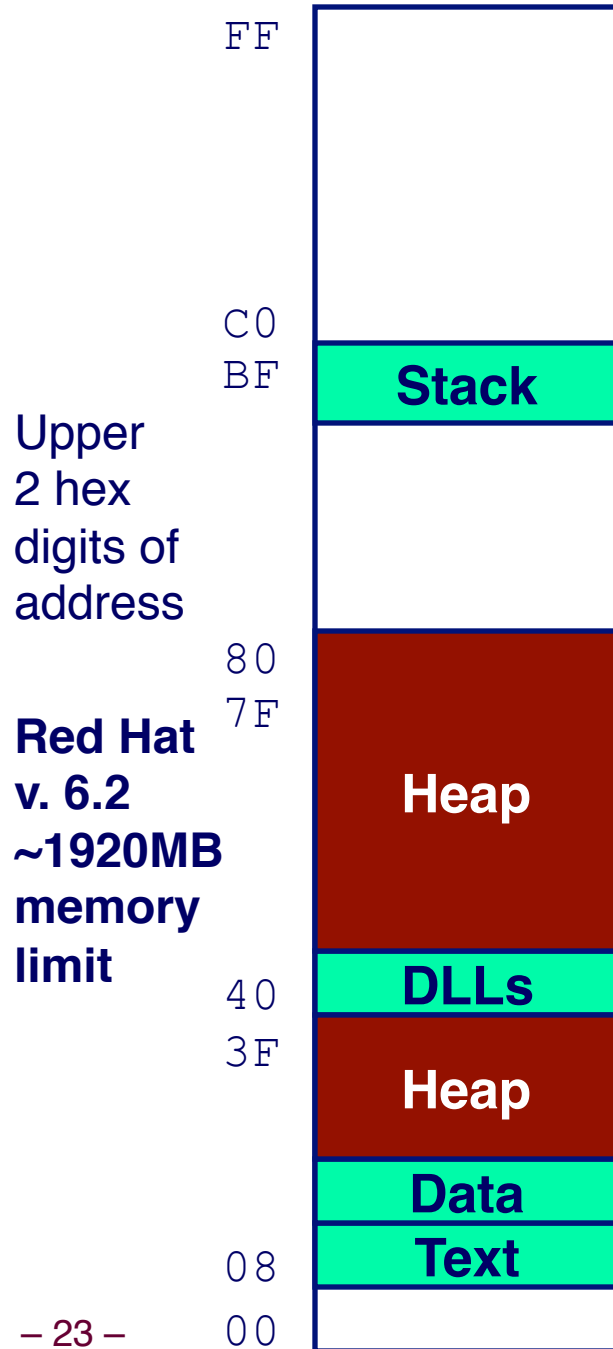
- Different free lists for different size classes

**Method 4: Blocks sorted by size**

- Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

# Supplementary Slides

# Linux Memory Layout



## Stack

- Runtime stack (8MB limit)

## Heap

- Dynamically allocated storage
- When call `malloc`, `calloc`, `new`

## DLLs

- Dynamically Linked Libraries
- Library routines (e.g., `printf`, `malloc`)
- Linked into object code when first executed

## Data


- Statically allocated data
- E.g., arrays & strings declared in code

## Text


- Executable machine instructions
- Read-only

# Linux Memory Allocation


**Initially**




**Linked**



**Some  
Heap**



**More  
Heap**





# Constraints

## Applications:

- Can issue arbitrary sequence of allocation and free requests
- Free requests must correspond to an allocated block


## Allocators

- Can't control number or size of allocated blocks
- Must respond immediately to all allocation requests
  - *i.e.*, can't reorder or buffer requests
- Must allocate blocks from free memory
  - *i.e.*, can only place allocated blocks in free memory
- Can't move the allocated blocks once they are allocated
  - *i.e.*, compaction is not allowed
- Must align blocks so they satisfy all alignment requirements
  - 8 byte alignment for GNU malloc (`libc` malloc) on Linux boxes
- Can only manipulate and modify free memory

# Assumptions

## Model the Heap memory as a sequence of words

- Memory is word addressed (each word can hold a pointer)



# Performance Goals: Throughput

Given some arbitrary sequence of malloc and free requests:

- $R_0, R_1, \dots, R_k, \dots, R_{n-1}$

Want to maximize throughput and peak memory utilization.

- These goals are often conflicting

Throughput:

- Number of completed requests per unit time
- Example:
  - 5,000 malloc calls and 5,000 free calls in 10 seconds
  - Throughput is 1,000 operations/second.

# Performance Goals: Peak Memory Utilization

Given some sequence of malloc and free requests:

- $R_0, R_1, \dots, R_k, \dots, R_{n-1}$

**Def: Aggregate payload  $P_k$ :**

- `malloc(p)` results in a block with a *payload* of  $p$  bytes..
- After request  $R_k$  has completed, the *aggregate payload*  $P_k$  is the sum of currently allocated payloads.

**Def: Current heap size is denoted by  $H_k$**

- Assume that  $H_k$  is monotonically nondecreasing

**Def: Peak memory utilization:**

- After  $k$  requests, *peak memory utilization* is:
  - $U_k = (\max_{i \leq k} P_i) / H_k$